

The Art of Fluid Animation



Jos Stam



CRC Press

Taylor & Francis Group

AN A K PETERS BOOK

The Art of Fluid Animation

The Art of Fluid Animation

Jos Stam



CRC Press

Taylor & Francis Group
Boca Raton London New York

CRC Press is an imprint of the
Taylor & Francis Group, an **informa** business
AN A K PETERS BOOK

CRC Press
Taylor & Francis Group
6000 Broken Sound Parkway NW, Suite 300
Boca Raton, FL 33487-2742

© 2016 by Taylor & Francis Group, LLC
CRC Press is an imprint of Taylor & Francis Group, an Informa business

No claim to original U.S. Government works

Printed on acid-free paper
Version Date: 20150921

International Standard Book Number-13: 978-1-4987-0020-7 (Paperback)

This book contains information obtained from authentic and highly regarded sources. Reasonable efforts have been made to publish reliable data and information, but the author and publisher cannot assume responsibility for the validity of all materials or the consequences of their use. The authors and publishers have attempted to trace the copyright holders of all material reproduced in this publication and apologize to copyright holders if permission to publish in this form has not been obtained. If any copyright material has not been acknowledged please write and let us know so we may rectify in any future reprint.

Except as permitted under U.S. Copyright Law, no part of this book may be reprinted, reproduced, transmitted, or utilized in any form by any electronic, mechanical, or other means, now known or hereafter invented, including photocopying, microfilming, and recording, or in any information storage or retrieval system, without written permission from the publishers.

For permission to photocopy or use material electronically from this work, please access www.copyright.com (<http://www.copyright.com/>) or contact the Copyright Clearance Center, Inc. (CCC), 222 Rosewood Drive, Danvers, MA 01923, 978-750-8400. CCC is a not-for-profit organization that provides licenses and registration for a variety of users. For organizations that have been granted a photocopy license by the CCC, a separate system of payment has been arranged.

Trademark Notice: Product or corporate names may be trademarks or registered trademarks, and are used only for identification and explanation without intent to infringe.

Visit the Taylor & Francis Web site at
<http://www.taylorandfrancis.com>

and the CRC Press Web site at
<http://www.crcpress.com>

To Pam and Gillian



This is a small replica of the northern portal of the Urnes Stavkyrkje in Norway. This church was built in 1130, near the end of the Viking domination of Northern Europe. Vikings were highly skilled artists, not just long-haired brutish warriors. I took this photograph at the British Museum in London, England.

Contents

Preface, xi

Acknowledgments, xix

Author, xxi

CHAPTER 1 ■ Introduction	1
CHAPTER 2 ■ Observations, Equations, and Numbers	11
2.1 BEAVERS, CAVE PERSONS, AND FIRE	11
2.2 FROM CAVES TO GREEKS: ARCHIMEDES, GOLD, AND MEDALS	13
2.3 LONG, CURLY HAIRRED MATHEMATICIANS, THE ABYSS AND THE AIRBRUSH	16
CHAPTER 3 ■ Euler–Newton Equations or Navier–Stokes Equations	35
3.1 LEONARDO DA VINCI	35
3.2 EULER AND CONTINUITY	36
3.3 INCOMPRESSIBILITY, CONTINUITY, HELMHOLTZ, AND HODGE THEORY	44
3.4 EULER AND THE MOTION OF FLUIDS	49
3.5 NEWTON AND VISCOSITY	55
3.6 NAVIER AND STOKES AND THEIR EQUATIONS	64
3.7 BOUNDARIES, BOUNDARIES, BOUNDARIES, BOUNDARIES, AND BOUNDARIES	69

CHAPTER 4 ■ The Early Days of Computational Fluid Dynamics	81
CHAPTER 5 ■ Kolmogorov and Turbulence	91
CHAPTER 6 ■ Introduction to Fluid Animation	97
6.1 DISCRETIZE! BUGS, GRIDS, AND BUGS MOVING THROUGH GRIDS	98
6.1.1 Bugs	99
6.1.2 Grids	103
6.1.3 Bugs Moving through Grids	107
6.1.4 Semi-Lagrangian	108
6.1.5 PIC	112
CHAPTER 7 ■ Intermezzi	117
7.1 INTERMEZZO UNO: LINEAR SYSTEMS	118
7.2 INTERMEZZO DUE: THE GENERAL SOLUTION OF A LINEAR SYSTEM	124
7.3 INTERMEZZO TRE: CIRCULANT MATRICES AND THE FOURIER TRANSFORM	132
7.4 INTERMEZZO QUATTRO: NUMERICAL SOLUTION OF LINEAR SYSTEMS	141
CHAPTER 8 ■ A Simple Fluid Solver	155
8.1 A MATH HORROR FLICK: OPERATOR SPLITTING	159
8.2 CODE PLEASE?	160
8.2.1 Moving Densities	164
8.2.2 Moving Fluids That Change Themselves: Nonlinearity	170
8.3 BUGS CRAWLING ON DONUTS, THE FFT, AND ~60 LINES OF C CODE	176
8.4 FOUR-DIMENSIONAL TURBULENT VECTOR FIELDS AND TURBULENCE	186
8.5 DECORATING FLUIDS	189

CHAPTER 9 ■ The Little Computers Who Can Handle Fluids	195
CHAPTER 10 ■ The Smart Phones That Can Handle Fluids	199
CHAPTER 11 ■ Fluid FX: Version 2.0 of Autodesk Fluid	203
CHAPTER 12 ■ Show Time! MAYA Fluid Effects	205
CHAPTER 13 ■ Fluids on Arbitrary Surfaces	213
CHAPTER 14 ■ Control Freaks! How to Make Fluids Do What We Want	219
14.1 SHOOTING CANNONBALLS IN TWO DIMENSIONS	222
14.2 COMPUTER OPTIMIZERS	223
14.3 THE AD, THE JOINT, AND THE PATH BACK TO THE OPTIMIZER	233
CHAPTER 15 ■ Real Experiments, Computer Experiments, and Validation	239
15.1 SPHERES ARE SUCH A DRAG	240
15.2 CURLY FLOWS BEHIND SPHERES, WAVY FLOWS IN TUBES, AND TURBULENT PLUMES BETWEEN PLATES	242
15.2.1 Experiment Number One: Curly Flows behind Spheres	243
15.2.2 Experiment Number Two: Wavy Flows in Tubes	243
15.2.3 Experiment Number Three: Turbulent Plumes between Plates	245
CHAPTER 16 ■ Epilogue: Let's Call It Quits	249
INDEX, 251	

Preface

What I cannot create, I do not understand.

RICHARD FEYNMAN (FAMOUS AMERICAN PHYSICIST
AND NOBEL LAUREATE)

Lectures which really teach will never be popular; lectures which are popular will never teach.

MICHAEL FARADAY (FAMOUS ENGLISH
EXPERIMENTALIST AND SCIENTIST)

A mathematical theory is not to be considered complete until you have made it so clear that you can explain it to the first man whom you meet on the street.

DAVID HILBERT (FAMOUS GERMAN MATHEMATICIAN)

I have no formal background in fluid dynamics. I am not an engineer nor do I have a specialized degree in the mathematics or physics of fluids. I am fortunate that I did not have to carry that baggage around. On the other hand, I *do* have degrees in pure mathematics and computer science and have an artsy background. More importantly, I have written computer code that animates fluids.*

I wrote code. That is the bottom line.

* Other stuff as well of course. Chess games, Pac-Man rip offs, writing code to make money during summer jobs, ray tracers, particle system simulations, surface modeling, and more recently a unified dynamics solver called *Nucleus*. That is the beauty of writing code. The computer can be taught to do all sorts of things and you can do it anywhere. All you need is a laptop with an Internet connection. Sort of like mathematics. A mathematician is basically a clever mechanism that turns coffee into theorems. Coding and mathematical work can be done anywhere: the Copacabana Beach of Rio Janeiro (Brazil) or a remote cabin in Northern Canada.

I did not just download some code from the Internet and mash it together. Therefore, I think I understand what I am talking about. Of course, it is based on previous work. I did not invent the theories and concepts behind it. Research is a process.

You cannot fool a computer.

You can fool students, colleagues, or friends but not a computer. Why? A computer is like the best pet ever, it is wickedly fast and always obeys. Creating computer code is the ultimate test whether you understand something or not. Teach it to the computer and you will understand it. If you work for a company, you also have to deal with marketing, public relations, sales, and the customer feedback cycle.

Customer feedback is brutally honest. They pay you money for your software and they expect it to work. If it doesn't work, they will let you know. It either works or it doesn't. Postmodern literature debates, on the other hand, are not like that. Everyone is right and everyone is wrong at the same time. But these debates fueled by coffee or wine might be more fun than spending hours fueled by coffee fixing computer code that has to work.

I like writing code. And I like fixing code as long as it is my code.

This book is written by a dilettante of some sorts. However, by writing computer code, I came to understand the dynamics of fluids. This book is not your usual fluid dynamics textbook full of clever equations. I know equations quite well and I love them. They helped me to understand fluids. But to most people, they look like strange hieroglyphs from some other universe.

That is how Chinese characters appear to me.

That is why fancy equations are not very helpful to most people. When I get an e-mail written entirely in Chinese characters, I just ignore it. What? The alternative is that I take a course to learn Chinese characters. That would be cool and help me the next time I travel to China. But that is going to be a lot of work and my time is limited. Besides, it is fun to be a stranger in a strange land.

I want this book to be accessible to people who are not experts in mathematics or physics. I also want to make this a fun book. That is how I do research. I like to have fun even at *work*. Some parts of this book have mathematics. My hope is that the math will be somewhat accessible to most people who are willing to *go along*. This is a common tactic in mathematics when reading a mathematical paper that is not in your area of

expertise and trying to understand it. At first, you just try to get the gist of it. Then if it sounds cool, you go through the details. And if you really get excited, then you should try to write some code that implements the content of the paper.

I also provide one-paragraph summaries of the material after each section in this book. I found this to be a good practice when I was learning math and computer science at the university. What did I learn in a nutshell after a lecture? I picked up this practice from a math professor, Pierre de la Harpe, at the University of Geneva who started each lecture with a summary of the material he covered in the previous lectures.

This methodology really helps if you are taking tons of classes. I also tend to understand some material using different tools than what most people are used to. I like to repeat things from different points of view. I like to argue with myself. That way, I can be both right and wrong, just not at the same time. You can learn some interesting stuff from this process. It is also a good tactic before giving a talk. You will be ready to face most questions. No question is stupid, and some can point you in new directions. I am always open to exploring new ideas, and learning. You can never be the smartest or the most creative person in the room.

The main goal of this book is to show how to create computer code that animates the motion of fluids. Computer code will be included in this book. Readers can download the accompanying code and run it on their own computers.

My goal is actually more ambitious. I want programmers to use these codes as a starting point to create their own apps, games, and so on. In fact, some of my code has been available for over 10 years and many people have used it as a starting point to create their own games, fun demos, and apps. Even better would be if programmers rewrite the code in a completely different manner or in a completely different language. That would be so cool. I would like to challenge any reader to write a shorter version of my code in C that is still readable.

I did not want to write a *Fluid Dynamics for Dummies* style of book. This is because I think my readers are smart and creative people who want to know how fluids are simulated on a computer to create nifty animations. In fact, I want smart and creative programmers to read this book and extend the code to create novel applications. Basically, I want this book to inspire people to do their own stuff. To create and not take anything for granted.

Be a rebel.

Not to destroy, but to create.

I love it when I get an e-mail out of the blue from someone far from my home in Toronto, like India, pointing me to a web-based application they created that combines fluids and reaction-diffusion processes. And “wow” the program runs in a web page. And the person thanks me, too. How cool is that. My day is certainly made.

Since there are so many good technical books on fluid dynamics, I want this book to be different and less technical. I want this book to be a bridge between my two favorite fluid dynamics books:

An Album of Fluid Motion (Stanford University, California, 1982), assembled by Milton Van Dyke*

A Mathematical Introduction to Fluid Mechanics (Springer, 2000), written by Alexandre Chorin and Jerrold Marsden

The first book is a collection of photographs of famous fluid experiments. It is the perfect geek coffee table book.

The second book is a brilliant and concise mathematical introduction to fluid dynamics. This book clarified many obscure aspects of fluid dynamics for me through a rigorous mathematical treatment. Consider the following statement: “Pressure is the Lagrange multiplier derived from the divergence free constraint.” If this statement makes sense to you, then Chorin and Marsden’s book is for you. But please continue reading this book.

This book is also a personal account of how I have dealt with fluid dynamics. Obscure at first, leading to enlightenment followed by writing computer code.

I want to make the narrative of this book somewhat interesting. Research is not just a collection of impersonal facts that seem to come out of nowhere. The science of fluids was created or discovered, whatever, by real people. I think it is important to give homage to them.

* I have to tell this story. I ordered this book from Milton Van Dyke’s website and did not have to pay for it when I ordered it. Only when I got the book through regular mail, did I see that there was a notice to send a check of some ridiculously low amount of money to a certain Milton Van Dyke living in Palo Alto, California. No one I know would rip off a person who put together such a beautiful book. Of course, I immediately sent the check to the address with “Thanks!” written on the back of the check.

Don't worry.

This book is not going to describe how many wives, husbands, lovers, or kids these scientists had. There are plenty of good books out there documenting their personal lives. I am not an expert in these matters anyway nor am I particularly interested in their extracurricular activities. It is all about their scientific achievements.*

This book is based on many talks I have given over the last 15 years. Which explains the somewhat informal colloquial style of the book. Also, I assume my readers have access to search engines as there is no exhaustive list of references. If you like something in the book and want to know more about it, research it. That is how I work these days. I do miss the days I had to go to libraries, however, that is just nostalgia.

The first time I gave a talk on this subject was at an annual computer graphics conference called SIGGRAPH[†] in Los Angeles in 1999. The paper I presented was called "Stable Fluids." That was a crazy conference for me. I had other talks on completely different subjects,[‡] one in the same afternoon, and I had to lug my hardware, an SGI Octane, using a cab from my hotel room at the Westin Bonaventure up on Figueroa Avenue to the convention center at the bottom of Figueroa.

But it worked out.

This was one of the first times that fluid simulations were shown to react to user input in real time. There was applause. And then there was applause again two years later, again in Los Angeles, when I gave my demonstration of real-time fluids on a Pocket PC.[§] I wrote this just for fun. The Pocket PC fits in your pocket (hence the name). Consequently, I was able to show my fluid creations everywhere: at parties, on the subway, or to my family living abroad in Europe.

Only later did people tell me that I disproved some skeptics, that my fast demonstrations two years earlier were only due to using fancy hardware.

* But still. Leonhard Euler, one of the heroes of this book, had 13 kids. Legend has it that he did his finest math while holding one of his babies on his lap.

[†] SIGGRAPH is an acronym for Special Interest Group on GRAPHics and Interactive Techniques. Since 1974, this conference has been held yearly in various places across North America. It is the Mecca for graphics guys. It is the most prestigious place to publish a paper in computer graphics. In 1999, there were over 40,000 attendees.

[‡] "Subdivision Surfaces" and "Diffraction Shaders."

[§] The Pocket PC was released by Microsoft in 2000 and powered by an operating system called Windows CE. I showed my demo on an iPAQ, which was created by Compaq at the time. In 2007, Apple came out with the iPhone/iPod. I will say more about that later in the book.

When people disagree and challenge your work, that is actually a good sign. If no one cares about your work one way or the other, then why bother doing it.

The skeptics thought my demonstrations were fast because I showed them running on an SGI Octane.* I only needed the Octane because it had three-dimensional hardware texture rendering capabilities. Oh, and also because I got it for free and it was the best workstation at the time. But man, was it ever heavy to carry around.[†] It wouldn't fit in your pocket. The only reason I wrote apps for a mobile device was because it was a cool thing to do. In the end, it showed that software sometimes overpowers brute-force hardware.

I did this fun stuff when I was living in Seattle, Washington, in the late 1990s. Just after the *grunge scene* left town and moved to Los Angeles.[‡] At the same time, our animation software called MAYA was about to take off. We used to have an office in sunny Santa Barbara, California, where everyone is always happy. When I went there on a trip from rainy Seattle and showed them the real-time demos, they immediately wanted to put it into our MAYA software. This was in 2000. After a roller-coaster ride, the fluid solver finally made it into our MAYA 4.5 software in 2002 under the name of “Fluid Effects.” This release would have been impossible without the help of many people from the MAYA development team in Toronto, Canada. They did most of the work, getting my research code into a real piece of software and adding their own secret sauces and spices.

Putting research code into a product takes time and a lot of effort. Trust me. But in the end, it made it into MAYA! We shipped it. And people are using it. We didn’t waste our time on some vaporware to prove we were busy and hardworking. And more importantly, we had fun doing it. Well, at least I had fun doing it.

In 2008, I wrote a fluid app for the iPhone/iPod. I did this mainly because the iPhone supports OpenGL ES, has an accelerometer, and has a multitouch interface. All the things that I wanted to have in an

* At the time, I worked for SGI who actually made the Octane. At the time, it cost over \$60,000. Now, you can buy one on eBay for about \$200. Oh right: SGI stands for Silicon Graphics Incorporated. It used to be the hottest company in Silicon Valley, California. They are pretty much defunct right now. Google has taken over their old campus in Mountain View.

[†] An Octane weighs 24.3 kg (54 lb).

[‡] There were still a lot of good bands like the “Murder City Devils” and “Rorschach Test,” just to name some.

interactive fluid app in 2000. I will say more about this experience later in the book. In 2010, we released Fluid FX, an improved version of the original app.

To summarize: I planted a seed in the area of fluid animation. Others have too. But this is my story.

Jos Stam
Toronto, Ontario, Canada

Acknowledgments

WRITING A BOOK AND DOING RESEARCH IS NOT A LONELY ENDEAVOR. There are many people I want to thank. First, thanks to my wife of 24 years, Pamela, for all her support and love. She is my best friend. To my daughter Gillian because ... because she is my daughter and she is sweet, smart, and beautiful. To my older brother Sim, who introduced me to science and computers. Without him, I would possibly have become a hobo artist in Amsterdam. My parents, Jos and At, and my two sisters, Nel and Go, of course gave me a wonderful kick-start in life. We have short Dutch names in our family. I do not even have a middle name.

I thank CRC Press, especially Rick Adams, for publishing this book. And of course my company Autodesk, which gave me the time to write the book. I especially thank Gordon Kurtenbach, Azam Khan, Francesco Iorio, and the rest of our research group. It is great to be part of such a cool team. Special thanks also to Cory Mogk for creating the cover using MAYA Fluid Effects. Thanks to the Fluid FX team here in Toronto for creating these cool apps, based on my solver after I created my own. I want to single out Dan Pressman and Sergey Buyanov. I also thank the MAYA team we have in Toronto and teams we used to have in Santa Barbara and Seattle. This was a huge team, and if I have to single out one individual, it would be Duncan Brinsmead.

I thank my high school and university friend Marcus Grote for all the stimulating intellectual discussions we used to have. But more related to the current work he pointed me to Chorin and Marsden's book. This book led me out of the confusion and darkness. Since his dad worked at CERN, we got a VIP tour of the future of computing in the early 1980s. Thanks to my final high school year math teacher Nicolas Giovaninni, who taught me the difference between fun math and boring math. All of a sudden, I was getting top grades in math! Thanks to Professor Eugene Fiume for helping me to get accepted at a top North American graduate

school: the University of Toronto. Also thanks for always being available and supportive. Professor Fiume was also a good supervisor, letting me work on whatever I wanted and not minding that I showed up at noon and left at five in the afternoon.

Thanks to the cities of Toronto, Paris, Helsinki, and Seattle. That is where I did most of the research described in this book despite all the distractions these cities have to offer. I find it is a good thing to live in different countries. You get rid of any *national pride* and *provincialism* you might have.* That stuff gets in the way of free and creative thinking. Moving around also makes you humble as you have to learn a new language and adapt to a new culture. It also makes you intellectually richer in the end. Just be willing to learn.

I also want to thank the SIGGRAPH Computer Graphics Technical Achievement Award Committee for presenting me with an award in 2005, partly for the work described in this book. Thanks also to the Academy of Motion Picture Arts and Sciences for presenting a Technical Achievement Award for our MAYA implementation of fluids. I am of course leaving out many other people that have influenced me. You know who you are if you are reading this. Thanks to all of you.[†]

* Unless of course Oranje is playing.

[†] If you feel left out, check out the slide, yes one slide, of my award acceptance speech at the SIGGRAPH Conference in 2005. The link is http://www.autodeskresearch.com/pdf/talks/jos_award05.pdf. If you are still left out, I will put on my Canadian tuck hat and say, “Sorry, buddy, but thanks anyway.”

Author

Jos Stam was born in the Netherlands and educated in Geneva, where he received dual bachelor degrees in computer science and pure mathematics. In 1989, he moved to Toronto, where he completed his master's and PhD in computer science. After that, he pursued postdoctoral studies as an ERCIM fellow at INRIA in France and at VTT in Finland. In 1997, Dr. Stam joined the Alias Seattle office as a researcher and stayed there until 2003, when he relocated to Alias's main office in Toronto. He now works as a senior research scientist as part of Autodesk's acquisition of Alias in 2006.

Dr. Stam's research spans several areas of computer graphics: natural phenomena, physics-based simulation, rendering, and surface modeling, especially subdivision surfaces. He has published papers in all of these areas in journals and at conferences, most notably at the annual SIGGRAPH conference. In 2005, Dr. Stam was presented with one of the most prestigious awards in computer graphics: the SIGGRAPH Computer Graphics Achievement Award. In addition, he won two Technical Achievement Awards from the Academy of Motion Picture Arts and Sciences: in 2005 for his work on subdivision surfaces and in 2007 for his work on fluid dynamics. Dr. Stam was also featured in a January 2008 *Wired* magazine article.

Introduction

Simple play is also the most beautiful. How often do you see a pass of forty meters when twenty meters is enough? Or a one-two in the penalty area when there are seven people around you and a simple wide pass around the seven would be a solution? The solution that seems the simplest is in fact the most difficult one.

JOHAN CRUIJFF (DUTCH SOCCER LEGEND)

Problems can be complicated. Solutions cannot.

ADVERTISEMENT SPOTTED AT THE COPENHAGEN
AIRPORT AFTER GIVING A TALK ON FLUID DYNAMICS
AT DANSIS

There are times when we're testing an actual explosion, and then there are times when we blow stuff up just because we can.

JAMIE HYNEMAN (CO-HOST OF THE POPULAR SHOW
MYTHBUSTERS)

Visually, fluids are everywhere and nowhere.

That might seem like a weird statement. But let me explain.

Fluids like water, honey, maple syrup, fire, explosions, and oil are everywhere. We can see, smell, and sometimes taste and touch them. On the other hand, air is nowhere to be seen. Air is also a fluid. We usually see the effect of motion of the air on other things like insects, dust, clouds,

and the motion of trees. Air is this invisible thing that makes things move. Some fluids are like that. The astonishing fact is that all these fluidlike effects such as liquids and air can be described by a single formal framework called *fluid dynamics*.

We will get to that.

What is a fluid? It really is a technical word that defines a substance that can change its shape in a continuous manner. So, a rock is not really a fluid but when it is heated it becomes one, like lava spurting from a volcano.* It also releases heat that will affect the motion of the fluid that is air. Lava definitely can take many different shapes. In fact that is why fluids are cool. A lava lamp is a good example. The legendary Toronto rock band “Sucker Punch” even has a song called “Cool Like a Lava Lamp.” No really, this song rocks.

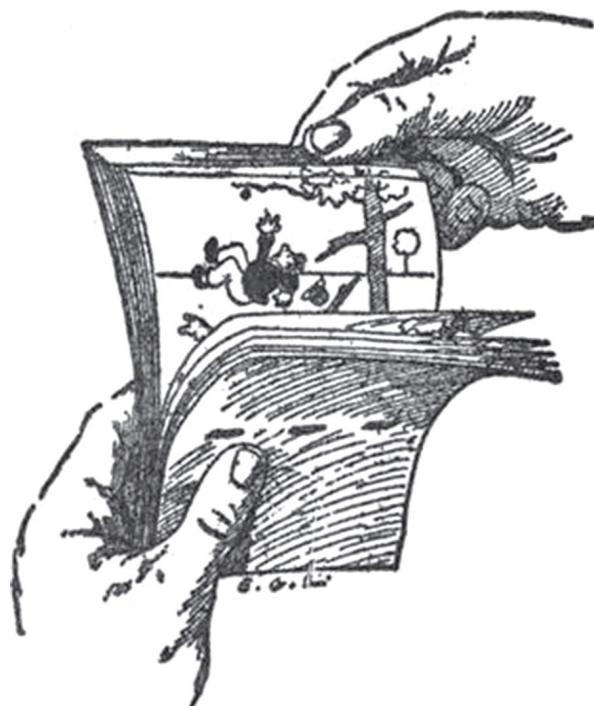
Since fluids are everywhere and nowhere, why do we want to animate them? There are many reasons actually. First, some fluid effects are hard to recreate, like nuclear explosions, tsunamis, and volcanic eruptions. In these cases, computer-generated fluids can create a multitude of catastrophic effects without anyone being hurt. Except, of course, for the animators who must put in long hours creating these effects, and often end up separating from their significant others. These animations can also create effects that have never been seen before or are impossible to create physically. In the special effects industry, they call actual physical shots with no computer graphics *practical*. I always found this to be an awkward word. Of course, some of these practical shots might be fixed “in post” with computer tools to get the perfect look. Practical means you actually blow stuff up, like they do in the popular television show *MythBusters*. Adam Savage and Jamie Hyneman, the two main hosts of this show, actually used to work for Industrial Light and Magic (ILM), which is a major special effects company based in San Francisco. Think *Star Wars*.

In addition to safety issues, there is another reason to use computers to create fluid animations: artistic control. In movies, people want fluids to behave in specific ways, which can be impossible or near impossible using *practical* methods. How would you get liquid horses coming out of a stream of water like in the *Lords of the Rings* movie using real footage? It might be possible but very improbable to achieve. When you simulate a fluid using computer animation you are playing God. You can vary control parameters

* On the other hand, rock formations can be considered in a sense as being very, very, very slow fluids. Think about the drift of continents that cause volcanoes and earthquakes.

to make the fluid do something you have in mind. Of course, this practice of animating fluids has been around since the early days of animation. Think of the classic Disney movies that featured fluids. However, this was a rather tedious method to create a convincing animation of a fluid. Back then every shot had to be hand drawn to create a series of what is called *key-frames*. To get an idea of the motion, artists would use a *flip book* basically flipping rapidly through the pages of a book sequentially. However, in this case, each page has a hand-drawn picture on it created by the artist. By a process of trial and error and intuition, a final fluidlike animation is created. Like that. Some artists were amazing at doing this.

One of the earliest flip books created is shown in Figure 1.1. It was invented in 1868 by John Barnes Linnett, an Englishman from Birmingham. He called it the *kineograph*, a fancy Latin word for *moving picture*. He filed a patent on it too. Why not? In those days there was a tremendous interest in animated pictures. It was not until December 28, 1895,* that the *Frères Lumière* from France showed their first motion



The kineograph

FIGURE 1.1 The first “flip book.”

* Exactly 70 years to the date before I was born.



FIGURE 1.2 The Lumière brothers.

picture to the public. They are both portrayed in Figure 1.2 and one of the frames of their first movie is shown in Figure 1.3. Legend has it that the audience watching the movie for the first time was frightened. The sight of the train entering the station was so scary that most of the audience ran out of the cinema screaming. This short clip also had a fluid effect in it: the steam emanating from the locomotive. Interestingly, the brothers' last name *Lumière* means *light* in French, yes light, as in bright, not as in not being heavy.

The flip book and the cinema are good metaphors for computer animation. Indeed, the goal of computer animation is to create a sequence of pictures that give the illusion of movement. The bottom line is to fill an array of picture elements called *pixels* on a screen for every snapshot. Traditionally in computer graphics, things are animated in a virtual 3D world and then projected onto a screen. This is called the *rendering pipeline* and is depicted in Figures 1.4 and 1.5. The first figure is an engraving by the German Leonardo Albrecht Dürer. This is an early depiction of how perspective projection works and how it is used in art. Figure 1.5 shows the basic methodology of computer graphics. The ball,

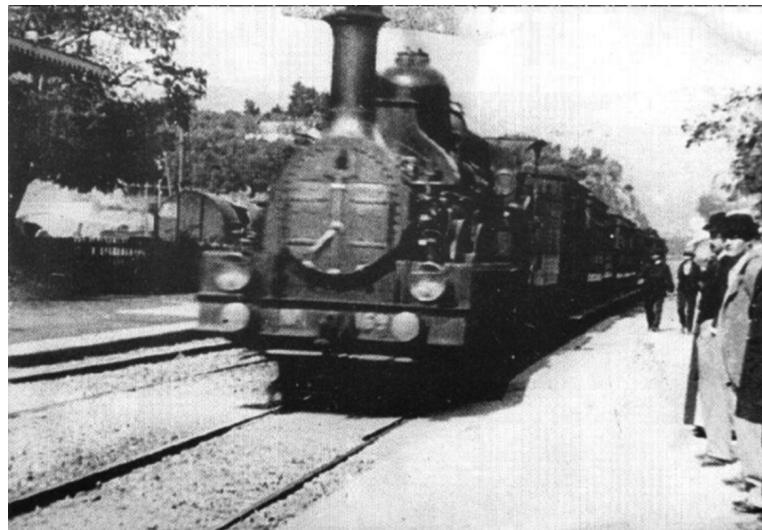


FIGURE 1.3 The first movie shown publicly by the Lumière brothers.

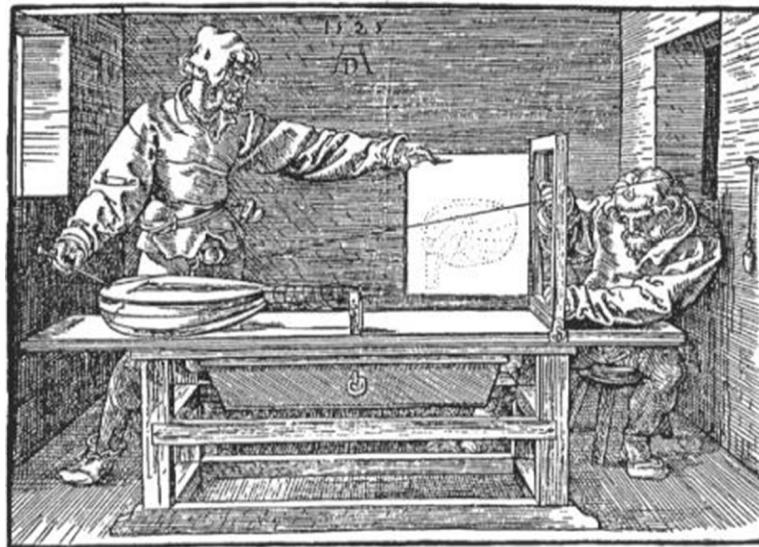


FIGURE 1.4 Engraving by Albrecht Dürer.

pyramid, and cylinder shown in Figure 1.5 are all illuminated by a virtual light source and then finally projected onto a plane: a 2D picture. This picture is then digitized into pixels and sent either to a screen or captured by a camera as shown in Figure 1.5. This process is called the rendering pipeline: 3D stuff goes in on one end of the pipe and *ta da* out comes an image. This book is not about the rendering pipeline and how it works. There are many references that explain this process in

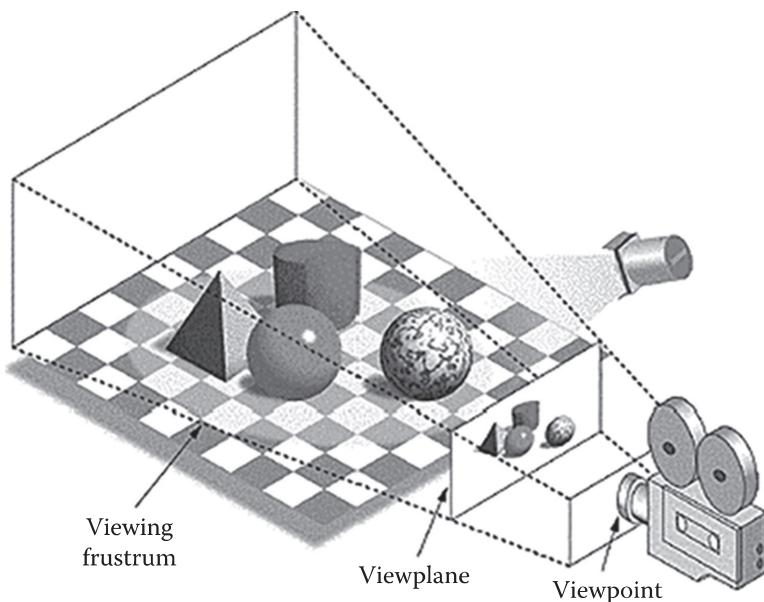


FIGURE 1.5 The rendering pipeline.

great detail.* However, it is important to grasp this basic concept in order to understand the content of this book.

To summarize so far: Our goal is to compute snapshots of a fluid, which are then projected onto a screen to create an animation.

We work in a virtual world that eventually gets projected onto a screen. The screen could also be a standard monitor or a tiny iPhone screen. You name it. But the starting point is a description of a fluid in a space rather than a screen. This brings to mind the famous allegory of “The Cave” by the famous Greek philosopher Plato (424/423–348/347 BCE). In Plato’s allegory, slaves are chained and forced to see only the shadows of 3D objects, not the actual objects themselves.[†] They see only a projection of a reality unbeknownst to them as shown in Figure 1.6, they do not even know about their own *reality*. Plato’s point is of course to show that our senses only perceive a projection of some sort of higher ideal reality. In computer graphics, things are modeled in 3D space and then rendered onto a *2D screen*.

Animation adds another dimension: *time*. Things move in three dimensions, so really our space is four dimensional. A fluid that doesn’t move

* Just type in “rendering pipeline” your favorite search engine.

[†] Plato, *The Republic*. I do not buy it: life is life. The Dutch have a great saying: “dood is zeker en zeker is dood.” Roughly in English: “death is certain and certainty is dead.”

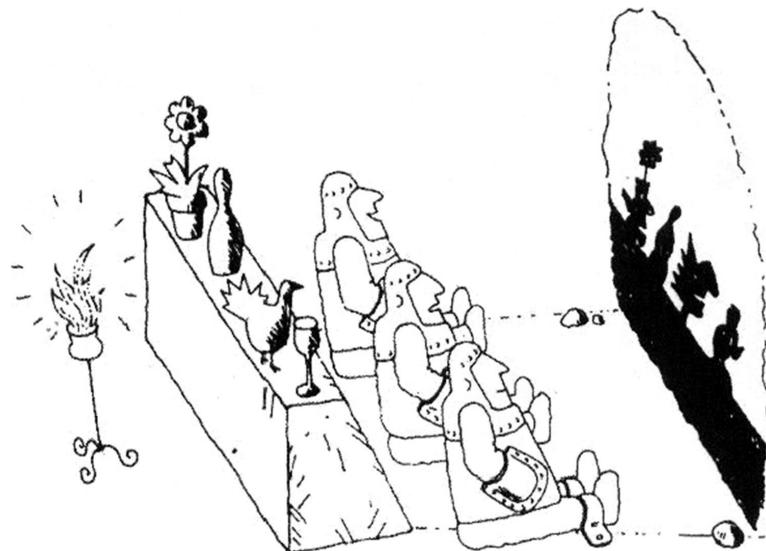


FIGURE 1.6 Plato's allegory of the cave.

around is boring. The objects in Figure 1.5 are usually animated, and different snapshots are rendered on a 2D digital canvas. This process results in an animation: a sequence of frames.

The goal in fluid animation is to compute different states of the fluid over time and render these states as flat two-dimensional representations. The artists manipulate and control the fluid's properties in this more complex 4D space. The physics automatically takes care of the motion. The bottom line is: how to get from an artist's conception of a fluid effect to the generation of a set of pixels on a screen. This is what this book is about.

A final thought before we go on. Whether we work in two-, three-, or four (or higher?)-dimensional spaces, the bottom line is that everything gets translated into *bits*. The bits are really all our speedy friend or foe called a computer can ultimately understand.

A bit is a thing that can only be in two states. These are usually called *zero* and *one* or 0 and 1. But it could also be called *false* and *true* or *Pokemon* and *Papa Smurf*. The famous British mathematician George Boole (1815–1864) came up with a formal way to deal with these bits. He used his theories to write a philosophical book entitled *The Laws of Thought*.*

* Actually, the real title is *An Investigation of the Laws of Thought on Which Are Founded on the Mathematical Theories of Logic and Probabilities*. He was one of the first people to turn logic into pure arithmetic.

The programming language called *Pascal*^{*} even has a *type* named in Boole's honor called Boolean. In another programming language called C++ there is a similar type called *bool*. C++ coders like to type in less characters than Pascal coders. Don't get me started on the *COBOL* programming language.[†]

As an example let's look at the following program written in Pascal computer language.[‡] I know this is an absurd example but that is the whole point.

```
Program PokemonIsNotPapaSmurf;
Var dude      : Boolean;
     dudeName: String;
Begin
  Writeln('Type Pokemon not Papa Smurf');
  Readln(dudeName);
  If dudeName = 'Pokemon' then dude:= True else dude:=
    False;
  If dude then Writeln('Good Job') else
    Writeln('Wrong Dude, Dude');
End.
```

Welcome to the fun world of programming. What this program does is tells the computer to perform a set of precise instructions. The programmer who types in the instructions is the **Master** and the computer is the **Slave**. The **Human** or his **Pet** is the one that uses the program.[§] Why do you think so many good programmers live in their mom's basements or in remote Scandinavian locations? That is actually not true. Programmers like to have fun too. Geez, we even have "brogrammers": Bros who like to party and write code when they are bombed. Do not download their code. More seriously, a lot of programmers are creative types like artists, musicians, builders, mechanics, and so on. Remember you have to **create** code.

^{*} We will say a bit more about this programming language below.

[†] COBOL stands for **C**OMMON **B**USINESS-**O**RIENTED **L**anguage. This is a very verbose language and the brainchild of DARPA.

[‡] The characters in boldface are keywords of the language and the others are just made up by the programmer. No language would have *dude* as a keyword. But who knows? Maybe some brogrammers will come up with such a language.

[§] Sometimes the **Master** has to use his own program. It is called "eating your own dog food" or "dogfooding" in the software industry. Actually, not just for software. I guess people who make dog food were forced to eat their own concoctions themselves, hence the name.

At any rate, the absurd program pretty much works as follows in layman terms:

Master to Slave: My brilliant piece of code is going to be called `PokemonIsNotPapaSmurf`. Got it?

Slave to Master: Yeah whatever.

Master to Slave: I am going to use two locations in your memory. One I am going to call `dude`, and it can only take two values `False` and `True` because it is of type `Boolean`. Pay attention. The next location is called `dudeName` because I say so. This location contains characters and can have any length and so it can take up a lot of your memory. Your job is to fill it up. Got it?

Slave to Master: Yeah I got it. Can we move on a little?

Master to Slave: When **Human** or his **Pet** launches the program, then first print out “Type Pokemon not Papa Smurf” on whatever display device that **Human** or **Pet** is using. You will figure it out.

Slave to Master: Okay, I just did that. They are using a MacBook Pro BTW. You are welcome.

Master to Slave: Read whatever the **Human** or **Pet** types in through whatever input device, I don’t care, and return it to me as soon as possible in a location in your memory called `dudeName`.

Slave to Master: Not even a please huh? **Human** or **Pet** have arrived and are typing in characters. I will store them at a location starting at `0x0000ab345f1d5ecf*` in my memory and I will label it `dudeName` because you said so. Okay, here you go. Have fun.

Master: Alright, I got the information I wanted. Thanks, you piece of dumb hardware. At least you didn’t hang or crash on me this time.

Master: Now let’s check it out. Option 1: Our **Human** or **Pet** typed in `Pokemon`. Good job. I will ask that piece of hardware to write out `Good Job`. Option “anything else really” and I will ask that same piece of hardware to write out `Wrong Dude, Dude`. Yes, even in the case if the **Human** or his **Pet** typed in `Eyjafjallajökull` as long as it is not `Pokemon`.

* Assuming that **Slave** is using a 64-bit architecture.

Human or Pet: Whoa? What is the point of this dumb piece of software that I downloaded for free?

Slave to Human or Pet: Don't blame me. Intermediate slave named **Compiler** translated all of **Master's** nonsense into the following string of bits: 000010101010010101010000000111101010101010
1001100
000000010111111100000000111111111101010101010101010010
1010101001111111111110000000011100101010101010101010100
0000111111110000000101110000111001010100011111110010110
1010110000011010110100101010101010100101.

Slave to Himself: That is all I understand. I know I am mildly autistic just like my **Master**.

To summarize: Fluid animations are created with the help of computers. They generate a sequence of frames that result in an animation: a sequence of pictures. A computer is a speedy, dumb piece of hardware. But it makes for a good pet.

Observations, Equations, and Numbers

Someone told me that each equation I included in the book [*A Brief History of Time*] would halve its sales.

STEPHEN HAWKINS (FAMOUS ENGLISH PHYSICIST)

2.1 BEAVERS, CAVE PERSONS, AND FIRE

80,000 years ago, man's survival in a vast uncharted land depended on the possession of fire. For those early humans, fire was an object of great mystery, since no one had mastered its creation. Fire had to be stolen from Nature; it had to be kept alive—sheltered from wind and rain, guarded from rival tribes. Fire was a symbol of power and a means of survival. The tribe who possessed fire, possessed life.

J.H. ROSNY FROM THEIR BOOK *QUEST FOR FIRE*

Since fluids are everywhere, humans have always been fascinated by them and also feared them. Not only humans by the way. Think of beavers and their dam-building capabilities. No wonder they have a beaver on the logo of the University of Toronto where I got my PhD degree in computer science (see Figure 2.1).



FIGURE 2.1 The official logo of the University of Toronto.

Humans, just like beavers, learned how to deal with fluids early on. Probably the first biggest achievement on the human side was the taming of fire. This is well documented in the epic movie *Quest for Fire* by Jean-Jacques Annaud based on the book by the author “J.H. Rosny.”*

Humans back when they lived in caves found many ingenuous ways to create and preserve fire. Fire is a fluid too. There are of course many other caveman (I mean caveperson) stories involving other fluids like water, sweat, mud, blood, and so on.

Dealing with the effects of fluids in practice is one thing but how do you describe them precisely and control them?

To summarize: Humans and animals had to deal with fluids very early on. They did not rely on equations however.

* I always wondered about the “J.H.” part of the name: could it be Jean-Henri or Jaques-Hubert perhaps? Who knows? It turns out it is a pseudonym made up by two brothers named Joseph Henri Honoré Boex (1856–1940) and Séraphin Justin François Boex (1859–1948). Basically, the “Boex Bros.” They both lived in the city of Brussels in Belgium. They also wrote books together. They wrote a lot of early science fiction books. At some time, there was a rift between the two brothers and the elder one wrote the Fire book under his new pseudonym *J.H. Rosny, ainé*. The last word means “the elder one” in French.

2.2 FROM CAVES TO GREEKS: ARCHIMEDES, GOLD, AND MEDALS

Eureka!

ARCHIMEDES (LEGENDARY GREEK MATHEMATICIAN)

Let us fast forward from cave land to ancient Greece. Archimedes (circa 287–circa 212 BCE) is probably the first person to come up with a principled approach to fluids.

According to legend, this is his story.

A golden crown had been made for King Hiero II. The king was suspicious, however, that some silver had been substituted for the gold by a dishonest blacksmith, but how could he prove it? The king turned to his favorite mathematician, Archimedes.

Let him figure it out.

While bathing one day, Archimedes noticed that the level of water went up as he entered the tub. *Eureka!* Legend again has it that he ran down the streets of Athens naked after his discovery. Back then it was alright to walk or run around naked. Why not? It is hot in Greece. If Archimedes made his discovery in Northern Canada during the winter, he would be running around in the snow in a bear-made fur coat yelling “Awesome dude!”

“Geez Louise I guess crazy Archie next door came up with another discovery, eh?”

To prove that the King had been ripped off and since the crown couldn’t be smelted, Archimedes came up with this clever solution. He asked for a piece of pure gold (guaranteed by a trustworthy source) that had the same weight as the King’s crown. He then submerged the crown first and then the sample of gold in his bathtub. Since water is highly incompressible, the amount of displaced water determines the volume of the submerged object exactly. Therefore, once you know the volume you know the density. The density of pure gold is known.* So, given that the masses of the crown and the sample are fixed and assuming that they are both pure gold, they should displace the same amount of water. This turned out not to be the case. The crown’s gold had been laced with silver, a material cheaper and less dense than gold. The crown displaced more water than the sample made of pure gold. Hence, the crown had more volume than the sample and therefore the crown was less dense than pure gold.

* To be precise: 19.30 g/cm^3 .

This is logical and mathematical reasoning at its best. The difference was slight however. The blacksmith was a smart man. No one knows what happened to the clever blacksmith after the Eureka moment of Archimedes.

Just to hammer in this argument, imagine a huge ball of Styrofoam having the same weight as a small lead marble. It is pretty obvious that they are not made of the same material. No need to immerse them in your bathtub.

The reason I mention this story is that Archimedes discovered an essential fact about fluids. A fluid is highly incompressible. This is not only true for water but it is also true for air. Yes, even air can be considered incompressible in most practical cases; more about that follows later. If you do not believe this fact, clap your hands and ask yourself where the air went between your hands.

Archimedes was a genius and made many early contributions to mathematics. This is why his is the figure on the *Fields Medal* shown in Figure 2.2. In this depiction, he doesn't look like someone who would readily run naked through the streets of Athens. Why do they make him look so serious? Mathematicians I know are fun people who do fun stuff.



FIGURE 2.2 The Fields Medal has a depiction of Archimedes on it.

The medal is named after the Canadian mathematician, John Charles Fields (1863–1932). The medal is awarded every four years to a few mathematicians who did some outstanding work. One caveat: they have to be under the age of 40. No Fields Medals for old geezers.

The awards ceremony is held at the “International Congress of Mathematics.” A gathering of the brightest minds in mathematics: a mathematical Mecca. Fields organized the congress in Toronto in 1924 and lobbied for an award at the meeting. The Fields Medal was eventually first awarded at the 1936 Congress in Oslo, Norway, four years after Fields had died.

The Fields Medal is the equivalent of a Nobel Prize for mathematics. Legend has it that Nobel didn’t like mathematics because his wife had an affair with a mathematician. Knowing mathematicians personally this seems very unlikely. Most likely, Nobel was more like the infamous Barbie Doll from the early 1990s that said, “math is hard.”

Another one of Archimedes contributions related to fluids, is the idea of *buoyancy*.^{*} It is rightly called *Archimedes’ principle*. It goes as follows:

Any object, wholly or partially immersed in a fluid, is buoyed up by a force equal to the weight of the fluid displaced by the object.

This is pretty obvious right. A heavy marble will sink and a rubber ducky will float. The genius element here is that he found a practical principle or law that can predict the exact upward force on the immersed object. It is exactly related to the weight of the fluid that was displaced. It predicts exactly why an object floats or sinks. It explains why a human who cannot swim will die when thrown in a lake. Gravity in this case wins over buoyancy. A hypothetical human made of something less dense than water will survive on the other hand. Actually, it is a bit more complicated. If you are thrown into highly salted water you might float without any effort. People float in the Dead Sea in Israel. This is because the weight of displaced salty water is higher than pure water. Also, people with a lot of body fat have a net density close to water and they might easily float in the ocean. Muscle is denser than fat. So, if you cannot swim it is better to be fat than muscular.

^{*} It is derived from the word *buoy* which is derived from the Dutch word *boei*. You know, the things you see floating in harbors.

To summarize: Some ancient Greeks like Archimedes liked to run down the street naked. Make sure your kids learn how to swim from an early age. Bathtubs are a good place to relax and do science.

2.3 LONG, CURLY HAIR ED MATHEMATICIANS, THE ABYSS AND THE AIRBRUSH

Before you know it, the Renaissance will be here and we'll all be painting.

WOODY ALLEN'S CHARACTER PLAYING THE FOOL
IN HIS MOVIE *EVERYTHING YOU ALWAYS WANTED TO
KNOW ABOUT SEX BUT WERE AFRAID TO ASK* (1972)

The Dutch customs once thought my pictures were photos. Where on earth did they think I could have photographed my subjects? In Hell, perhaps?

HANS RUDOLF GIGER (LEGENDARY SWISS
AIRBRUSH ARTIST)

After exponential quantities the circular functions, sine and cosine, should be considered because they arise when imaginary quantities are involved in the exponential.

LEONHARD EULER (FAMOUS SWISS MATHEMATICIAN)

We move from ancient Greece to Renaissance France. Blaise Pascal^{*} (1623–1662) is another genius who studied fluids. Not only that but he was one of the first people to build a computing machine. A computer language created by Niklaus Wirth[†] called *Pascal* is named in his honor. Since I grew up in Geneva, Switzerland, this was one of the first computer languages I had to code in. It is a somewhat overly formal language. I still remember how one of my teachers spent a lot of time explaining what a *pointer* is. In another language called C it is completely obvious, just the location of something in the physical memory of a computer. That's it. A company called Borland founded by the Frenchman Philippe Kahn,

^{*} I always joke that he should be from Texas. Pascal is a popular first name in French. In Texas there are many people with last names like “Joe,” “Paul,” “Randy,” etc.

[†] He is a professor at the *Eidgenössische Technische Hochschule Zürich*. Gesundheit: the short version we use is ETHZ.

created a wickedly cool compiler released in the 1980s called *Turbo Pascal*. It was cheap, fast, and user friendly. I loved it. It was fast because everything was kept in slave's memory when it was translated to 0 and 1, not on those dinky floppy drives. Of course, this was all cool until your computer crashed. In that case, oh well. Rewrite your code. At any rate, Version 2.0 usually ends up being cleaner.*

Now, back to the real Pascal and his work on fluids.

One of Pascal's more memorable experiments is his *crève-tonneau*, which literally means "barrel buster" in French. It sounds like something out of a Renaissance *MythBuster* episode. Figure 2.3 shows a portrait of Blaise Pascal (he looks like a dude you want to chill out with) and Figure 2.4



FIGURE 2.3 Blaise Pascal.

* Geek story. When I coded at VTT in Espoo, Finland, I lost one month worth of writing code because I was using the "vi" editor. I hit the ":wq" keys and my computer was low on disk storage. It couldn't do the "w" operation but then executed the "q" anyway. I lost everything. Lessons learned: (1) break up your code into smaller files, (2) type in ":w" first and then type in ":q," and stay away from ":q!", and (3) do not use "vi" unless you are a hipster buying vinyl. My version 2 of the code was actually better. I was careless and should have known better.

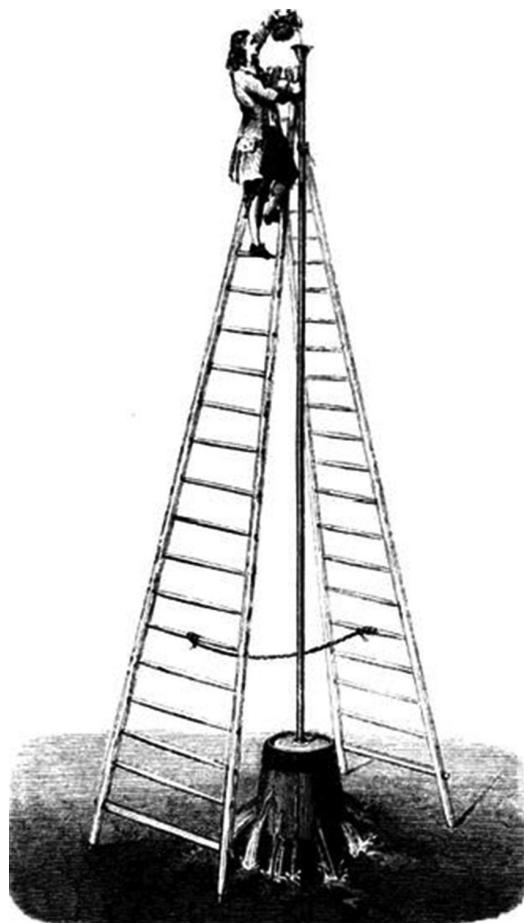


FIG. 45.—Hydrostatic paradox. Pascal's experiment.

FIGURE 2.4 The crève-tonneau experiment.

shows his famous *tonneau* experiment. So, why would anyone climb up a 10-m ladder to make a barrel burst? Of course because it makes science look cool. Science is fun *Messieurs, Dames et Mademoiselles*. Pascal was making a point. It confirmed his principle.

Pressure applied to an enclosed fluid is transmitted undiminished to every part of the fluid, as well as to the walls of the container.

This is directly related to the fact that a fluid is highly incompressible. This is how it works. Pascal connected a vertical 10-m pipe to a barrel already filled with water. Then he climbed the ladder and poured water into the pipe. This raises the overall pressure in the barrel. And since the barrel was poorly made, it burst: *crevé*. The burst factor depends both on

the strength of the barrel and on the height of the pipe. More concretely, the necessary height of the pipe to burst the barrel is proportional to the strength of the barrel.

No wonder the unit of pressure is named after Pascal. It is equal to the amount of force exerted per area. This is the reason why it is a bad idea to walk in stiletto heels on Copacabana Beach. Too much force on a tiny area. On the other hand, it is a good idea to wear snowshoes also known as *racket shoes* when walking on soft snow in Northern Canada. The rackets make your feet bigger and redistribute the force caused by your weight to create a lower pressure.

The pressure of the atmosphere at sea level is 101,325 Pa. That is huge: a lot of curly, long-haired French philosophers. Consequently, if the inside of your body were a vacuum, you would shrink to a freaky mess. The pressure at the bottom of oceans is even higher. Humans wouldn't survive there but many fish and other creatures whose body density is almost the same as water are doing just fine.

The deepest known location in the world's oceans is the so-called *Mariana Trench* near the island of Guam in the western part of the Pacific Ocean. It is about 11 km below sea level and the pressure there is a thousand times higher than the atmospheric pressure. It is deeper than how tall the Mount Everest is: 11 versus 8 km. Humans first reached the bottom of the trench using a specially built *bathyscaphe* called the *Trieste* in 1960 (see Figure 2.5). This is sort of an Everest climb in reverse. Figure 2.6 shows some of the bizarre creatures you might encounter at those depths.* This layer is known as the *abyss*. Marine biologists are slowly discovering the life in this fascinating world below us. National Geographic even claims that 86% of all life on earth is still unknown to humankind.[†] Looking at the creatures in Figure 2.6, one can only imagine what else is down there.

Thus far we have dealt only with static fluids: fluids that do not move. Until Pascal's barrel burst that is. Let's look at an example that involves moving fluids. Let's introduce the airbrush.

* For more pictures see Claire Nouvian's book, *The Deep: The Extraordinary Creatures of the Abyss*. Of course it is always easier to Google it. These creatures can resist the high pressure because their tissues are composed of essentially a nondeformable liquid that is also gelatinous; this allows these creatures to balance the external pressure with internal pressure.

[†] <http://news.nationalgeographic.com/news/2011/08/110824-earths-species-8-7-million-biology-planet-animals-science/>

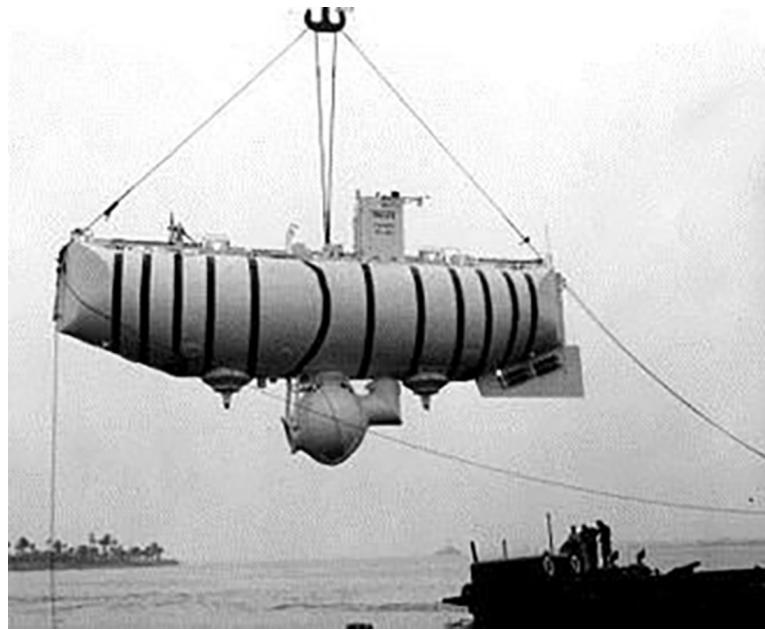


FIGURE 2.5 The bathyscaphe called the *Trieste* that made it to the bottom of the Mariana Trench. (Courtesy of the U.S. Naval Historical Center, Washington, DC.)



FIGURE 2.6 Some of the creatures you might encounter in the abyss. This is for real. These are not computer generated.

I spent most of my teenage years painting with the airbrush until I realized computers could create pictures as well. The airbrush enables you to create wicked pictures but first you need to master the technology in order to create them. The airbrush is basically a spray can with controls over the outflow and the color of the paint. I usually liked acrylics and watercolors over oil paints. Acrylics are cool but they dry really fast so you have to make sure to clean everything really rapidly. Watercolors are user friendly but they are not as vivid as acrylics. Oil paints are the ultimate paint when using a traditional brush. They take forever to dry however. I remember painting outside in the summer and bugs would get stuck on the oil paint. My experience with oils and the airbrush were not

very good: big mess. Turpentine to the rescue! I am always amazed what *graffiti artists* can produce with their clunky spray cans. They are highly skilled badass artists.

Let's see how the airbrush really works since it involves fluid *kinematics*. Kinematics is a fancy word for fluids that actually move and this is related to fluid animation. We will get to the *dynamics* of fluids below, which describes how the kinematic properties of a fluid, like velocities, evolve over time.

Figure 2.7 is a photograph of my old Paasche airbrush and Figure 2.8 shows a schematic depiction of the hardware involved in using an airbrush. It is not as easy as going to the local art store and purchasing paints, brushes, and canvases or paper. First, you need a supply of compressed air. Air can be compressed when trapped in a container such as a cylinder.* Usually, airbrush artists prefer a *compressor*: a device that is also useful to pump bicycle and car tires. It fills up a tank of compressed air now and then when the pressure is below a certain threshold. Some industrial



FIGURE 2.7 My Paasche airbrush.

* My mom drove me to an industrial part of Geneva where they rented these cylinders. You only pay for the air inside the cylinder and you rent the cylinder. The guys at the warehouse were puzzled why a tall, gangly teenager and his mom would want to rent one of them. They had never heard of airbrush art. The cylinder was cool but it ran out of gas even before I finished my second painting. That is when I convinced my parents to buy me a compressor. After that I convinced my dad to buy me a Sinclair QL computer and an Amiga computer.

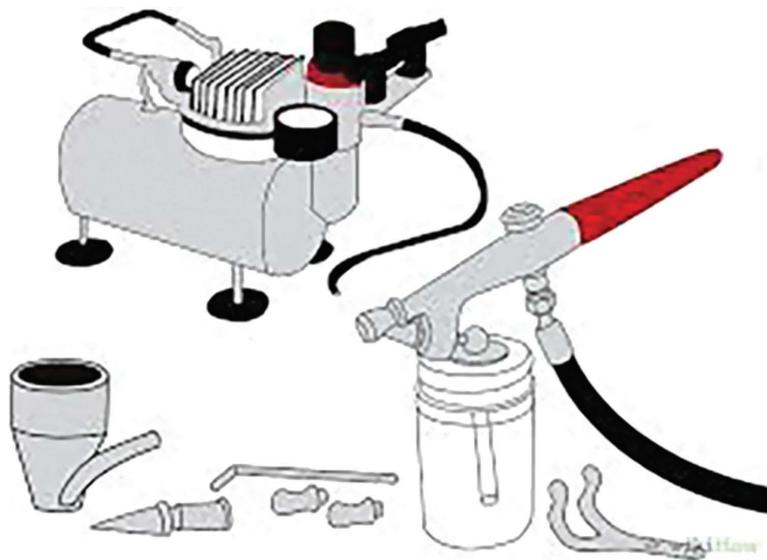


FIGURE 2.8 A cartoon of an airbrush setup.

compressors can be very noisy, but the ones built for artists typically emit only a small hum. They are however more expensive and smaller.* The compressor is connected via a hose to the airbrush, usually at the bottom. That is the input of air. The input of paint comes via a container attached to the airbrush. Where the fluid kinematics comes in is when both air and paint are mixed and ejected out of the nozzle. Sort of like colored water coming out of a garden hose set on *spray* or *mist* but not on *soak*. Anyone who has a backyard will know what I mean.

The basic mechanism of how an airbrush works goes back to the work by the Swiss mathematician Daniel Bernoulli. Figure 2.9 shows him and the cover of his famous book on fluid dynamics. Back then it was okay for mathematicians to wear a curly wig. He comes from a famous family of mathematicians. There are Bernoulli numbers, Bernoulli distributions, Bernoulli principles, and so on.

One member of the Bernoulli family even proved that the following sum does not have a finite value:

$$H = 1 + \frac{1}{2} + \frac{1}{3} + \dots$$

* I once used an industrial compressor and it almost gave my friend's cat a heart attack when the cat was awakened from a deep sleep because the compressor automatically turned itself on. We got away with some nasty scratches.

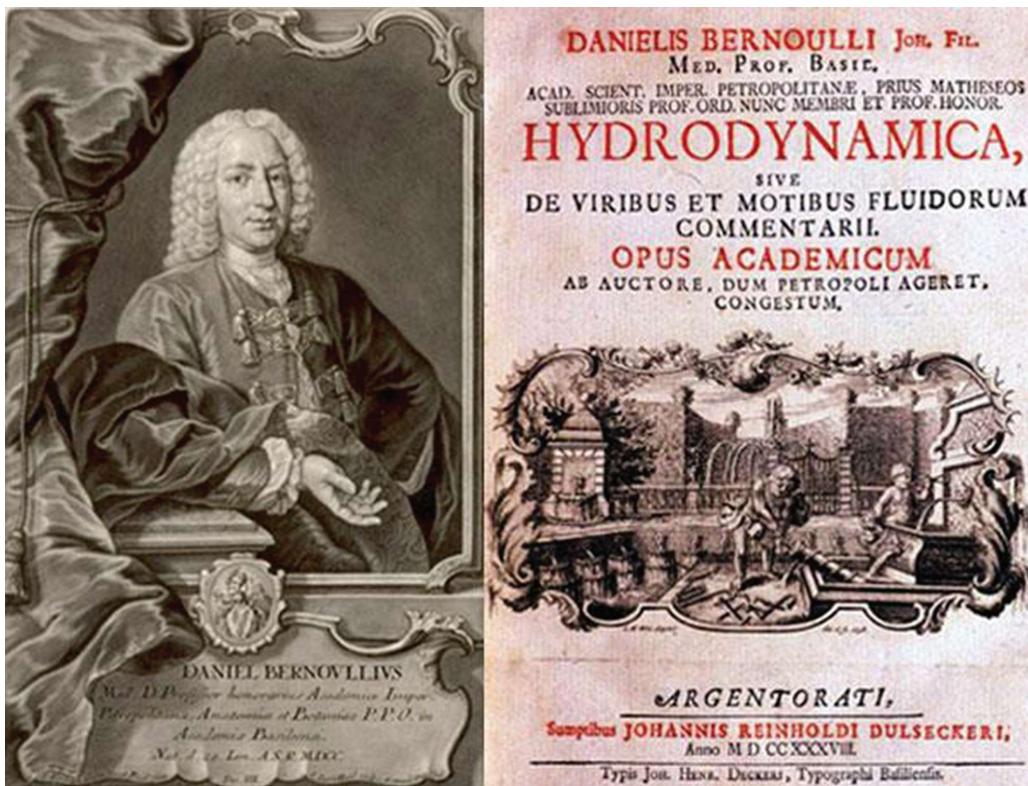


FIGURE 2.9 Daniel Bernoulli (left) and the cover of his treatise of fluids (right).

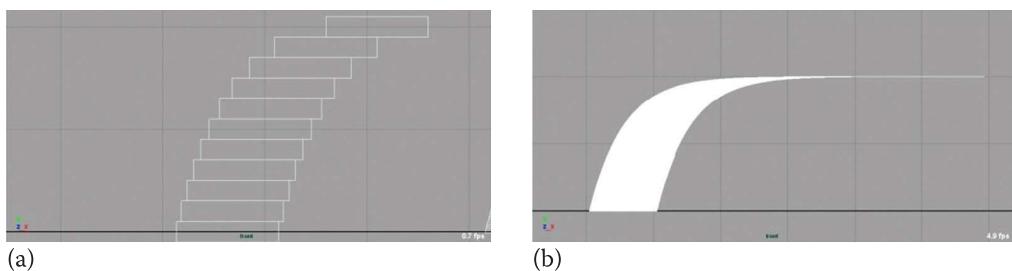


FIGURE 2.10 How to stack 10 books (a) and 10,000 books (b).

It is infinite. The “...” of course means that the sum involves an infinite number of terms. The next one is $1/4$, then comes $1/5$, and so on. It would take an infinite amount of paper to write down the entire sum. This sum is called *harmonic*. It crawls toward infinity very slowly however.

One fun practical way to visualize this phenomenon is through the *book stacking problem*. Here is the problem: given identically sized books how would one stack them up in order to reach the furthest distance possible? Figure 2.10 shows how to stack 10 books and 10,000 books.* In principle,

* I created these figures in our MAYA animation software using an MEL script.

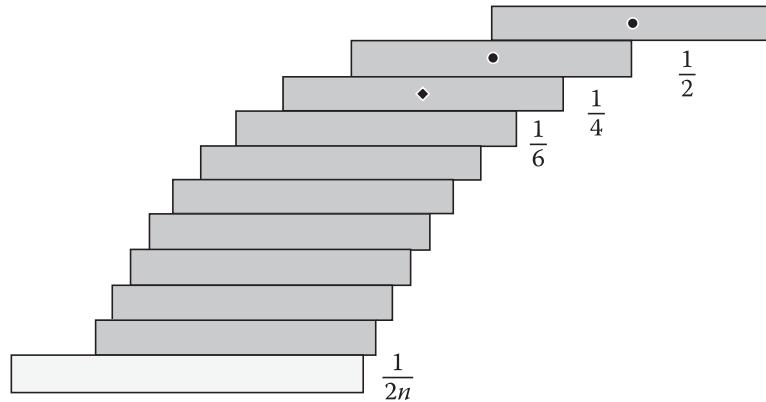


FIGURE 2.11 Optimal *classic* solution of the book-stacking problem.

you could reach any distance given enough books, there is an optimal manner of doing this.

Do try this at home.

The optimal stacking formula is shown in Figure 2.11. It shows that the distance between the adjacent books is in a decreasing series:

$$\frac{1}{2}, \frac{1}{4}, \frac{1}{6}, \dots$$

The total distance from the table is therefore equal to the sum of the terms of the series:

$$\text{Overhang Distance} = \frac{1}{2} + \frac{1}{4} + \frac{1}{6} + \dots$$

This sum also goes to infinity. This means that any distance can be reached given enough books. Again, try this at home.

Actually, a stack of playing cards will do a better job. You would have to go to your local “Dollar Store” and be ready to spend about a thousand dollars or more and have a lot of time on your hands to stack them properly. Also, make sure you close all windows and have your pet locked up in another room. This will get you the stacking shown in Figure 2.10b.

Not all infinite sums have infinite values even if you need an infinite amount of paper to write them down. For example, check out the following series:

$$1 = \frac{1}{2} + \frac{1}{4} + \frac{1}{8} + \frac{1}{16} + \frac{1}{32} + \dots$$

This is the sum of the inverses of all the powers of two and it is equal to one. That is why Zeno's turtle or rabbit can cross a road. "Zeno's Paradox"^{*} roughly states that the turtle will never reach the other side of the road because there is always another halfway to go. Actually, in the original story it is that the rabbit can never catch up to the turtle once he arrogantly took a nap.

Zeno's Paradox is only paradoxical if you do not accept infinities. And if you do, *no problema amigo*.

Welcome to the world of infinities.

The sum of an infinity number of halves and halves of themselves eventually converges to a finite number. So no problem; our turtle is able to cross the road, unless it is run over by a car. Many interesting results like this came from the Bernoulli family.

Infinite series are wicked. Consider this astonishing result:

$$1 + 2 + 3 + 4 + 5 + 6 + \dots = -\frac{1}{12}.$$

This is crazy right? Nevertheless, it is used in Bosonic String Theory[†] and implies that space time has to have 26 dimensions. Look it up.

Let's get back down to earth and Daniel Bernoulli's work on fluids. His principle roughly states that.

When airspeed increases, pressure decreases.

This explains how an airbrush works. The paint gets sucked up from the container through the pipe and gets mixed with the air and eventually leaves the nozzle of the airbrush to create a cool painting.

Daniel Bernoulli actually contributed a lot more to science than this principle; he gave us an equation.

Equations are mathematical objects which help to think clearly about problems. Equations also help to communicate results in a universal language: Mathematics. As an example, let's look at one of the most beautiful equation in mathematics:

^{*} The paradox is named after the Greek philosopher Zeno of Elea (-490 | -430). He is not to be confused with Zeno of Citium of course. A fun paradox is: "This sentence is false." A geekier version is due to the British philosopher Bertrand Russell that states: "Let A be the set of all sets that do not contain themselves. Is A an element of A ?" This statement destroyed Gottlob Frege's theory of Logic that he had worked on for many years.

[†] The term *Bosonic* has no relation to *Bozo the Clown*. It is named after the famous Indian physicist Satyendra Nath Bose (1894–1974).

$$e^{i\pi} + 1 = 0.$$

This equation combines five fundamental constants of mathematics into one equation. Any mathematician will acknowledge its beauty regardless of his or her cultural background or mother tongue.* It is a universal truth and it is elegant. This equation was discovered or created, whatever, by another famous Swiss mathematician named Leonhard Euler. We will discuss his fundamental contributions to fluids as follows.

The equation states that Napier's constant e raised to the power of the complex number i times π plus one is equal to zero. These five numbers show up all over the place in mathematics and Euler combined them all in one equation.

This is pure genius.

What it really says in plain language is: "if you rotate a vector counter-clockwise in the plane by 180° it will point in the opposite direction. And if you take the sum of the original vector and the rotated one it will be zero." That seems pretty obvious. Euler's genius was to concoct an equation from this obvious fact that involves very different mathematical constants. They show up all over the place in different mathematical fields but here they are all united by Euler's formula. I think any high school math course should mention this wonderful equation. It will teach students about awesome numbers, more about these following numbers.

Of course, any book that mentions equations has to state Einstein's famous equation stating that energy is related to mass via the speed of light squared. Here you go:

$$E = mc^2.$$

I just lost one-quarter of the readers at this point according to Hawking's quote at the beginning of this chapter. No worries. We do not need these equations to animate fluids. But there will be more equations to come. Too bad that equations scare people as in most cases they actually make understanding things simpler, not more complicated. Some people abuse equations; however, they are called math snobs. It makes them look clever

* This is kind of like *soccer*, *voetbal*, *futebol*, ... In almost any country I have traveled to, you just have to mention a team's name or a soccer legend to have an instant bond in the foreign land. When I was in South Korea I mentioned that I was born in Holland and most people immediately responded *Guus Hiddink!*

but they are understood only by a few mortals. I went to many math talks and usually everyone is lost after the first two or three slides. Then some expert in the audience wakes up at the front of the audience and asks a really clever question at the end of the talk. In my experience, older mathematicians give clearer presentations. The famous mathematician, Stephen Smale, comes to mind: he proved the Poincaré Conjecture for dimensions larger than four on the beach of Rio and got a Fields Medal for it. His talks, at least the ones which I had the privilege to attend, were crystal clear. Look it up.

There is, however, one difference I want to point out between these two equations. Euler's equation is a *truth* which has a widely accepted proof: a series of logical arguments that starts with known facts and then ends with the equation in the final step.* You could say mathematics is a beautiful and clever solitary game. When I say solitary I do not mean that it is not competitive. Everyone wins when a result is proven. That is progress. Of course there is some competition between mathematicians about who will find the first proof or the most elegant one. But in the end it's all good. They all go to the pub to celebrate a new *truth*.†

Einstein's equation is different and much more ambitious as it tells us something about the way nature works. It is based on physical principles which are then condensed into an equation. The speed of light denoted by c is different from a mathematical constant like π . The speed of light has to be measured while the constant π can be defined exactly as, are you ready, "half the length of the perimeter of an (ideal) circle having a radius equal to one." Or, and this is pure mathematical magic, you could define it as the number π that makes Euler's equation true. From that you can then derive the perimeter of a circle: magic. Mathematics is all about starting with some facts and then deriving other facts from them from different angles of attack. Most often, the simplest argument is preferred. This principle goes by the name of *Occam's razor*. When there are alternate theories to explain a phenomenon, just pick the simplest one.

Now back to Einstein's equation. In 1983, the speed of light became a constant through a clever trick. The unit of length (the meter) was redefined so that its value was fixed once and for all. The meter used to be defined as the length of a bar of metal called *l'étalon* held in some safe

* Mathematicians usually put *QED* at the end of their proofs. *QED* is an acronym for *Quod Erat Demonstrandum*. This is Latin for "which had to be demonstrated."

† David Hilbert famously wrote that: "Mathematics is a game played according to certain simple rules with meaningless marks on paper."

location in Paris, France (not in Texas). After many measurements, the international community decided that length *au contraire* should be defined by the speed of light and not the other way around. So, they came up with the following definition: “a meter is the distance traveled by light in vacuum during a time interval of $1/c$ seconds.” The speed of light is a rather ugly looking constant determined by experiments:

$$c = 299,792.458 \text{ km/s.}$$

This approach seems to do the job so far.

You cannot prove a physical theory, you can only disprove it.

In a nutshell that is the difference to me between mathematics and physics. Of course, this is oversimplifying the dichotomy. Anyone can always disprove a mathematical result. But you do not need an experiment just a counter-mathematical argument.

Speaking of physical equations and elegance, I cannot stop myself from mentioning the *Lagrangian of the Standard Model* used in particle physics. It is depicted in all its full glory in Figure 2.12. It makes for great wallpaper. Seriously, it does work, apparently. One of the terms captures the recently discovered *God particle* also known as the *Higgs Boson*.^{*} Try to spot the Higgs term. Mathematicians call these equations not elegant, inelegant, ugly, a disgrace, garbage, and “that is why I am not doing physics.” Too bad the elegant equations do not always work. Einstein and Newton’s equations are exceptions rather than the rules.

Let’s get back to fluids.

Daniel Bernoulli’s equation precisely relates the speed (velocity) of the flow and the pressure. From his principle, we know they are inversely related. However, we have to keep in mind that this result is valid only for simple fluids. What that means is that we ignore dynamics and things like turbulence. The cool stuff we will deal with later in this book. But it works in many cases where it is a good approximation.

To establish a quantitative relation between the velocity and pressure, we will use a somewhat unconventional derivation called *dimensional analysis*. The following is going to be a bit geeky but please bear with me. I promise it will not involve any calculus, even though I love calculus.

^{*} A *Boson* is named after the same Satyendra Nath Bose mentioned earlier. Peter Ware Higgs (1929) is a British physicist who was awarded the Nobel Prize in 2013.

$$\begin{aligned}
& -\frac{1}{2}\partial_\mu g_\mu^a \partial_\nu g_\nu^a - g_s f^{abc} \partial_\mu g_\mu^a g_\mu^b g_\nu^c - \frac{1}{4}g_s^2 f^{abc} f^{ade} g_\mu^b g_\nu^c g_\mu^d g_\nu^e + \\
& \frac{1}{2}ig_s^2 (\bar{q}_j^a \gamma^\mu q_j^a) g_\mu^a + \bar{G}^a \partial^2 G^a + g_s f^{abc} \partial_\mu \bar{G}^a G^b g_\mu^c - \partial_\nu W_\mu^+ \partial_\nu W_\mu^- - \\
& M^2 W_\mu^+ W_\mu^- - \frac{1}{2} \partial_\nu Z_\mu^0 \partial_\nu Z_\mu^0 - \frac{1}{2c_w^2} M^2 Z_\mu^0 Z_\mu^0 - \frac{1}{2} \partial_\mu A_\nu \partial_\mu A_\nu - \frac{1}{2} \partial_\mu H \partial_\mu H - \\
& \frac{1}{2} m_h^2 H^2 - \partial_\mu \phi^+ \partial_\mu \phi^- - M^2 \phi^+ \phi^- - \frac{1}{2} \partial_\mu \phi^0 \partial_\mu \phi^0 - \frac{1}{2c_w^2} M \phi^0 \phi^0 - \beta_h [\frac{2M^2}{g^2} + \\
& \frac{2M}{g} H + \frac{1}{2}(H^2 + \phi^+ \phi^0 + 2\phi^+ \phi^-)] + \frac{2M^4}{g^2} \alpha_h - ig_{cw} [\partial_\nu Z_\mu^0 (W_\mu^+ W_\nu^- - \\
& W_\nu^+ W_\mu^-) - Z_\nu^0 (W_\mu^+ \partial_\nu W_\mu^- - W_\mu^- \partial_\nu W_\mu^+) + Z_\mu^0 (W_\nu^+ \partial_\nu W_\mu^- - \\
& W_\nu^- \partial_\nu W_\mu^+)] - ig_{sw} [\partial_\nu A_\mu (W_\mu^+ W_\nu^- - W_\nu^+ W_\mu^-) - A_\nu (W_\mu^+ \partial_\nu W_\mu^- - \\
& W_\mu^- \partial_\nu W_\mu^+) + A_\mu (W_\nu^+ \partial_\nu W_\mu^- - W_\mu^- \partial_\nu W_\mu^+)] - \frac{1}{2} g^2 W_\mu^+ W_\mu^- W_\nu^+ W_\nu^- + \\
& \frac{1}{2} g^2 W_\mu^+ W_\nu^- W_\nu^+ W_\mu^- + g^2 c_w^2 (Z_\mu^0 W_\mu^+ Z_\nu^0 W_\nu^- - Z_\mu^0 Z_\mu^0 W_\nu^+ W_\nu^-) + \\
& g^2 s_w^2 (A_\mu W_\mu^+ A_\nu W_\nu^- - A_\mu A_\nu W_\mu^+ W_\nu^-) + g^2 s_w c_w [A_\mu Z_\mu^0 (W_\mu^+ W_\nu^- - \\
& W_\nu^+ W_\mu^-) - 2A_\mu Z_\mu^0 W_\nu^+ W_\nu^-] - g_0 [H^3 + H \phi^0 \phi^0 + 2H \phi^+ \phi^-] - \\
& \frac{1}{8} g^2 \alpha_h [H^4 + (\phi^0)^4 + 4(\phi^+ \phi^-)^2 + 4(\phi^0)^2 \phi^+ \phi^- + 4H^2 \phi^+ \phi^- + 2(\phi^0)^2 H^2] - \\
& g M W_\mu^+ W_\mu^- H - \frac{1}{2} g \frac{M}{c_w^2} Z_\mu^0 Z_\mu^0 H - \frac{1}{2} ig [W_\mu^+ (\phi^0 \partial_\mu \phi^- - \phi^- \partial_\mu \phi^0) - \\
& W_\mu^- (\phi^0 \partial_\mu \phi^+ - \phi^+ \partial_\mu \phi^0)] + \frac{1}{2} g [W_\mu^+ (H \partial_\mu \phi^- - \phi^- \partial_\mu H) - W_\mu^- (H \partial_\mu \phi^+ - \\
& \phi^+ \partial_\mu H)] + \frac{1}{2} g \frac{1}{c_w^2} (Z_\mu^0 (H \partial_\mu \phi^0 - \phi^0 \partial_\mu H) - ig \frac{s_w^2}{c_w^2} M Z_\mu^0 (W_\mu^+ \phi^- - W_\mu^- \phi^+) + \\
& ig s_w M A_\mu (W_\mu^+ \phi^- - W_\mu^- \phi^+) - ig \frac{1-2c_w^2}{2c_w^2} Z_\mu^0 (\phi^+ \partial_\mu \phi^- - \phi^- \partial_\mu \phi^+) + \\
& ig s_w A_\mu (\phi^+ \partial_\mu \phi^- - \phi^- \partial_\mu \phi^+) - \frac{1}{4} g^2 W_\mu^+ W_\mu^- [H^2 + (\phi^0)^2 + 2\phi^+ \phi^-] - \\
& \frac{1}{4} g^2 \frac{1}{c_w^2} Z_\mu^0 Z_\mu^0 [H^2 + (\phi^0)^2 + 2(2s_w^2 - 1)^2 \phi^+ \phi^-] - \frac{1}{2} g^2 \frac{s_w^2}{c_w^2} Z_\mu^0 \phi^0 (W_\mu^+ \phi^- + \\
& W_\mu^- \phi^+) - \frac{1}{2} ig^2 \frac{s_w^2}{c_w^2} Z_\mu^0 H (W_\mu^+ \phi^- - W_\mu^- \phi^+) + \frac{1}{2} g^2 s_w A_\mu \phi^0 (W_\mu^+ \phi^- + \\
& W_\mu^- \phi^+) + \frac{1}{2} ig^2 s_w A_\mu H (W_\mu^+ \phi^- - W_\mu^- \phi^+) - g^2 \frac{s_w^2}{c_w^2} (2c_w^2 - 1) Z_\mu^0 A_\mu \phi^+ \phi^- - \\
& g^1 s_w^2 A_\mu \phi^+ \phi^- - \bar{e}^\lambda (\gamma \partial + m_\lambda^0) e^\lambda - \bar{\nu}^\lambda \gamma \partial \nu^\lambda - \bar{u}_j^\lambda (\gamma \partial + m_u^\lambda) u_j^\lambda - \bar{d}_j^\lambda (\gamma \partial + \\
& m_d^\lambda) d_j^\lambda + ig s_w A_\mu [-(\bar{e}^\lambda \gamma e^\lambda) + \frac{2}{3}(\bar{u}_j^\lambda \gamma u_j^\lambda) - \frac{1}{3}(\bar{d}_j^\lambda \gamma d_j^\lambda)] + \frac{iq}{c_w^2} Z_\mu^0 [(\bar{\nu}^\lambda \gamma^\mu (1 + \\
& \gamma^5) \nu^\lambda) + (\bar{e}^\lambda \gamma^\mu (4s_w^2 - 1 - \gamma^5) e^\lambda) + (\bar{u}_j^\lambda \gamma^\mu (\frac{4}{3}s_w^2 - 1 - \gamma^5) u_j^\lambda) + \\
& (\bar{d}_j^\lambda \gamma^\mu (1 - \frac{8}{3}s_w^2 - \gamma^5) d_j^\lambda)] + \frac{iq}{2\sqrt{2}} W_\mu^+ [(\bar{\nu}^\lambda \gamma^\mu (1 + \gamma^5) e^\lambda) + (\bar{u}_j^\lambda \gamma^\mu (1 + \\
& \gamma^5) C_{\lambda\kappa} d_j^\kappa)] + \frac{iq}{2\sqrt{2}} W_\mu^- [(\bar{e}^\lambda \gamma^\mu (1 + \gamma^5) \nu^\lambda) + (\bar{d}_j^\lambda C_{\lambda\kappa}^\dagger \gamma^\mu (1 + \gamma^5) u_j^\kappa)] + \\
& \frac{iq}{2\sqrt{2} M} [-\phi^+ (\bar{\nu}^\lambda (1 - \gamma^5) e^\lambda) + \phi^- (\bar{e}^\lambda (1 + \gamma^5) \nu^\lambda)] - \frac{iq m_\lambda^0}{2 M} [H (\bar{e}^\lambda e^\lambda) + \\
& i \phi^0 (\bar{e}^\lambda \gamma^5 e^\lambda)] + \frac{iq}{2M\sqrt{2}} \phi^+ [-m_d^\kappa (\bar{u}_j^\lambda C_{\lambda\kappa} (1 - \gamma^5) d_j^\kappa) + m_u^\kappa (\bar{u}_j^\lambda C_{\lambda\kappa} (1 + \\
& \gamma^5) d_j^\kappa)] + \frac{iq}{2M\sqrt{2}} \phi^- [m_d^\lambda (\bar{d}_j^\lambda C_{\lambda\kappa}^\dagger (1 + \gamma^5) u_j^\kappa) - m_u^\kappa (\bar{d}_j^\lambda C_{\lambda\kappa}^\dagger (1 - \gamma^5) u_j^\kappa)] - \\
& \frac{iq m_\lambda^0}{2 M} H (\bar{u}_j^\lambda u_j^\lambda) - \frac{iq m_\lambda^0}{2 M} H (\bar{d}_j^\lambda d_j^\lambda) + \frac{iq m_\lambda^0}{2 M} \phi^0 (\bar{u}_j^\lambda \gamma^5 u_j^\lambda) - \frac{iq m_\lambda^0}{2 M} \phi^0 (\bar{d}_j^\lambda \gamma^5 d_j^\lambda) + \\
& \bar{X}^+ (\partial^2 - M^2) X^+ + \bar{X}^- (\partial^2 - M^2) X^- + \bar{X}^0 (\partial^2 - \frac{M^2}{c_w^2}) X^0 + \bar{Y} \partial^2 Y + \\
& ig c_w W_\mu^+ (\partial_\mu \bar{X}^0 X^- - \partial_\mu \bar{X}^- X^0) + ig s_w W_\mu^+ (\partial_\mu \bar{Y} X^- - \partial_\mu \bar{X}^+ Y) + \\
& ig c_w W_\mu^- (\partial_\mu \bar{X}^- X^0 - \partial_\mu \bar{X}^0 X^+) + ig s_w W_\mu^- (\partial_\mu \bar{X}^- Y - \partial_\mu \bar{Y} X^+) + \\
& ig c_w Z_\mu^0 (\partial_\mu \bar{X}^+ X^+ - \partial_\mu \bar{X}^- X^-) + ig s_w A_\mu (\partial_\mu \bar{X}^+ X^+ - \partial_\mu \bar{X}^- X^-) - \\
& \frac{1}{2} g M [\bar{X}^+ X^+ H + \bar{X}^- X^- H + \frac{1}{c_w^2} \bar{X}^0 X^0 H] + \frac{1-2c_w^2}{2c_w^2} ig M [\bar{X}^+ X^0 \phi^+ - \\
& \bar{X}^- X^0 \phi^-] + \frac{1}{2c_w^2} ig M [\bar{X}^0 X^- \phi^+ - \bar{X}^0 X^+ \phi^-] + ig M s_w [\bar{X}^0 X^- \phi^+ - \\
& \bar{X}^0 X^+ \phi^-] + \frac{1}{2} ig M [\bar{X}^+ X^+ \phi^0 - \bar{X}^- X^- \phi^0]
\end{aligned}$$

FIGURE 2.12 The Lagrangian of the Standard Model of particle physics. (Courtesy of Thomas D. Gutierrez.)

Physical quantities have units. There are three fundamental units: length (L), time (T), and mass (M). Their actual values in this derivation do not matter since they depend on convention. Most people use the metric system: meters for length, seconds for time, and kilograms for mass. The United States, on the other hand, uses inches for length, seconds for time, and pounds for mass. This doesn't matter as we can convert from one to the other. For example, 1 in. equals 0.0254 m. In Canada where I live, height is usually measured in feet and inches even though they adopted the metric system in the 1970s. But in the end, it doesn't matter as we have these

conversion factors. No one but the insane or a mathematician will say: “Hey, I ordered a foot-long veggie hot dog and you gave me a 0.3048 one.”*

Velocity (V) has the units of length divided by time: $[V] = [L/T]$. The latter is not really a strict equation but relates the units of a quantity to the fundamental units, hence the brackets. For example, $[A] = [-A]$ makes sense for any physical quantity A . Usually in math, $A = -A$ implies $A = 0$.† Not the case with the brackets. Pressure (P) has units of mass divided by length and also divided by the time squared: $[P] = [M/L/T^2]$.‡ This follows from the fact that pressure is equal to the force divided by area. As most of you know: “force is equal to em times ah.” The units of “em” are mass [M] and the units of “ah” are that of an acceleration [L/T^2]. So, the units of force are $[F] = [M L/T^2]$. Therefore, the units of pressure are $[P] = [M L/T^2]/[L^2] = [M/L/T^2]$.

This was a bit involved and technical. But the bottom line is that:

The units of pressure are equal to mass divided by length and divided by time squared.

Given these relations, we want to derive an equation that captures exactly the principle that says “velocity is inversely proportional to pressure.” We can try: $V = -P$. Well that does not work since the units of $[V]$ are not equal to the units of $[P]$. Next, we could try $V^2 = -P$. In this case, the units of the velocity squared [L^2/T^2] are still not equal to the units of pressure [$M/L/T^2$]. But we are getting there.

Enter the density of the fluid. What is a density of a fluid? The units of density (D) are the amount of mass per volume: $[D] = [M/L^3]$. Air is less dense than water, for example. In fact, the ratio between these two densities is one to one thousand. Water is a thousand times denser than air. Imagine two bins of equal volume. One bin contains air and the other bin contains water. Guess which one is heavier? The one containing water of course because it is denser. Now, let’s try the following equation: $D V^2 = -P$. Bingo. The units now match: $[D] [V] = [M/L^3] [L^2/T^2] = [M/L/T^2] = [P]$.

Dimensional analysis is a powerful tool to find equations. It is also a good tool to verify if equations are correct. If the units do not match on

* When people ask me how tall I am here in Canada, I usually respond “six eight.” A European would have absolutely no clue.

† Good luck trying to find another number that has that property. Unless “numbers” live in a “field” with only two elements. Let’s call them “0” and “1.” In that case we have that: $0 + 0 = 0$, $0 + 1 = 1$, $1 + 1 = 0$, $0 \times 0 = 0$, $0 \times 1 = 0$, and $1 \times 1 = 1$. Mathematicians denote this field by \mathbb{F}_2 . The French call it \mathbb{F}_{un} . This is a bad pun.

‡ T^2 stands for T multiplied by itself: $T^2 = T \times T$ and L^3 stands for $L \times L \times L$.

both sides of the equal sign of an equation, then you know there is something fishy going on. It is then time to go back to the black or whiteboard.

We can now write Bernoulli's equation in fancy math symbols.*

$$\rho \frac{1}{2} v^2 + p = 0.$$

You might ask about that one-half factor in front of the velocity squared. This doesn't follow from our dimensional analysis. One way to determine this constant could be through experimentation. There are of course more conventional derivations which use kinetic and potential energy to get at this exact result. This is what you find in most standard textbooks. To take into account effects due to gravity, Bernoulli's equation has to be slightly changed to

$$\rho \frac{1}{2} v^2 + \rho g h + p_0 = \text{constant}.$$

It takes into account the effects of gravity through the number g which at sea level and in standard units is approximately equal to

$$g = 9.81 \text{ m/time squared.}$$

Bernoulli's equation also includes the effect of height denoted by h on the total pressure. The constant pressure denoted by p_0 is usually equal to the atmospheric pressure. The constant on the right-hand side is not really important as the equation is often used to relate pressures and velocities at different points in the flow so the constant is canceled out.

We can use these results to make some predictions about fluids. That is what science is all about. If the practical measures do not agree with the predictions of your theory, then too bad, your theory is *wrong* and useless.

The following is a textbook example that shows how to get concrete predictions out of Bernoulli's equation. It is normally known as *Bernoulli's tank*. Figure 2.13 shows the setup (with Bernoulli's original drawing on the right). We have an arbitrarily shaped tank filled up with some liquid,

* The symbol ρ is Greek for "r" and is pronounced "rho" and is often used to denote density, do not ask me why.

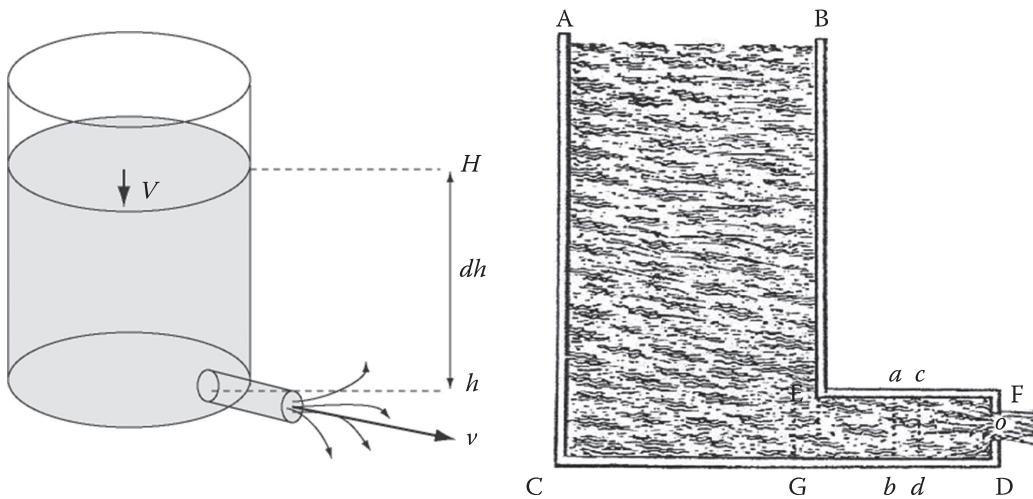


FIGURE 2.13 Bernoulli's tank.

let's say water. At the bottom, there is an opening that lets the water exit from the tank. As expected, the level of water will drop and eventually exit the tank. The question is: "at a given level of water what is the velocity of the water flowing out of the hole at the bottom?" We can answer it by using Bernoulli's equation. We have to make some simplifying assumptions, however. Both at the top of the liquid and at the bottom where the water flows out, Bernoulli's equation holds and therefore:

$$\rho \frac{1}{2} V^2 + \rho g H + p_0 = \rho \frac{1}{2} v^2 + \rho g h + p_0.$$

We get a simple equation by assuming that the velocity of the flow at the top is much slower than the bottom. It follows that we can assume (approximately) that $V = 0$. Let the difference between the heights be denoted by $dh = H - h$. Then we have that

$$g dh = \frac{1}{2} v^2 \text{ and consequently } v = \sqrt{2 g dh}.$$

Notice that we canceled equal terms from both sides of the equation like atmospheric pressure which is the same at the bottom and the top of the tank. The similar is true for the density which is also the same at both ends. The velocity at the outflow decreases with the height of the liquid.

This is a common experience when one is relieving oneself in a urinal. This might happen after *shotgunning* a can of beer. I had never done a shotgunning beer contest until 2003 after giving a keynote talk at a conference in Northern Sweden. I have heard that this practice is popular at North American colleges. Here is the experiment: you hold the beer can horizontally and punch a hole near the bottom of the can facing you. Then you slowly tilt the can with your mouth on the hole. When the can is vertically positioned, crack the can open and there you go. Amazingly, I was able to take it all in. In Sweden, they do this with half-liter beer cans. Some contestants were not so lucky.

Thanks to Bernoulli, you will know exactly what the velocity of the beer flowing out of the can is. It is wickedly fast! Kids do not try this at home.

More seriously, the velocity at the bottom hole is equal to the velocity of an object dropped from the same height (neglecting forces like drag).

To summarize: The Bernoulli family did rock. They did so much cool math and came up with a lot of cool equations. Kids do not try the shotgunning experiment until after you give a keynote presentation at some remote location in Northern Europe.

Euler–Newton Equations or Navier–Stokes Equations

WHAT AN AMOUNT OF famous name-dropping in a single title. All these math stars are depicted in Figure 3.1. The reason for naming them in the title is that collectively, they gave us a principled description of how fluids evolve over time. How cool is that?

But first.

3.1 LEONARDO DA VINCI

The merit of painting lies in the exactness of reproduction. Painting is a science and all sciences are based on mathematics. No human inquiry can be a science unless it pursues its path through mathematical exposition and demonstration.

LEONARDO DA VINCI (ITALIAN RENAISSANCE MAN)

Of course, one has to mention the quintessential *Renaissance man*, Leonardo Da Vinci (1452–1519), who painted the famous portrait of Mona Lisa on display at the Louvre museum in Paris, France. It is almost a *cliché* really to mention him in an artsy science book. He studied the motion of fluids and derived the first version of the *equation of continuity*, which is depicted in Figure 3.2 alongside his magnificent drawing of a turbulent flow. Visually, he understood that *incompressibility* leads to turbulent



FIGURE 3.1 From left to right, Newton, Euler, Navier, and Stokes.

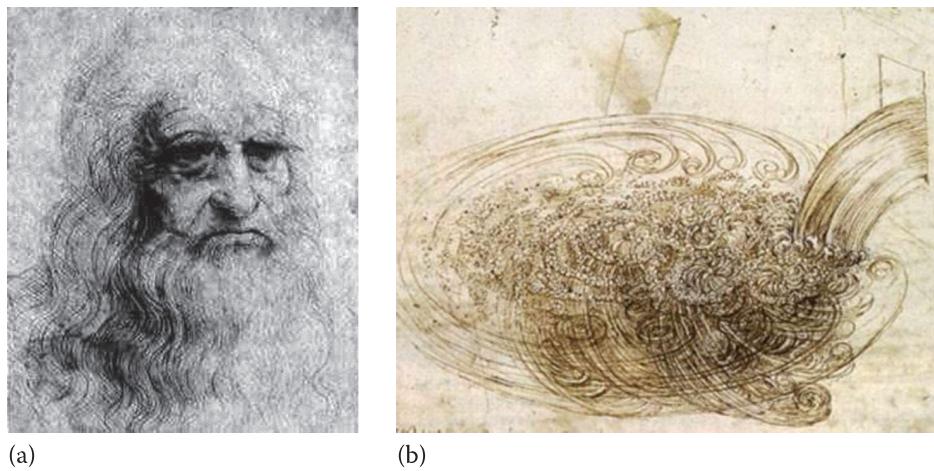


FIGURE 3.2 Leonardo Da Vinci (a) and his epic drawing depicting turbulence in fluids (b).

swirling looking fluids. Actually, he understood that *continuity* leads to swirling flows, which is the same thing, and will be explained later in this book. This is a key visual feature of any fluid.

Da Vinci was a hypercreative smart person. I feel bad only spending one paragraph on his achievements related to fluids.

To summarize: Da Vinci was a true Renaissance man and understood both the beauty and the fundamental principles of fluids. He could not animate them. But he could draw them beautifully and explain some of their key properties.

3.2 EULER AND CONTINUITY

Il cessa de calculer et de vivre.*

MARQUIS DE CONDORCET (FRENCH MATHEMATICIAN)
ABOUT EULER'S DEATH

* “He ceased to calculate and to live.”

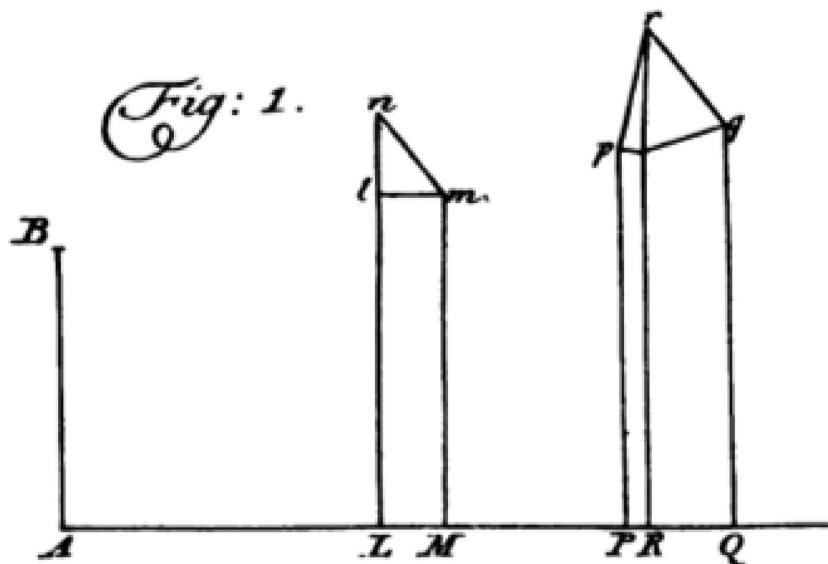


FIGURE 3.3 Euler's derivation of the continuity equation. (From pleasemakenote.blogspot.com/.)

As mentioned earlier, Euler and his amazing equations combine five fundamental mathematical constants. Personally, I felt that Euler should get most of the credit for finding the basic equations of fluid dynamics. It is no wonder why I chose his picture in Figure 3.1, looking away from the other three geniuses who are more somber looking. His contributions are a big piece of the puzzle. Euler's first insight was to derive an equation for the velocity of fluid at any point in space, assuming that the fluid is incompressible. A hand-drawn figure of Euler's original derivation is depicted in Figure 3.3.

What led Euler to formalize the motion of fluids into equations? First, he formulated an equation for the continuity or equivalently the incompressibility of the fluid. As he was primarily a mathematician, he first considered an *ideal fluid*, that is, a fluid that does not dissipate. If a fluid is set in motion, it will move forever, like the elusive perpetual motion machines. Figure 3.4 shows three examples of the so-called *perpetual motion machines*. Of course, all of them would eventually grind to a halt due to wear and tear and the effects of friction and damping. Even solar-powered devices will halt before the sun runs out of energy. And if not, they will halt at some point, when the sun eventually runs out of energy.

The beauty of mathematics is that we can just assume that such ideal fluids exist. What does continuity or incompressibility mean? At an intuitive level, it is possible to change the shape of liquid, but this does not mean that its volume changes, that is, a liquid takes the shape of its container but



FIGURE 3.4 Three examples of perpetual motion machines that have been proposed.

its volume does not change. Since it is easier to depict flat figures, consider a 2D fluid. Assume that a soapy foam residue floats on the surface of water in the bathtub. It is a good approximation of a 2D fluid. From the studies of Archimedes, it was observed that problem-solving methods were carried out in bathtubs. Figure 3.5 shows two shapes. If the fluid is ideal, then the area of the shape on the right has to be equal to the area of the shape on the left.*

The property of preserving the area is *global*. It involves both the entire shape and the entire fluid. Euler on the other hand came up with a *local* property which guarantees that the area is preserved for small pieces of the fluid. Piecing together these local conditions results in a global condition in practice. More details about this connection will be given later.

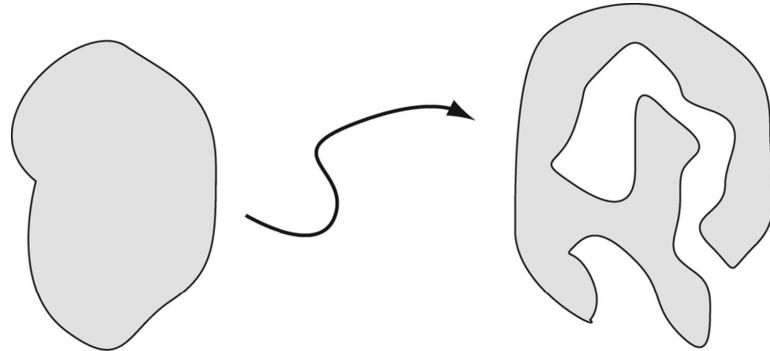


FIGURE 3.5 The area of two shapes is preserved exactly over time in an ideal continuous fluid.

* I tried my best in Adobe Illustrator to make the two shapes have an equal area. But hopefully you get the idea.

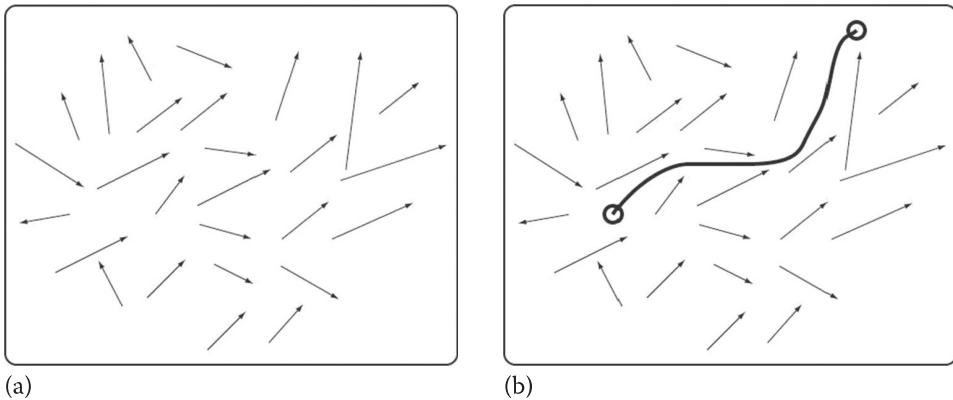


FIGURE 3.6 A 2D vector field (a) and a path of a particle through it (b).

At this point, we will introduce the key concept of a *velocity field*. Mathematicians call it a *vector field*. It can be observed initially from this book that visually some fluids such as air are nowhere to be seen. We cannot view the velocity of a fluid, but can only see or feel its effect on things. The way that mathematicians deal with this is to assign a fluid's velocity to every point in space. Imagine an arrow sitting at every point in space. Its direction indicates the direction of the flow and its magnitude indicates the speed of the flow. Figure 3.6a shows a vector field for a 2D fluid. Of course, the figure only shows a subset of all the vectors of the flow. Drawing all of the vectors would just result in a black mess. A complete black picture is not very helpful.

Imagine a tiny particle dropped in the fluid. The particle will then follow a path in time guided by the velocity field as shown in Figure 3.6b. Mathematicians would say “the path of the particle is tangential to each vector.” The velocity of the particle on its path depends on the magnitude of the vectors. You can imagine a cyclist that has to follow signs saying “go here,” “go there,” “slow down,” and “go faster.” But not “stop now.” It would have to be a fit cyclist, cycling forever.

Figure 3.7 shows a collage of photographs that I took in Venice, Italy. I felt like a *particella* trying to navigate myself using vectors in order to find the famous Rialto Bridge and the famous Piazza San Marco (or St. Mark's Square in English). I found both of them. Notice that the arrow in the center picture guiding me to San Marco is not very helpful. In this case, the flow is divergent and not incompressible. But I got to the square anyway! They should have a convergent arrow when you actually get to the square. These arrows were very helpful since it is extremely easy to get lost in Venice. I wish they had arrows guiding me to my hotel. This is when



FIGURE 3.7 Venetians have a vector field to send you to both the Piazza San Marco and the Rialto Bridge.

human communication comes in handy. Yes, I know that I could have used my smart phone, but data roaming rates are expensive when you do not live in Italy. Besides, you get to meet friendly people eager to help. No one sent me off in the wrong direction. I got to my hotel: *grazie!*

We all know what an arrow is. But how do you describe it exactly? A 2D vector is decomposed into two components, namely a sideway one and an upward one. These are both numbers that can be positive and negative. In three dimensions, there would also be an additional third number.

Figure 3.8 shows some examples of vectors in two dimensions. Each vector is defined by their sideway and upward direction. That is two numbers, $\text{arrow} = (\text{side}, \text{up})$. This is called a *tuple*. Just a fancy name for two numbers separated by a comma and squeezed in between two parentheses. *Side* and *up* are called the *coordinates* of the arrow. They are sometimes called Cartesian coordinates, named after the French philosopher, René Descartes (1596–1650). He also famously wrote “*Cogito ergo sum*”: “I think, therefore I am.”* Reality is an illusion fabricated by my mind. He was a solipsist of some sorts.

* I am. I am, I exist, I think, therefore I am; I am because I think, why do I think? I don't want to think any more, I am because I think that I don't want to be, I think that I... because... ugh!” Excerpt from the French existentialist philosopher Jean-Paul Sartre's (1905–1980) novel called *Nausea*.

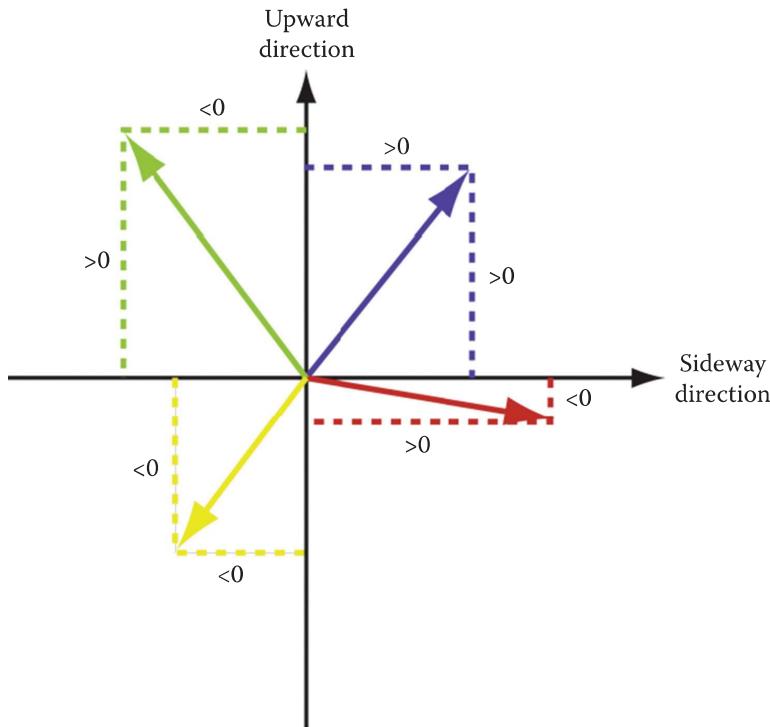


FIGURE 3.8 Four vectors in a 2D plane and their coordinates.

The numbers *side* and *up* can have any value both positive and negative. Consider again Figure 3.8. If *side* is positive, the arrow point to the right (the red arrow and the purple arrow). If *side* is negative, the vector points to the left (the green arrow and the yellow one). If *up* is positive, the vector points upward (the green one and the purple one). If *up* is negative, it points downward (yellow guy and the red guy).

The length of a vector is given by a result first stated by Pythagoras*:

The length of the arrow squared equals the length of its coordinate *side* squared plus its coordinate *up* squared.

Because we are taking squares, even if either *side* or *up* is negative, their squares will add up to a positive number: the length squared. Therefore, we can always use the square root to get the length. No need for awesome numbers yet.

In Figure 3.8, it is pretty easy to spot which colored vector is the largest. I think it is the green one. Someone else might say “No, no, no it is the red one.” Oh really? Since we have the coordinates and we can compute the

* Pythagoras of Samos (−570 | −495) is another Greek philosopher and mathematician. His theorem is surely known to anyone who has finished high school. In math speak: $a^2 + b^2 = c^2$.

length of both of these vectors using Pythagoras' result, the argument can be settled once and for all. I was right this time. Just as easily, I could have been wrong this time. But through math, the question is settled once and for all.

By the way, vectors can live in higher dimensions than two and three.

In a perfect fluid as shown in Figure 3.5, the area is preserved. What Euler showed is that it is equivalent to saying: "What flows in a tiny piece of the fluid has to flow out." What goes in has to come out. No sinks or sources are allowed in an ideal fluid for now. This is illustrated in Figure 3.9a for an arbitrary piece of fluid. To get an equation, we consider a *tiny* square piece of fluid as shown in Figure 3.9b. In this case, we only consider the flow through the horizontal and the vertical sides of the square. Vectors that point in a sideway direction are positive, so do the vectors that point in the upward direction. In Figure 3.9b there are four vectors and the continuity condition states that they have to sum up to zero:

$$a - b + d - c = 0.$$

Euler generalized this result. If u denotes the sideway velocity and v denotes the upward velocity, then this equation can also be written as follows:

$$u_{right} - u_{left} + v_{top} - v_{bottom} = 0.$$

In the case of Figure 3.9b, we have that: $u_{left} = a$, $u_{right} = -b$, $v_{top} = d$, and $v_{bottom} = -c$.*

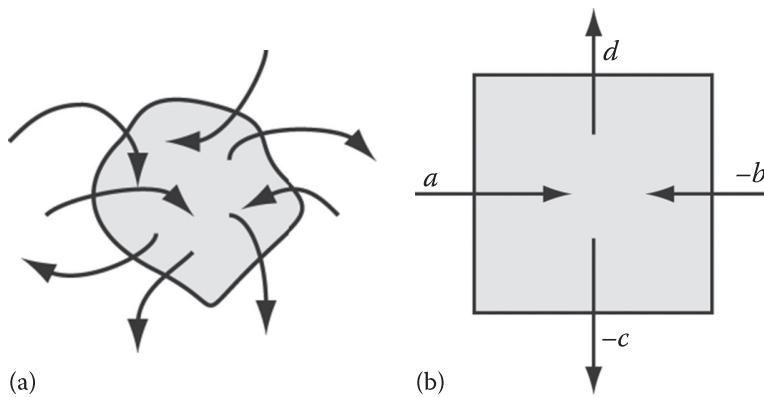


FIGURE 3.9 An arbitrary piece of fluid (a) and an axis-aligned square piece of fluid (b).

* If the reader got to this point she might have substituted the variables and noticed it is a different equation from the previous one, but it still holds for Figure 3.9b!

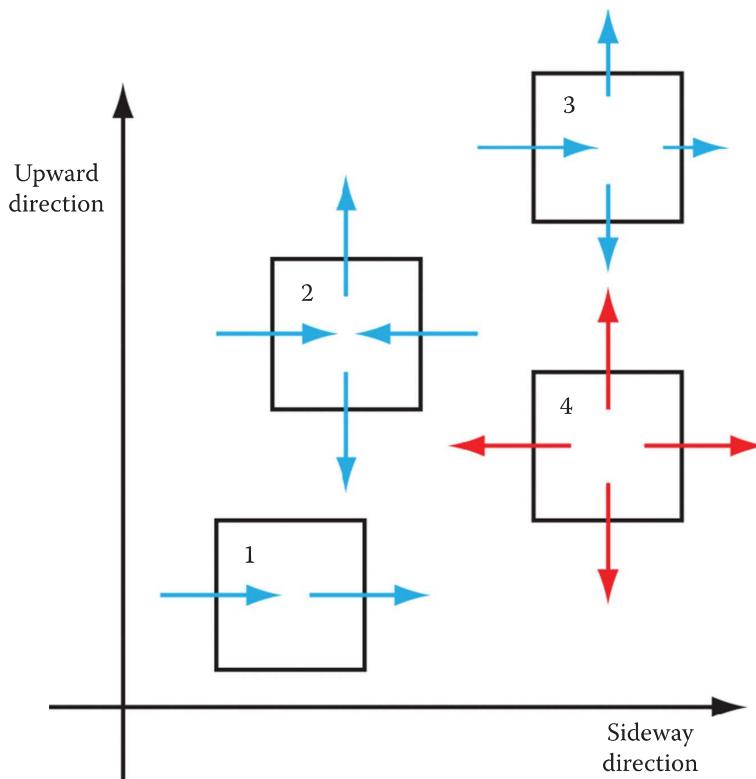


FIGURE 3.10 Euler's local continuity condition in various parts of the fluid.

For an incompressible flow, this relation has to hold for every square piece of a fluid. Figure 3.10 depicts four square pieces of the flow denoted as “1,” “2,” “3,” and “4.” Pieces “1,” “2,” and “3” satisfy the continuity condition while piece “4” does not. Piece “4” has all vectors pointing out of the square, and therefore, the square acts like a source of flow. This means that the flow expands, and therefore, it is not incompressible.

What about a 1D flow that is incompressible? In this case, the *squares* are line segments, as shown in Figure 3.11. In this case, the flow has to be constant. This is obvious because there is nowhere for the flow to go but in one direction. What comes in on one side of the segment has to exit on the

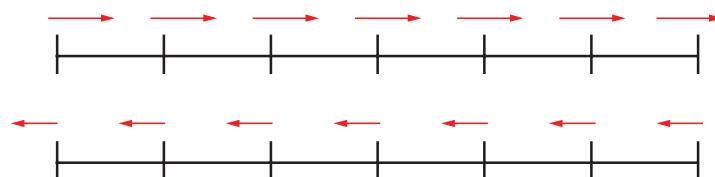


FIGURE 3.11 One-dimensional incompressible fluids are *trivial*: they must have a constant velocity everywhere.

other end of the segment. The flow is stuck in one dimension. Specifying the velocity at any point determines the flow everywhere and it is equal to that velocity. Mathematicians call these cases trivial, obvious, banal, easy, or of no interest.*

That was Euler's first contribution: what goes in has to come out. This result leads to trivial flows in one dimension but leads to highly complex flows in two and three dimensions as shown in Leonardo Da Vinci's drawing.

To summarize: Euler is one of the key players in animating fluids. We discussed continuity of a fluid and vector flow fields. These fields tell us where things should move and we want to preserve their mass. Euler gave us a condition on the vector flow field that guarantees mass preservation. It is very simple: what flows into a piece of the fluid has to come out. Beware though, it is a local condition not a global one.

Now, we will give a more geometrical interpretation of the continuity condition using what is called the *Helmholtz–Hodge decomposition* of a vector field.

3.3 INCOMPRESSIBILITY, CONTINUITY, HELMHOLTZ, AND HODGE THEORY

Ski racing, especially downhill, is a dangerous activity and there are many accidents. It would be really too bad to lose everything because of a crash.

HERMANN “THE HERMINATOR” MAIER (LEGENDARY
AUSTRIAN DOWNHILL SKIER)

The Helmholtz–Hodge theory of vector fields is a monumental result that sheds light on the solution of the incompressibility problem of fluids.

This theory is the brainchild of two geniuses: the German scientist Hermann von Helmholtz (1821–1894) and the Scottish mathematician William Vallance Douglas Hodge (1903–1975). They are both depicted in Figure 3.12.[†]

I like Helmholtz's work. He has also done some fundamental research in explaining visual perception as well. This is very relevant to computer

* Russian mathematicians call many results *trivial* even though it takes other mathematicians many pages to prove the result.

[†] At some point, mathematicians stopped wearing long, curly haired wigs.

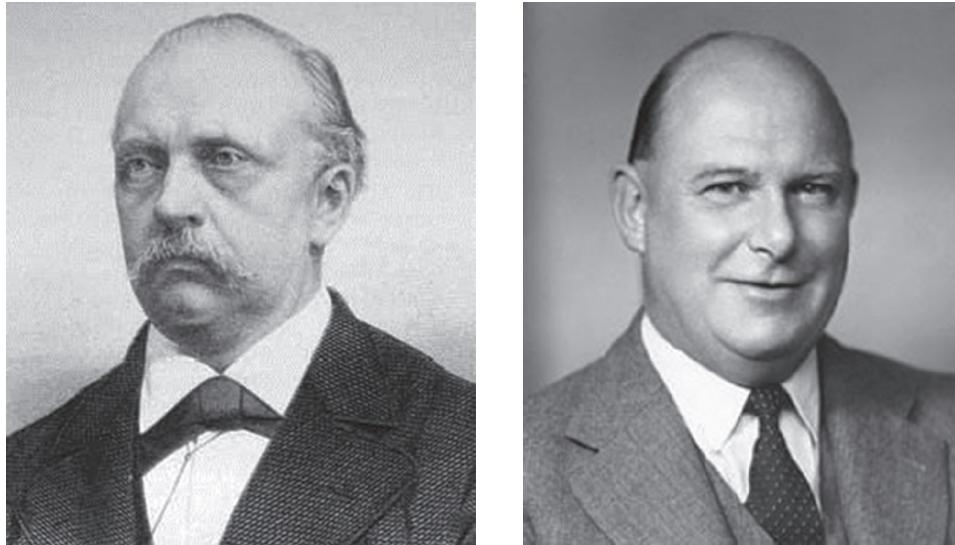


FIGURE 3.12 Helmholtz (a) and Hodge (b).

graphics. Our perception is not like a camera, it is more subtle. But I digress. He also worked in many other areas.

Hodge's name shows up in one of the problems that are part of the 10 Millennium Clay Prizes that we will mention later. The conjecture goes as follows:

The Hodge conjecture asserts that for particularly nice types of spaces called projective algebraic varieties, the pieces called Hodge cycles are actually (rational linear) combinations of geometric pieces called algebraic cycles.

I am not going to explain what this means because I do not know what it means exactly. Well sort of but not well enough. But if you solve this problem, you get a million bucks and fame. Good luck. Some people who call themselves mathematicians actually care about this stuff.

We will present the Helmholtz–Hodge decomposition in two dimensions. Remember that we want the flow that enters a piece of fluid to be equal to the fluid that comes out.

The problem simply stated is how do we turn an arbitrary vector field into an incompressible one?

The Helmholtz–Hodge decomposition solves this problem.

The result loosely states that:

Any vector field is the sum of an incompressible field and a gradient field.

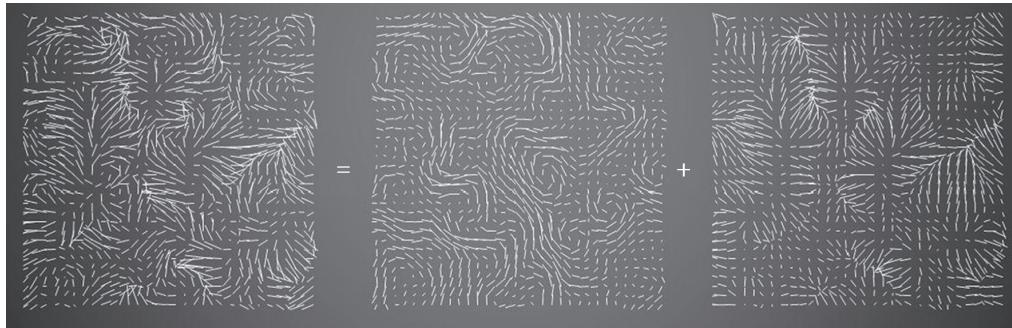


FIGURE 3.13 Any vector field is the sum of an incompressible field and a gradient field.

This is illustrated in Figure 3.13. The vector flow in the center is what we are striving for. It could be noted that it has nice whorls and vortices just like in Da Vinci's wonderful depiction of a flow in Figure 3.2b. The gradient vector field on the right-hand side of Figure 3.13 has exactly the opposite property: the flow is convergent and divergent at pretty much any point.

Bad flow! But still, have a treat, because we need you.

The strategy is to remove the bad part of the vector field to get a nice looking incompressible flow.

The procedure in order to get the nice flow is shown in Figure 3.14: original flow minus bad boy flow equals cool dude flow. The key ingredient in this decomposition is to figure out how to compute the (bad) *gradient field* from the original vector field.

How do we achieve that?

What is a gradient field anyway?

Consider the height field depicted in Figure 3.15. A height field is like a landscape. The height is just the height of that landscape measured from

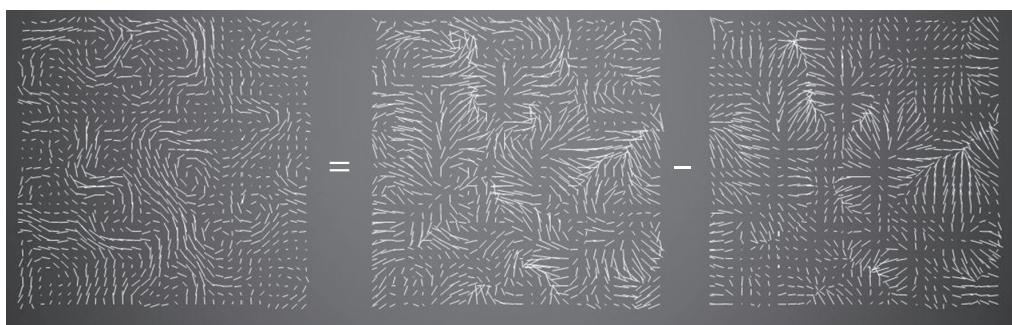


FIGURE 3.14 The incompressible field is obtained by subtracting the gradient field from the arbitrary vector field.

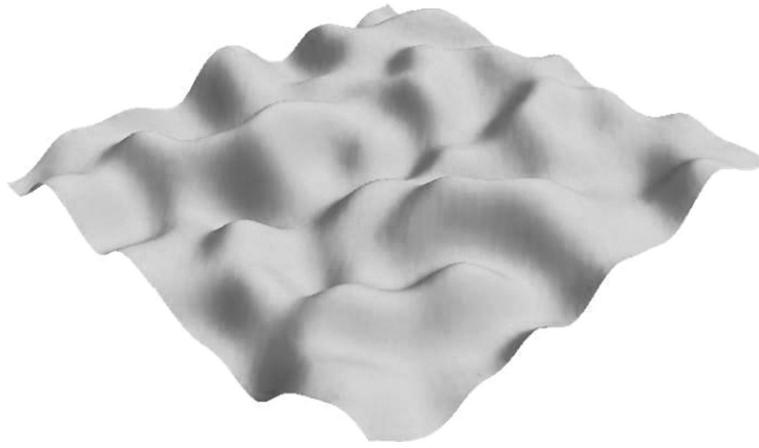


FIGURE 3.15 A height field. It is also periodic by the way.

some arbitrary chosen fixed elevation, for example, sea level. It does not matter because it is all about the change of height not the absolute values of height. There are peaks and valleys. If you wander in this landscape, there is always a pair of directions that will be the most efficient for you to move upward or downward. Hikers and skiers (if there is snow) usually do not choose those directions unless they are crazy and are into fast down-hill skiing or extreme mountain climbing.

The gradient field is handy for the crazies who are trying to achieve the biggest thrill of their lives in the least amount of time. I am not of one of them anymore. In Chamonix, France, they have a semilegit epic slope called *Le Couloir Poubelle*.^{*} There is even a gondola that takes you there. You make one mistake and you are going to be injured for sure. Not for the faint at heart. But if you get to the bottom alive and in one piece you definitely have *ski-cred*. Wear a headcam to prove you did it. I have gone down crazy slopes but not that crazy. I did not usually follow the gradient: the direction of steepest descent.

The gradient field is the opposite of the incompressible field. At most points, the flow is divergent or convergent as shown in Figure 3.13 (right).

Epic ski antics aside, the crucial problem now becomes: how does one compute the bad gradient field from the flow field we started with? The answer is that we have to find the height field first. Then, it is easy that we just choose the direction of steepest descent or steepest ascent. Some

^{*} *Poubelle* means garbage can in French. They used to drop garbage from the top lodge down this particular *coulloir*. *Coulloir* basically means passageway in French. It was the fastest way to get garbage down the hill. You have to climb over a fence to get to the poubelle. Look it up on YouTube.

crazy people do both and they climb up a mountain slope along the shortest path with their skis on their back first and then go down skiing along the steepest path.

The height field turns out to be the pressure of the fluid flow.

What?

Remember pressure is a single number not a vector. If we can compute the pressure, then its gradient is the bad as field we are looking for.

The true role of pressure in computing incompressible fluid flows was an epiphany for me: pressure is not an independent quantity for *incompressible* flows. That's it.

Pressure is just an intermediate quantity to enforce incompressibility in a fluid. For an incompressible flow, it is always possible to extract the pressure from the flow vector field if you like. Pressure of course is real. But from a computational point of view, it is just an aid to animate incompressible flows. That is how I understand it. When I give talks I always leave pressure out of equations. This sometimes makes for animated conversations afterward with some computational fluid dynamics experts wearing their backpacks.

“You are using the wrong equations, you missed the pressure term!”

“No, I did not since it is just a Lagrange multiplier to enforce the divergence free condition.”

“That is heresy!”

“But the animations looked pretty cool, no?”

“Yeah, but you are still wrong!”

“Whatever dude. Have a nice day.”

To summarize: Helmholtz and Hodge showed us how to get an incompressible field from an arbitrary field by computing the pressure of the velocity field. This is followed by a step correcting the velocity field by subtracting the gradient field of the pressure from it. We will see below how to compute this pressure field using a *fish equation*.*

* This is bad humor. It is actually named after the French mathematician called Poisson. *Poisson* means fish in French. Coders in the early 1980s were onto this and called their libraries written in FORTRAN77 “**FISHPACK**” to solve “elliptic partial differential equations.” The Poisson equation is a special case of these types of equations.

3.4 EULER AND THE MOTION OF FLUIDS

Strange fascination, fascinating me/Changes are taking the pace
I'm going through.

LYRICS FROM THE SONG “CHANGES” BY DAVID BOWIE
(FAMOUS BRITISH GLAM ROCK STAR)

Plus ça change, plus c'est la même chose.

FRENCH PROVERB

The velocity of a fluid tells us how things like points immersed in the fluid move over time. Now, we need an equation for the change of velocity over time. Welcome to fluid *dynamics*. The quantity that describes how the velocity changes over time is called the *acceleration* of fluid. It determines the change of what causes the change of the motion of a point in the fluid. Got that? Imagine the arrows in Figure 3.6 moving over time. Good thing that does not happen in Venice. Then, obviously, the point will take a different path over time than the one shown in Figure 3.6. The motion of the arrows over time is determined by the acceleration.

We now turn to Euler's second contribution to fluids: an equation of how the velocity of an ideal fluid evolves over time. An essential property of this description is that it is *nonlinear*.^{*} Usually, pretty much everything that happens around us is nonlinear. Linearity, to be linear, is actually the exception. But scientists like linear problems because they are easy to solve. Very often a nonlinear problem is locally cast into a linear one and *voilà!* The problem is solved approximately. But things can go horribly wrong, because by making a nonlinear problem linear you are essentially solving another problem.

What does it mean to be linear versus being nonlinear? Figure 3.16 illustrates the difference. The curves on the left and on the right cross the horizontal line. The left one is straight, and it is easy to figure out where it will cross the line. For this line, the problem is linear: *easy peasy*. Determining the line crossings for the curve depicted on the right-hand side of Figure 3.16 is more complicated. That curve even crosses the horizontal line multiple times. The curve is nonlinear. It is much harder to

* The term *nonlinear* has always annoyed me since nonlinearity is actually the norm; as in *non-not-normal*, double negative. Linearity should be called *not-normal* and *nonlinearity* should be called *normal*.

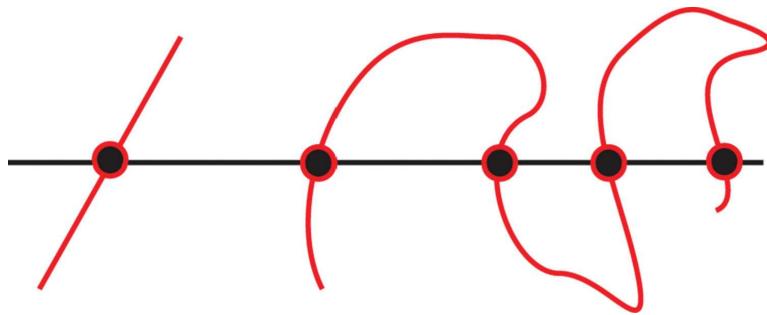


FIGURE 3.16 A linear curve (left) compared to an instance of a nonlinear curve (right).

figure out where it crosses the horizontal line than the linear case. Visually of course it is as easy. I was able to place the dots of intersection quite easily by visual inspection.

But computing their locations in a precise manner is much more difficult. There is no standard mathematical recipe that will work for all nonlinear cases. Choosing a method of solution is akin to a form of art. In many cases, one has to invent new techniques specific to a particular nonlinear problem. If you are lucky, these techniques will turn out to be useful for other nonlinear problems as well.

Personally, I like to explore very simple instances of a problem at first. I think and work like *a freak* apparently.* This is a good time to cite one of the most famous quotes from Albert Einstein. Drumroll.

Everything should be made as simple as possible, but no simpler.

There you have it. Einstein was a *freak*.

Euler's equations for the fluid's change in velocity over time are nonlinear.[†] Therefore, they are notoriously hard to deal with. The basic idea is to consider how things change under the influence of a vector field.

Euler's equations are all about change.

There are two ways to describe a change of something immersed in a fluid.

To illustrate this, let us start with a simple experiment involving a bathtub[‡] and temperature. Everyone who has taken a hot bath knows that it takes awhile for the hot water coming out of the faucet to reach the other

* Reference to the recent book by Steven D. Levitt and Stephen J. Dubner entitled *Think Like a Freak*.

[†] I will stick to the usual nomenclature and refrain from using the word *normal* instead of *nonlinear*.

[‡] Again! This time we will leave Archimedes out of the bath. No naked Greeks running down the street.

side of the tub. Before jumping in the tub you put your hand at the other side, doing so you can sense a change in temperature. *Bravo!* You just measured the change in temperature in an *Eulerian** framework. Why? You did not move your hand. You measured something like temperature in a fixed location.

Now let's imagine that you move your hand toward the input of hot water from the faucet. Obviously, the temperature will rise as you move your hand. *Bravo!* You just measured the change of temperature along a path in the bathtub in a *Lagrangian*[†] framework. Actually, this is a bit of a lie. In fact, Joseph Lagrange *only* allows you to move your hand along the flow of the water. So, a better analogy is that you measure the temperature of the water with your rubber ducky floating in the bathtub. But it is not as appealing as immersing your hand in hot water. I guess that we could attach a thermometer to the ducky and check out the reading by connecting the thermometer using Bluetooth to your iPhone as it bobs around in the bathtub. That is kind of geeky but it is Lagrangian.

In the first case, it took awhile for your hand to warm up, and in the second case, your hand got warmer way faster. This example shows that describing the change of something in a fluid is not so simple. Either you stay put or you move with the flow. But in both cases, you were measuring the change of something concrete over time: temperature. Just in a different way.

Mathematicians like to generalize stuff like that. There are other properties than temperature that can change in a fluid over time. For example, oil spills, the density of fish, the amount of foam, plankton, garbage, and so on. The motion of all these quantities changes over time while immersed in a fluid. Let us denote any of these quantities by the letter s . That is, the beauty of mathematical abstraction: s can stand for any of these things and much more. It can stand for pretty much anything that lives in a fluid really. This can even include the density of all the readers of this book.

In serious science, s has to be something you can measure, though. It cannot stand for the density of “three-headed and four-legged angels,” for example. On the other hand, if you are doing computer animation like we do you can pretty much immerse anything into a fluid including a “six-headed troll with one foot.”

* Named after Euler.

[†] Named after Joseph-Louis Lagrange (1736–1813). He was a genius French mathematician. Actually, he was born in Sardinia (now part of Italy) under the name Giuseppe Lodovico Lagrangia.

Let's get down to business, serious and geeky again.

We denote the Eulerian measure of change of some quantity labeled s by

$$\text{Euler}(s).$$

It measures the change of a quantity s at a single location in the fluid over time. Similarly, let us denote the Langrangian measure of change by

$$\text{Lagrange}(\nu, s).$$

Here, ν denotes the velocity of the fluid. The quantity s can be the temperature or the density of some dye injected into the fluid. I know that I said this before, but it is good to repeat it again.

And here comes the *coup de génie** of Euler.

No really.

Euler went one step further and considered s to be the velocity itself. This is when the nonlinearity business kicks in: the change of velocity depends on the velocity itself. He found an equation for

$$\text{Lagrange}(\nu, \nu).$$

The reason that this change is nonlinear is that the velocity shows up twice in the expression. The change is affecting the change. Sounds weird no? Nonlinearity is weird, and linearity is also weird because it is the exception. Therefore, nature is weird, but geniuses have found tools to harness some parts of nature's weird behavior.

Let's first consider the evolution of some dye of paint in an ideal fluid. Recall that ideal fluids do not dissipate, they roam forever, and thus, they preserve the area of anything dropped into them. Well, Euler formalized this by an equation that I will paraphrase as:

The density of paint will not change when following the flow of an ideal fluid.

And the math version is the *Lagrangian measure of change of the paint density flowing along the velocity is equal to zero*.

$$\text{Lagrange}(\nu, \text{paint density}) = 0.$$

* In English: *Stroke of genius*. It just sounds better in French.

The equation is more concise and saves us from typing in long complicated phrases.

Remember that Euler dealt with ideal fluids and also only dealt with an initial amount of paint. Consequently, the initial paint gets moved around mindlessly through the fluid without losing its initial density.

This is the ideal case with no external forces, sources of density, and sinks.

A good way to introduce external factors influencing fluids is through an artistic process called *paper marbling*. This is a technique that uses a liquid and various paints to create funky patterns. Figure 3.17 shows how they are created. And Figure 3.18 shows three examples. These *fluid textures* are usually created to decorate the first pages of some serious old

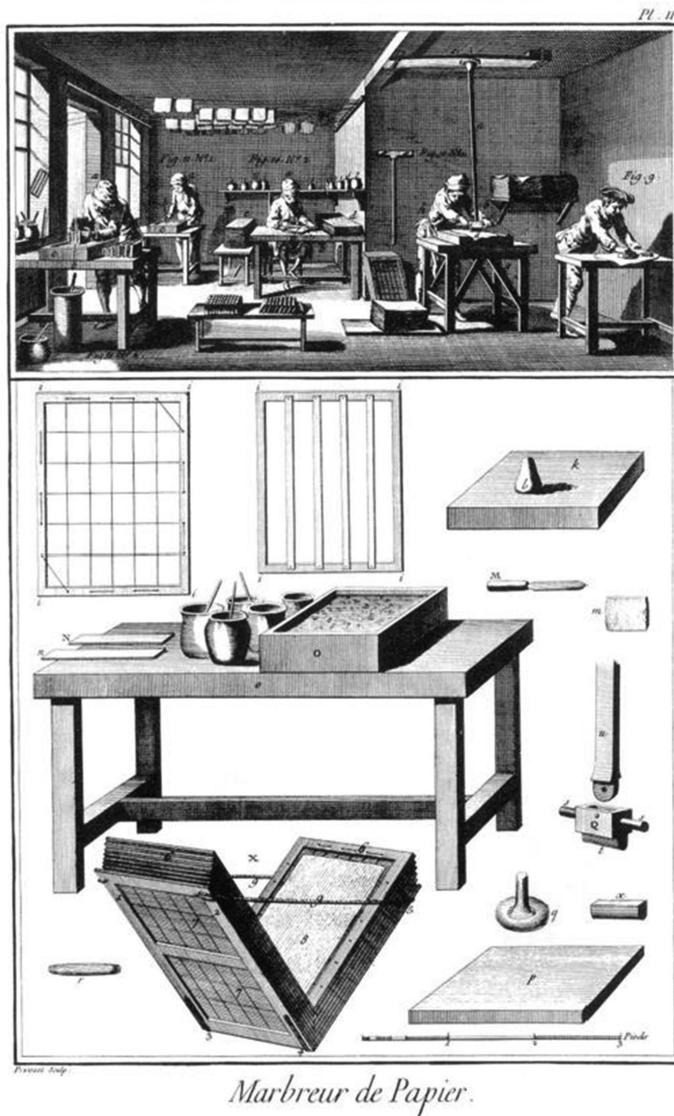


FIGURE 3.17 The process of creating marbling art.



FIGURE 3.18 Three examples of marbling art. Marbled endpaper from a book bound in France around (a) 1880 and (b) 1735 (detail). (c) Album cover “L’invitation au suicide” (1984 From New Rose Records, Paris, France).

books as shown in Figure 3.18a,b or the cover of the Punk/Goth French band from the 1980s called *Jad Wio* (Figure 3.18c). Alternatively, marbling can be practiced just for *art for art’s sake*.

Very briefly, this is how the marbling technique works.

A shallow tray is first filled with water as shown in Figure 3.17. Then, paints of different colors are injected into the fluid, normally with a brush. In many cases, chemicals are added to make sure the paints do not dissolve too fast. So things are closer to *ideal*. Patterns are then created by either blowing on the paint or using a fine filament like a swine’s hair to stir the fluid. The pattern is then transferred to a piece of paper or cloth.

The equation for the fluid’s motion in this case is like this:

$$\text{Lagrange}(\nu, \nu) = \text{blowing and stirring with a swine's hair}$$

It says:

Blow on and stir the fluid with a swine’s hair and the fluid’s velocity will evolve in a manner that we can predict.

Consequently, the distribution of the density of paints immersed in the fluid will be affected and create funky fluid patterns. Science meets art. Da Vinci would be thrilled.

More generally, Euler’s equation states that the change of velocity in an ideal fluid is caused by *external forces* like blowing, stirring with a swine’s hair, smashing, throwing stuff, spitting, cannonballing, and so on. All of these external effects are forces that we can add up and denote by a single vector field f . Therefore, Euler’s second equation can be written more generally as:

$$\text{Lagrange}(\nu, \nu) = f.$$

That is it for ideal fluids. This is not how Euler's second equation is usually written in mathematics, engineering, and physics books. It is my version. But this description is sufficient to explain how to write code that creates cool animations of fluids. I like to do things differently.

I cannot believe I haven't written a single partial differential so far in this book. It is so tempting.* We will show the usual notations when we go to the beach in Puerto Vallarta, Mexico below.

To summarize: Euler found the fundamental equation for the evolution of an ideal fluid over time. It involves a nonlinear process of velocities affecting velocities and being affected by external forces. Combined with his incompressibility condition, we have an exact description for an ideal incompressible fluid.

3.5 NEWTON AND VISCOSITY

I can calculate the motion of heavenly bodies, but not the madness of people.

SIR ISAAC NEWTON (FAMOUS ENGLISH
MATHEMATICIAN AND PHYSICIST)

Sir Isaac Newton (1642–1727) is of course famous for his three laws of motion and his invention of calculus. He made many more contributions in many other fields of mathematics and physics; most famously, his theory of gravity can predict the motion of planets and other bodies. Everyone knows the probably apocryphal story that he understood gravity when an apple fell on his head one day. This was another Eureka moment that changed science forever.

His contribution to the fluid puzzle is the introduction of the concept of *viscosity*: one of the things that make fluids nonideal. If you stir a fluid with a swine's hair, it will eventually come to a rest state. It depends on the fluid. It will take longer for water than honey. The difference is that water has a lower viscosity than honey. The less viscous a fluid is, the livelier its motion is. The more viscous a fluid is, the slower its motion is. Slam dancing is less viscous than line dancing. That is why I used to like slam dancing over line dancing.

* I feel like the French writer Georges Perec who wrote a novel called *La Disparition* that does not contain a single instance of the most popular letter “e” in the French language. It is actually a good read. This book has been cleverly translated into English under the title *A Void*. This footnote by the way used the letter “e” 39 times.

To describe viscosity, it is easier to start with a similar effect that involves simpler quantities like temperature or a density moving in a liquid.

And bugs!

Consider this simple experiment: drop a spot of ink in a glass of water. The dot of ink will slowly get fuzzier over time. Assuming that the spot of ink is red, then eventually the water in your glass will look sort of pinkish. The process of going from a concentrated dot of red paint that results in a pink mixture is called *diffusion*. This is a process, meaning that it describes something that evolves over time.

Sounds familiar right?

Diffusion is everywhere. Let us say you heat up a metal bar on one end. Then, eventually, the heat will diffuse to the other end and the bar will reach a uniform temperature. Smell is like that too. If a person with a strong perfume odor walks by you, it takes a certain amount of time before you can smell it.* It has to diffuse through the air. The closer you are to the person, the stronger the smell and the more immediate is that sensation.

The important point here is that perfume loses its potency when traveling through the air. In an ideal fluid, as we have seen earlier, it would just mix with the air but its potency would be unaffected.

Remember that the equation for some substance injected into an ideal fluid is

$$\text{Lagrange}(\nu, \text{perfume}) = 0.$$

This does not hold in the real world. We have to change that “0” to something else to account for diffusion.

How does one describe the process of diffusion?

Let's look at the motion of a single pesky bug crawling around your floor as depicted in Figure 3.19.[†] The bug just aimlessly crawls around in a seemingly erratic manner. Next, consider the case of many bugs all starting at a bounded location and set free to crawl as they wish. Their motions are erratic and independent of each other. These bugs are not ants, alright, with their complicated emergent behaviors. Our bugs are dumb, even collectively.

* Other examples include bad breath, the smell of grass after a rainstorm, or the scent of a fine Bordeaux wine when you sniff it.

[†] The usual example is that of a drunken mathematician attempting to get home from a bar.

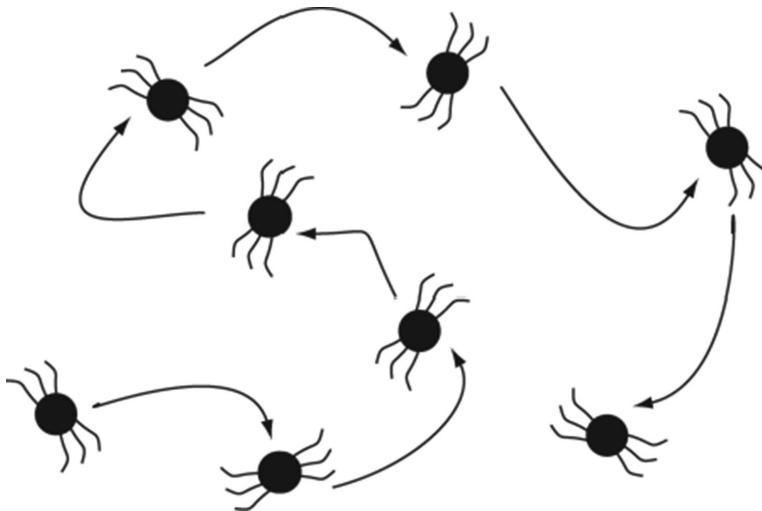


FIGURE 3.19 A bug crawling around in an incoherent pattern.

Mathematicians call the motion of these bugs as being *random*.^{*} They all perform a selfish *random walk* on their own: “Sorry excuse me, sorry excuse me, coming through. Thank you. I will just crawl over you.”

Now enters the concept of a *density of bugs*.

Where there are a lot of bugs the density is high, and where there are a few bugs the density is low.

Figure 3.20 shows the evolution of many bugs being initially dropped in a tiny location. Notice how the bugs disperse over time. For each snapshot of their random crawl, we show the circle that bounds all the bugs. No bug is outside of the circle. This is not some boundary from which they cannot escape. I did not say: “no bug is *allowed* outside of the circle.” It is just a measure of their spread. It is not a rubbery Petri dish. What one can observe is that the circle in general grows in size.

Consequently, since the number of bugs is constant, the density of bugs within the circle decreases. Yes, the *density* of bugs not the *number* of bugs. We are not killing any of them (yet). The bugs are just spread over a larger area.

More specifically, the number of bugs divided by the area of the bounding circle tends to decrease.

* According to the *Online Etymology Dictionary*: “having no definite aim or purpose,” 1650s, from at random (1560s), “at great speed” (thus, “carelessly, haphazardly”), alteration of Middle English noun *random* “impetuosity, speed” (c. 1300), from Old French *random* “rush, disorder, force, impetuosity,” from *randir* “to run fast,” from Frankish **rant* “a running” or some other Germanic source, from Proto-Germanic **randa* (cognates: Old High German *rennen* “to run,” Old English *rinnan* “to flow, to run”; see *run* [v.]).

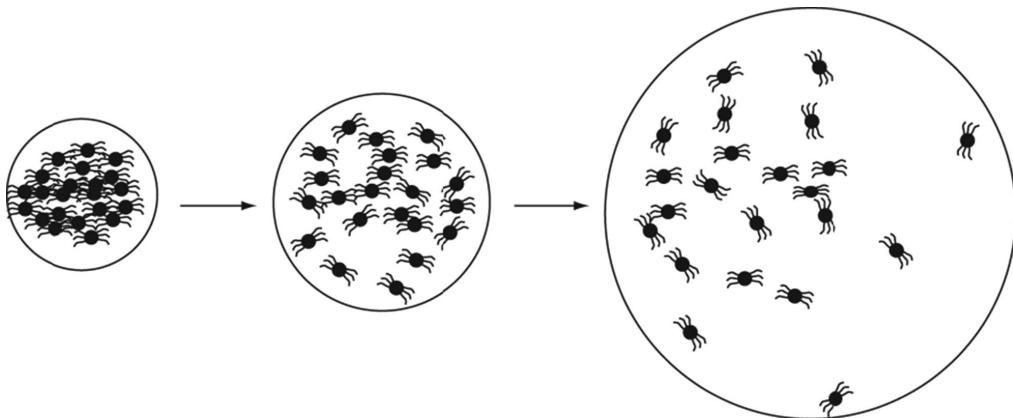


FIGURE 3.20 Many bugs crawling around over time.

That is diffusion in a nutshell.

Of course, we could have chosen drunkards, perfume particles, crack-heads, or dust particles instead of bugs, anything really that wanders about in a seemingly random manner.

There actually is a relation between the crawling of the bugs and the radius of the bug-bounding circle over time. What could affect the size of the radius? Obviously, a major factor is the speed of these pesky bugs that crawl around. If these bugs were all asleep nothing would happen. In that case, the circle would remain unchanged. Once the bugs wake up and start crawling around, the radius of the circle will start to expand at a rate related to their average velocity of crawling. Here is the actual relationship*:

$$R^2 = 4D_{thing} t.$$

where

R denotes the radius

D_{thing} characterizes the rate of diffusion of a *thing*

t denotes time, of course

It is instructive to find out what the units of D_{thing} are. From the equation, it follows that in terms of units $[L^2] = [D] [T]$. Therefore, the units of the rate of diffusion are $[D] = [L^2]/[T]$. Hence, D_{thing} characterizes how many *thingy*

* I am just stating this equation “as is.” Einstein derived it in 1905 in his paper where he made a good case for the existence of particles via *Brownian motion*. For 3D diffusion it is $R^2 = 6D_{thing} t$. In general, for a n -dimensional space the formula is $R^2 = 2n D_{thing} t$. This equation is probabilistic. It actually states that the *average* radius behaves in this way.

dinghies flow through an area over time. In this book, we will assume that it is a constant for a given phenomenon and it is the same for the entire population. It assumes that every bug behaves in the same manner.* This is less likely so for the hordes of drunken mathematicians being kicked out of the same bar after last call.[†] But one can always take an average of their behaviors and come up with a single diffusion rate. My guess is that

$$D_{\text{bugs}} > D_{\text{drunken mathematicians}}$$

Crawling bugs is the *microscopic* point of view of diffusion.

The *macroscopic* point of view on the other hand, only deals with densities and diffusion rates: of bugs, drunken mathematicians, or perfume particles and what not. This approach is typical in science: go from something *fundamental* and then move onto something more manageable by averaging or using some other form of approximation. This procedure sometimes goes by the fancy name of *homogenization*. Small-scale details, such as bugs crawling, are homogenized into a time-evolving density. Bye individual bugs and welcome to a bug milkshake.[‡]

At the macroscopic level, we do not care about the individual paths. We only care about the rate of diffusion D that characterizes the collective behavior of the individuals. When looking at clouds we do not see individual water droplets, rather we see white fluffy things floating and evolving in the sky.

The process of diffusion of a density over time is very akin to the condition of incompressibility that we mentioned earlier in this book. Remember incompressibility means: what flows in has to come out. The difference is that incompressibility is instantaneous. Diffusion at the macroscopic level is different. It is a process that evolves over time just like an ideal fluid.

As shown in Figure 3.21, each piece of space exchanges densities with its neighbors. It is a two-way process. The first piece is called *thing one*. It will give some density to another piece called *thing two* and *thing two* will give

* It doesn't have to be a constant. The diffusion rate can vary over time and spatial locations. Scientists can deal with it. But I want to focus on the essentials in this book without unnecessary complications.

[†] This bar would have to be in the middle of nowhere. Since sidewalks, buildings, cars, buses, trams, police officers, monuments, and garbage cans, for example, would get in the way of a truly drunken mathematically rigorous random walk.

[‡] Politics is like that too. Politicians don't care about a single individual. But they care about a large collection of the individuals who will vote for them.

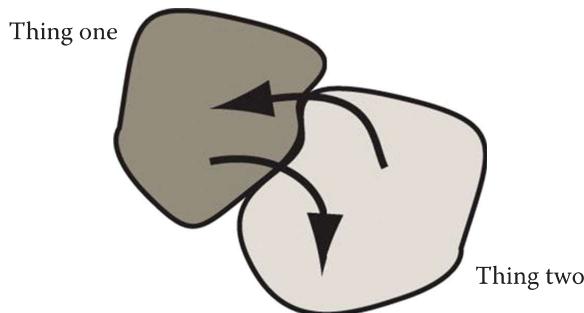


FIGURE 3.21 The densities of thing one and thing two exchange densities through diffusion.

some density to *thing one*. These exchanges go on over time, and eventually, every piece will have the same density. From then on, no exchanges will take place. That's it. The final state of the density is called the *steady state* and it is completely determined by the initial densities and the diffusion rate.

Calling these pieces *thing one* and *thing two* is of course a reference to the legendary book that was both written and illustrated by Dr. Seuss: *The Cat in the Hat Comes Back*. It is a perfect illustration of how pink liquid stuff does not lose mass but spreads all over the place. It diffuses in a playful manner throughout the story. It is one of the best introductions to diffusion that I can think of. It made a lasting impression on me ever since I read it as a kid.*

In between thing one and thing two, we can define a *flux* of density through their common boundary. Obviously, the flux depends on the diffusion rate and the difference in density between the pieces:

The flux of density equals the diffusion rate times the difference in density between the pieces.

This is actually known as *Ficks's law*. It is named after the German physiologist, Adolf Eugen Fick (1829–1901). In the following equation:

$$\text{Flux}_{12} = D \times (\rho_1 - \rho_2).$$

Consider two simple cases: if the density of thing two is zero and the density of thing one is nonzero, then the flux is in the direction of thing one to thing two. On the other hand, if the density of thing one is zero and the density of thing two is nonzero, then the flux is in the opposite direction. In general, the flux is a combination of these two cases. Consequently, there will be an exchange of densities between the two things.

* It also taught me basic English at a young age.

It is pretty intuitive what flux stands for. The name is aptly chosen. But then again what are the units of flux? Let us do the math: $[F] = [D][R] = [L^2/T][M/L^3] = [M/T/L]$. In other words, flux stands for the rate of transfer of the mass of bugs per unit length. So far so good. We have a description for the flux. What about the change in density caused by diffusion? That is, what we are looking for after all.

It is assumed that no density is lost, just exchanges of densities between things occur. The overall sum of densities remains the same. Remember that we do not kill our pesky bugs living in the microscopic world quite yet.

The flux determines what goes into and out of thing one. The same is happening for thing two. The loss or gain of thing one's density is determined by the flux.

Over some interval of time the density after equals the density before plus the flux between the pieces.

This is Fick's Second Law.

Naively, we can try this relationship:

$$\rho_1^{after} = \rho_1^{before} - \text{Flux}_{12}.$$

But this is not correct since flux does not have the units of density: $[M/L^3]$ does not equal $[M/T/L]$. It works out if we divide the flux by the area of the boundary A_{12} and multiply it by the time interval Δt between *before* and *after*. In that case, we get

$$\rho_1^{after} = \rho_1^{before} - \frac{\Delta t \times \text{Flux}_{12}}{A_{12}}.$$

This equation makes sense. The units now work out: $[M/L^3] = [T][M/T/L]/[L^2] = [M/L^3]$. The change of density depends on the previous density, the time interval, the area, and the flux.

And similarly, for thing two we have that

$$\rho_2^{after} = \rho_2^{before} + \frac{\Delta t \times \text{Flux}_{12}}{A_{12}}.$$

If we combine the two equations it is easy to notice that the total density is conserved:

$$\rho_1^{after} + \rho_2^{after} = \rho_1^{before} + \rho_2^{before}.$$

It is a fair game. This is true for all the other things that are adjacent to things one. It all adds up. That is the diffusion equation. Mass is conserved but it is redistributed between the many parts.

We can now introduce a diffusion function for thing one. It is a function of all the fluxes from all its neighboring things. Not just thing two. Let us say that thing one has n neighboring things. We are really interested in the change of density. A good way to measure change is to take the difference of the densities before and after of thing one over a time step.

$$\frac{\rho_1^{after} - \rho_1^{before}}{\Delta t} = -\text{Flux}_{12} - \text{Flux}_{13} - \dots - \text{Flux}_{1(n+1)} = D \times \text{Diffuse}(\rho_1).$$

Where we defined

$$\text{Diffuse}(\rho_1) = \frac{1}{A_{12}} \times \rho_2 + \frac{1}{A_{13}} \times \rho_3 + \dots + \frac{1}{A_{1(n+1)}} \times \rho_{(n+1)} - A \times \rho_1.$$

And

$$A = \frac{1}{A_{12}} + \frac{1}{A_{13}} + \dots + \frac{1}{A_{1(n+1)}}.$$

This was a kind of heavy going in order to describe diffusion mathematically. Intuitive phenomena are not always straightforward to describe precisely. Dr. Seuss did a much more fun job with his cat and the pink stain on the mother's dress.

These relations are for an arbitrary number of pieces exchanging densities. Let us specialize this to the particular situation shown in Figure 3.22. There the pieces are squares and have four neighbors they can exchange densities with. In this special case, $n = 4$. Concretely, the expression for the diffusion in this situation is as follows:

$$\text{Diffuse}(\rho_1) = \frac{1}{h^2} \times (\rho_2 + \rho_3 + \rho_4 + \rho_5 - 4 \times \rho_1).$$

This is a particular case. One can derive diffusion operators for many other configurations of things in a similar manner: 2D space divided into triangles, 3D cubical cells, 3D tetrahedra, 4D hypercubes, and so on.

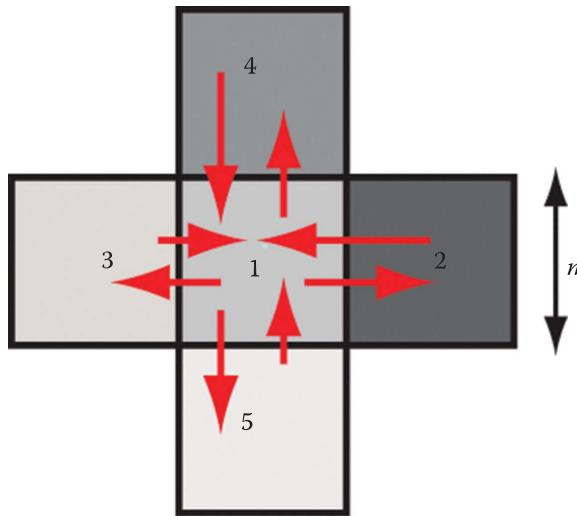


FIGURE 3.22 Diffusion between square pieces in two dimensions.

Combining the diffusion operator with the equation for the change of perfume in an ideal fluid we get

$$\text{Lagrange}(\nu, \text{perfume}) = D_{\text{perfume}} \times \text{Diffuse}(\text{perfume}).$$

So, perfume propagates both because it is carried along with the fluidlike air and because it diffuses. Both factors influence how we experience the smell of a person walking past us with a particular body odor. We replaced the “0” in our previous equation with the diffusion operator. This accounts for the decrease of potency of the odor.

Phew! We finally have something mathematicians call an *advection/diffusion* equation. Advection is another fancy word for measuring the change of something along the fluid’s velocity.

These equations describe the diffusion of quantities modeled by a single value over time and space. This is the macroscopic description of diffusion. No bugs.

What about viscosity that affects the velocity of a fluid? The difference with what we described earlier is that the velocity is a vector quantity. There are two components in two dimensions and three components in three dimensions.

Here is how Newton defined viscosity in his classic *Principia* from 1687:

The resistance which arises from the lack of slipperiness of the parts of the liquid, other things being equal, is proportional to the velocity with which the parts of the liquid are separated from one another.

The way I see it is that viscosity is simply the diffusion of each component of the velocity's field. Heresy! One can interpret the viscosity usually denoted by ν as the rate of diffusion of each component of the velocity.* Of course, this is a simple case. The viscosity might change over time or vary spatially. But again in the spirit of this book we want to keep things simple. In fact, our code uses the same method to animate diffusion for densities and to animate the fluid's velocity with viscosity.

It is the same thing really. Diffusion for temperatures and perfume particles tend to redistribute their properties over time. It is the same with the fluid's velocity. Any fluid that is stirred with a swine's hair will come to a perfect rest state.

So, without further ado here is what I call the *Euler/Newton equation*:

$$\text{Lagrange}(\nu, \nu) = \nu \times \text{Diffuse}(\nu) + f.$$

This equation pretty much describes most of fluid dynamics: the way fluids evolve over time. Of course, the condition of incompressibility has to be added to it. Remember: what flows in has to come out.[†]

To summarize: Euler came up with the equations for an ideal fluid and Newton added viscosity. Euler the mathematician and his perfect mathematical fluid model was complemented by the practical approach of Newton. Newton brought us back down to earth where things diffuse and dissipate.

3.6 NAVIER AND STOKES AND THEIR EQUATIONS

Science is a differential equation. Religion is a boundary condition.

ALAN TURING (BRITISH GENIUS MATHEMATICIAN
WHO LAID THE FOUNDATIONS OF COMPUTER SCIENCE)

So, what did Navier and Stokes contribute that Euler and Newton hadn't already done? The equation mentioned previously is usually called the *Navier-Stokes equations*. It is in the plural this time, because when they are written with fancy partial differentials there is a whole bunch of them. Three equal signs in this case for a 3D fluid. These equations are shown in Figure 3.23. They are the standard scriptures of fluid dynamics. They look

* This is another Greek symbol. This time it stands for *nu* basically the Greek version of n. Again do not ask me why they chose this symbol.

† At this point, some experts carrying their fluid mechanics luggage will ask: what happened to our dear pressure? If you want the pressure for an incompressible fluid, I can give it to you at any time. See below.

$$\begin{aligned}
 & \rho \left(\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} + v \frac{\partial u}{\partial y} + w \frac{\partial u}{\partial z} \right) = \\
 & \quad \rho g_x - \frac{\partial p}{\partial x} + \frac{\partial}{\partial x} \left[2\mu \frac{\partial u}{\partial x} + \lambda \nabla \cdot \mathbf{V} \right] + \frac{\partial}{\partial y} \left[\mu \left(\frac{\partial u}{\partial y} + \frac{\partial v}{\partial x} \right) \right] + \frac{\partial}{\partial z} \left[\mu \left(\frac{\partial u}{\partial z} + \frac{\partial w}{\partial x} \right) \right] \\
 & \rho \left(\frac{\partial v}{\partial t} + u \frac{\partial v}{\partial x} + v \frac{\partial v}{\partial y} + w \frac{\partial v}{\partial z} \right) = \\
 & \quad \rho g_y - \frac{\partial p}{\partial y} + \frac{\partial}{\partial y} \left[2\mu \frac{\partial v}{\partial y} + \lambda \nabla \cdot \mathbf{V} \right] + \frac{\partial}{\partial z} \left[\mu \left(\frac{\partial v}{\partial z} + \frac{\partial w}{\partial y} \right) \right] + \frac{\partial}{\partial x} \left[\mu \left(\frac{\partial u}{\partial y} + \frac{\partial v}{\partial x} \right) \right] \\
 & \rho \left(\frac{\partial w}{\partial t} + u \frac{\partial w}{\partial x} + v \frac{\partial w}{\partial y} + w \frac{\partial w}{\partial z} \right) = \\
 & \quad \rho g_z - \frac{\partial p}{\partial z} + \frac{\partial}{\partial z} \left[2\mu \frac{\partial w}{\partial z} + \lambda \nabla \cdot \mathbf{V} \right] + \frac{\partial}{\partial x} \left[\mu \left(\frac{\partial w}{\partial x} + \frac{\partial u}{\partial z} \right) \right] + \frac{\partial}{\partial y} \left[\mu \left(\frac{\partial v}{\partial z} + \frac{\partial w}{\partial y} \right) \right]
 \end{aligned}$$

FIGURE 3.23 The usual form of the Navier–Stokes equations in Cartesian coordinates.

pretty complicated right? Still they are simple compared to the *Lagrangian of the Standard Model* in particle physics. Remember Figure 2.12.

The Navier–Stokes equations are even scarier looking when the velocity of the fluid is described in *spherical coordinates*: each point is now described by two angles and a radius. The two angles determine a point on a sphere of radius one and the radius models by how much the point is removed from the center of the sphere. These equations are shown in Figure 3.24.

There is a more compact version however. I was able to write down the Navier–Stokes equations in the sand while on vacation on one of my many trips to the beach of Puerto Vallarta in Mexico as shown in Figure 3.25. I used part of a wooden stick I found to write a more elegant version of the Holy Scriptures. This is of course not my invention. Just a standard

$$\begin{aligned}
 r : \rho & \left(\frac{\partial u_r}{\partial t} + u_r \frac{\partial u_r}{\partial r} + \frac{u_\phi}{r \sin(\theta)} \frac{\partial u_r}{\partial \phi} + \frac{u_\theta}{r} \frac{\partial u_r}{\partial \theta} - \frac{u_\phi^2 + u_\theta^2}{r} \right) = - \frac{\partial p}{\partial r} + \rho g_r + \\
 & \mu \left[\frac{1}{r^2} \frac{\partial}{\partial r} \left(r^2 \frac{\partial u_r}{\partial r} \right) + \frac{1}{r^2 \sin(\theta)^2} \frac{\partial^2 u_r}{\partial \phi^2} + \frac{1}{r^2 \sin(\theta)} \frac{\partial}{\partial \theta} \left(\sin(\theta) \frac{\partial u_r}{\partial \theta} \right) - 2 \frac{u_r + \frac{\partial u_\theta}{\partial \theta} + u_\theta \cot(\theta)}{r^2} - \frac{2}{r^2 \sin(\theta)} \frac{\partial u_\phi}{\partial \phi} \right] \\
 \phi : \rho & \left(\frac{\partial u_\phi}{\partial t} + u_r \frac{\partial u_\phi}{\partial r} + \frac{u_\phi}{r \sin(\theta)} \frac{\partial u_\phi}{\partial \phi} + \frac{u_\theta}{r} \frac{\partial u_\phi}{\partial \theta} + \frac{u_r u_\phi + u_\phi u_\theta \cot(\theta)}{r} \right) = - \frac{1}{r \sin(\theta)} \frac{\partial p}{\partial \phi} + \rho g_\phi + \\
 & \mu \left[\frac{1}{r^2} \frac{\partial}{\partial r} \left(r^2 \frac{\partial u_\phi}{\partial r} \right) + \frac{1}{r^2 \sin(\theta)^2} \frac{\partial^2 u_\phi}{\partial \phi^2} + \frac{1}{r^2 \sin(\theta)} \frac{\partial}{\partial \theta} \left(\sin(\theta) \frac{\partial u_\phi}{\partial \theta} \right) + \frac{2 \sin(\theta) \frac{\partial u_r}{\partial \phi} + 2 \cos(\theta) \frac{\partial u_\theta}{\partial \phi} - u_\phi}{r^2 \sin(\theta)^2} \right] \\
 \theta : \rho & \left(\frac{\partial u_\theta}{\partial t} + u_r \frac{\partial u_\theta}{\partial r} + \frac{u_\phi}{r \sin(\theta)} \frac{\partial u_\theta}{\partial \phi} + \frac{u_\theta}{r} \frac{\partial u_\theta}{\partial \theta} + \frac{u_r u_\theta - u_\phi^2 \cot(\theta)}{r} \right) = - \frac{1}{r} \frac{\partial p}{\partial \theta} + \rho g_\theta + \\
 & \mu \left[\frac{1}{r^2} \frac{\partial}{\partial r} \left(r^2 \frac{\partial u_\theta}{\partial r} \right) + \frac{1}{r^2 \sin(\theta)^2} \frac{\partial^2 u_\theta}{\partial \phi^2} + \frac{1}{r^2 \sin(\theta)} \frac{\partial}{\partial \theta} \left(\sin(\theta) \frac{\partial u_\theta}{\partial \theta} \right) + \frac{2}{r^2} \frac{\partial u_r}{\partial \theta} - \frac{u_\theta + 2 \cos(\theta) \frac{\partial u_\phi}{\partial \phi}}{r^2 \sin(\theta)^2} \right].
 \end{aligned}$$

FIGURE 3.24 The Navier–Stokes equations in spherical coordinates.

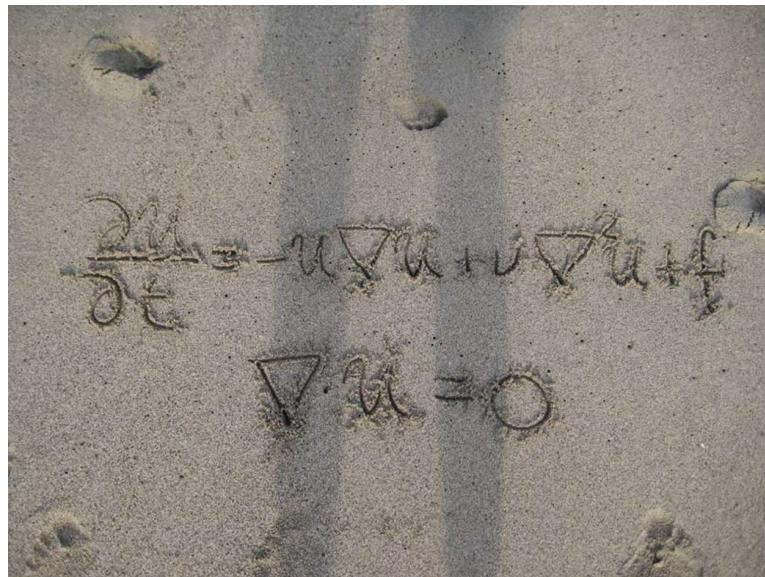


FIGURE 3.25 The Navier–Stokes equations written on the beach in Mexico.

condensed version. But it was fun to write these equations down on the beach of Puerto Vallarta.

Why did Navier and Stokes get all the credit and not Euler and Newton?

Claude-Louis Navier (1785–1836) was a French engineer who studied at the *École des Ponts et Chaussées*.^{*} He went onto become a mathematician, however, and he was the first person to combine Euler and Newton’s work to formulate the equations shown in Figure 3.23.

Independently, in Cambridge, England,[†] Sir George Gabriel Stokes (1819–1903) found the same equations but using more conventional techniques. His derivation is what you can find in standard textbooks unlike the *voodoo* derivations in this book that you are reading now.

There you go. That is the reason why the Holy Scriptures of fluid dynamics are called the *Navier–Stokes equations*. It is a perfect compromise between Continental French and Island British mathematical rivalry.

Just for fun I created a *bitstrips*[‡] cartoon depicted in Figure 3.26. This is a web-based cartoon creation tool that my daughter introduced me to. My daughter Gillian is featured in the cartoon alongside Lenny, Izzy, Claude, and George.

* In English: “School of Bridges and Pavements.” This is important stuff to study in order to help people move around in France: lots of rivers, mountains, lakes, monuments, and aggressive drivers.

† Stokes was born in Ireland, however.

‡ <http://www.bitstrips.com>.

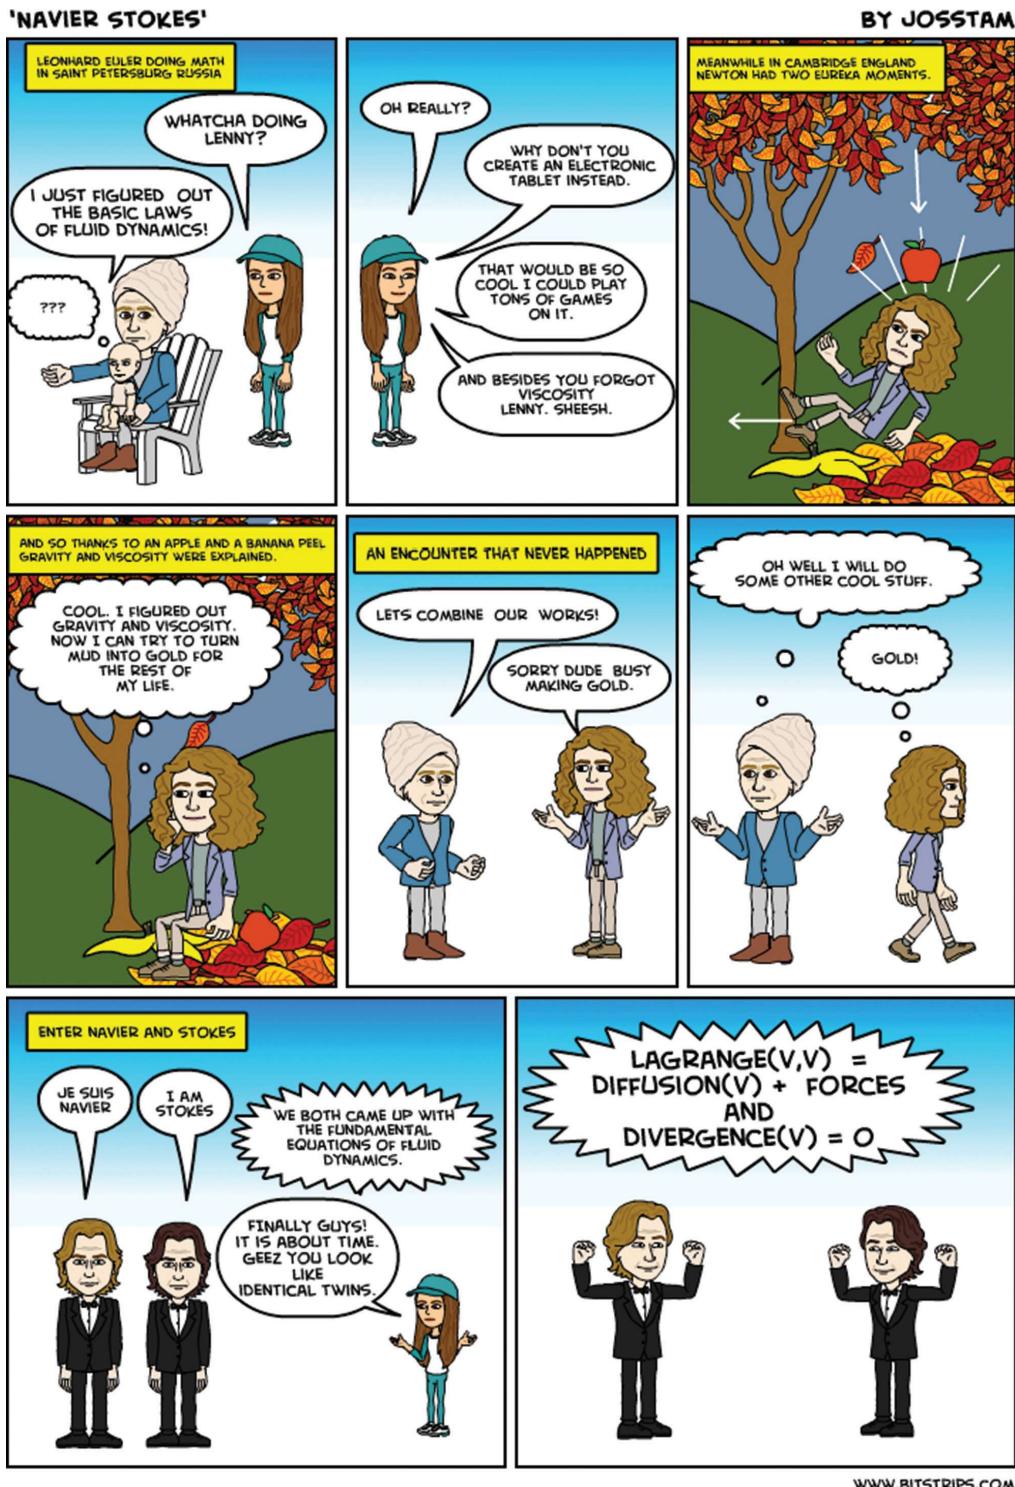


FIGURE 3.26 A brief history of the Navier–Stokes equations.

The Navier–Stokes equations can also make you rich and famous. In 2000, the Clay Mathematics Institute (CMI) proposed eight outstanding problems: the so-called millennium problems. Anyone who solves one of these problems will receive a million dollars.

Only one millennium problem has been solved thus far. This exploit is due to the eccentric Russian mathematician Grigori Perelman. He solved the so-called *Poincaré conjecture* which very roughly states that any 3D shape with no holes can be transformed continuously into a 3D sphere. This problem was previously solved for all other dimensions. Surprisingly, the Poincaré conjecture was easier to prove for higher dimensions. Higher dimensions provide more space and freedom apparently. There is more space to move stuff around.

Perelman did not care for the money and refused to accept the million dollars. Not only that but he also refused to accept the Fields Medal (remember the gold medal with the stern-looking Archimedes on it). Therefore, he did not get rich but definitely became famous. He is a humble man. As far as I can tell, his point is that he only closed some gaps in proofs that were based on a huge amount of previous work. Therefore, Perelman concluded that he should not get the prize. This is tricky. Research is never done in complete isolation. Apparently, Perelman now lives with his mom in a tiny apartment in St. Petersburg, Russia.

Another one of the millennium prizes concerns the Navier–Stokes equations. The problem on the official CMI web site* is stated as follows:

Prove or give a counter-example of the following statement.

In three space dimensions and time, given an initial velocity field, there exists a vector velocity and a scalar pressure field, which are both smooth and globally defined, that solve the Navier–Stokes equations.

In two dimensions, this problem was proven by the Russian mathematician, Olga Ladyzhenskaya, in 1969. Three-dimensional flows are inherently more difficult than 2D ones. Just as in the case of the Poincaré *spheres*.

The Navier–Stokes problem in three dimensions is still wide open. Recently, a mathematician from Kazakhstan named Mukhtarbay Otelbaev proposed a proof, but it is only available in Russian. Then, Terence Tao, a famous Australian mathematician who works at UCLA, proposed an

* <http://www.claymath.org/millennium-problems>.

argument that this problem could not be settled. Terrence Tao accepted the Fields Medal in 2006 for some other breakthroughs he worked on. The saga continues as of the writing of this book.

To summarize: Navier and Stokes put all the pieces together provided by Newton and Euler. If you can prove that these equations are well behaved, you might even reap a million bucks and be famous even if you are over 40 years old.

We now have to deal with the Achilles Heel of Fluid Dynamics: *boundary conditions*.

Boundary conditions are the hardest part of coding fluids simulations and getting it right. At least that is my experience.

In some of my other research, it is even trickier. The interaction of light between surfaces is all about boundaries: reflection, refraction, scattering, absorption, and so on. Those boundary conditions are much harder than the ones encountered in fluid animation. But that is not the topic of this book.

3.7 BOUNDARIES, BOUNDARIES, BOUNDARIES, BOUNDARIES, AND BOUNDARIES

I like pushing boundaries.

LADY GAGA (AMERICAN POP STAR)

It is interesting that viscosity, usually a damping mechanism, is here responsible for the generation of a geometrical progression of eddies.

KEITH MOFFATT (FAMOUS SCOTTISH FLUID
DYNAMICS EXPERT)

Thus far, we have described fluids without boundaries. We assumed that our bugs could crawl endlessly without bounds. Well guess what? The bugs can crawl all they want but they are confined to a room bounded by walls, humans, a dog, or a bed. This limits their ability to crawl. What follows is about what happens when the bugs hit the wall.

It turns out that there are many different boundary conditions for fluids. Let us list them straight ahead, and later, we will provide examples.

These are the five exclusive (XOR)* commandments for bugs at a specific boundary:

1. Thou shalt not cross this boundary.
2. Thou shalt not cross this boundary and be at rest when hitting this boundary.
3. Thou shalt be thrown into a fluid with a specific velocity through this boundary.
4. Thou shalt be free to escape the boundary.
5. Thou shalt be transported from this boundary to the opposite boundary.

In fluid dynamics speak, it is more like this:

1. SLIP
2. NOSLIP
3. INFLOW
4. OUTFLOW
5. PERIODIC

Alternatively, the different boundary conditions are illustrated in Figure 3.27 using pesky bugs.

In computer animation, the *slip* boundary condition rules. It just says that nothing can leave or enter a boundary. When you are pouring a liquid in a container you do not want the liquid to leak out of it. Unless of course you deliberately poke a hole in the container so that water will pour out on purpose. Remember Bernoulli's tank? Technically, when a bug encounters a wall it is only allowed to move sideways from the wall not through the wall.

* XOR stands for *exclusive or*. The way the word “or” is usually used is as an inclusive “or.” As in “Joe had a banana for breakfast” or “Joe had a latté.” Both can be true at the same time. Not so for the exclusive “or” denoted by XOR. For example: “Joe had a banana at exactly 10 a.m.” XOR “Joe had a pear at exactly 10 a.m.” Joe, like most people, cannot stuff a banana and a pear in his face at exactly the same time, especially at precisely 10 a.m.

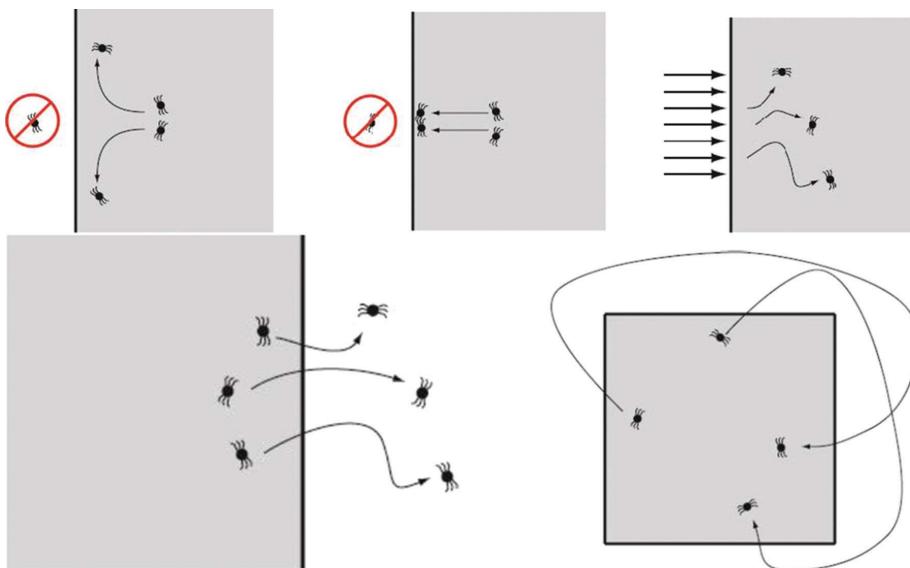


FIGURE 3.27 Boundary conditions for bugs. Slip, no slip, inflow, outflow, and periodic.

The *no slip* boundary condition means that the bugs will get stuck if they reach the wall. This is actually what happens in real fluids. Consider a ceiling fan. Check out the top of the blades. I can bet that there is dust on them. There is on mine. Even though the blades have been flapping around keeping you cool on those hot summer nights, they accumulate dust on their surface. When no slip rules, the dust particles like the bugs get stuck at the boundary. The velocity of the air at the boundary is zero. Hence, there is no escape. Until you dust the fan blades of course.

The no slip condition is the right physical thing to enforce. How come we do not use it in computer animation? The short answer is as follows: because it slows down the flow near boundaries. We do not like that. We want *lively* flows. In engineering and other *serious* applications, it is important however to enforce the no slip boundary condition.*

The *inflow* boundary condition means that at the boundary a certain velocity flows in. Think of a wind tunnel. It is a constant flow of air coming in from the boundary. Consequently, our bugs are blown all over the place as shown in Figure 3.27.

* Another fun story: the fluid feature of our MAYA software has a no slip boundary option. Almost 10 years after we released Fluid Effects, we did a test with no slip and it crashed the software. As far as I know, none of our customers had reported a crash in this situation. That is because no one went to the *Church of no slip*.

The *outflow* condition means that the bugs are allowed to leave the domain. Yes! We are free. However, little is known about their fate when they cross the boundary. Maybe their fate is like those poor lemmings that ran off a cliff. For incompressible flows with an inflow boundary condition, you need to have an outflow condition somewhere else. If not you will be in trouble, like when Pascal's barrel burst. But viscosity can come to the rescue.

The case of *periodic* boundary condition is interesting. It is a nonphysical condition, but it is useful in practice. In fact, I will describe a solver that entirely relies on this condition and that uses some very cool math later. In Figure 3.27 (bottom right), it seems that the bugs travel in outer space to reach the opposite boundary. This is not actually the case. All you have to do is to glue opposite sides of the square and now the bugs are free to crawl from one side to the other in a seamless manner. This situation is illustrated in Figure 3.28.

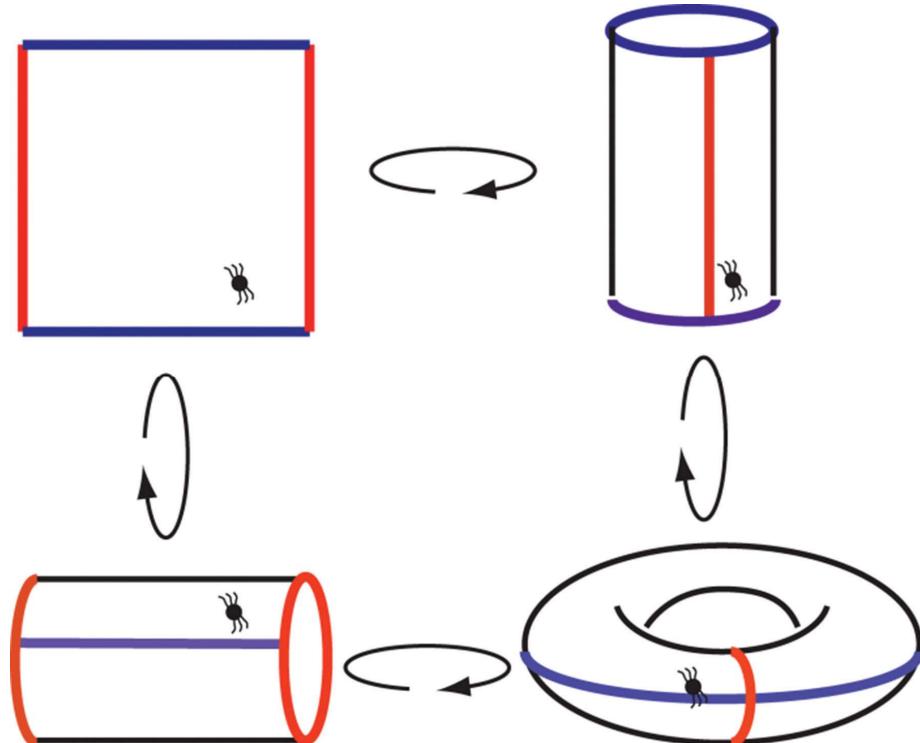


FIGURE 3.28 Periodic boundary conditions can be visualized by gluing together opposite boundaries. By first gluing two opposite boundaries, we get a cylinder (top right and bottom left). When gluing all opposite boundaries, we get a donut. The bug can wander freely.

Think of a donut with bugs crawling on it. Yikes. Mathematicians call a donut a *torus*. Any shape really with only one hole is a torus to a particular type of mathematician called a *topologist*.*

A torus and a donut are the same to a mathematician because they only have one hole.

There are other boundary conditions, of course, but these are the ones I have dealt with and turned into computer code.

The no slip condition is really the most realistic one. But it is not the one we use very often in computer animation. However, I want to spend some time describing one example that deals with a combination of boundary conditions called the *lid-driven cavity flow*. The setting is depicted in Figure 3.29.

This is an imaginary fluid made up by mathematicians. First, it is two dimensional. The velocity is denoted by the tuple (u, v) . Remember that it is a vector field described by two numbers. The left, right, and bottom boundaries have a no slip condition. And the top boundary has an inflow boundary condition. This is a bit of a misnomer, because in this case the flow is sideways. There is no outflow. But viscosity takes care of that. Viscosity will redistribute any inflow. In fact, it turns out that

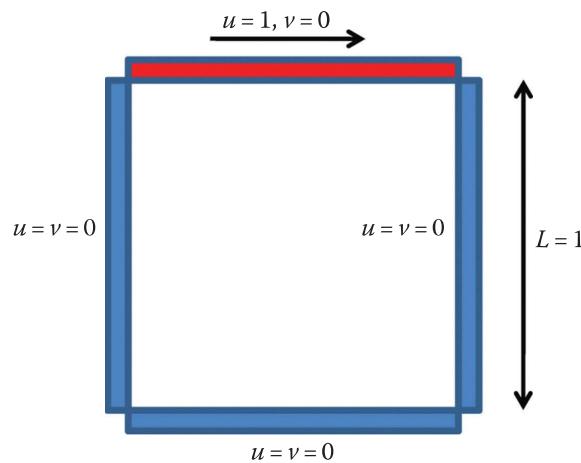


FIGURE 3.29 The setup of the lid-driven cavity problem.

* A popular albeit somewhat crude joke that mathematicians like to tell goes as follows. “A topologist is someone who can’t tell the difference between his ass and **one** hole in the ground, but who can tell the difference between his ass and **two** holes in the ground.” A respectable engineer at a math conference I attended started his keynote talk with this joke. LMFAO.

the flow reaches a *steady state*. After awhile, the flow settles into a stable pattern like we observed with diffusion. This pattern is a function of the viscosity of the fluid.

Ahem, not quite. We have to introduce a fundamental quantity that characterizes the behavior of a fluid first.

Sir George Gabriel Stokes (the Stokes in Navier–Stokes) introduced a *dimensionless* quantity called the *Reynolds number* named after Osborne Reynolds (1842–1912) who popularized the concept. Dimensionless, “that which has no units,” means no meters, no kilograms, and no seconds. It is just a *pure* number. Here is the definition:

$$\text{Reynolds number} = \frac{\text{typical velocity} \times \text{typical length}}{\text{viscosity}}.$$

Let’s verify that it really has no units. Velocity has units of [L/T], length has units of [L], and viscosity has units of [L²/T] (remember it is a diffusion rate). So if we put it all together, we get that [Reynolds number] = [L/T][L]/[L²/T] = 1. So, yes indeed the Reynolds number has no units. Why do we care about the Reynolds number? The reason is that this number allows us to compare different flows with the same Reynolds number. If they have the same Reynolds number, they sort of behave in the same manner.

To illustrate this concept, let us fill up a bathtub and let the water settle.* Now imagine filling up a gigantic pool the size of Switzerland with water. The viscosity of the bathtub and the pool are the same since they both contain water (no soap this time). This situation is shown in Figure 3.30, and on the right-hand side, there is an outline of the country of Switzerland and below it is a gigantic pool. Of course the bathtub and the pool are not drawn to scale. The bathtub is 2 m in length (I wish I had one of these) and the Swiss pool is approximately 200,000 m wide. So, 100,000 bathtubs lined up can cross Switzerland without accounting for the Alps of course. We assume that Switzerland is as flat as Holland in this case.

Everyone agrees that it would be quite a challenge to build a pool the size of Switzerland.[†] The beauty is that we do not have to. We can do experiments in our bathtub that can be applied to the gigantic Swiss pool as well.

* This time it will involve a Swiss person not a Greek.

[†] It might be possible in Saudi Arabia but not in Switzerland.

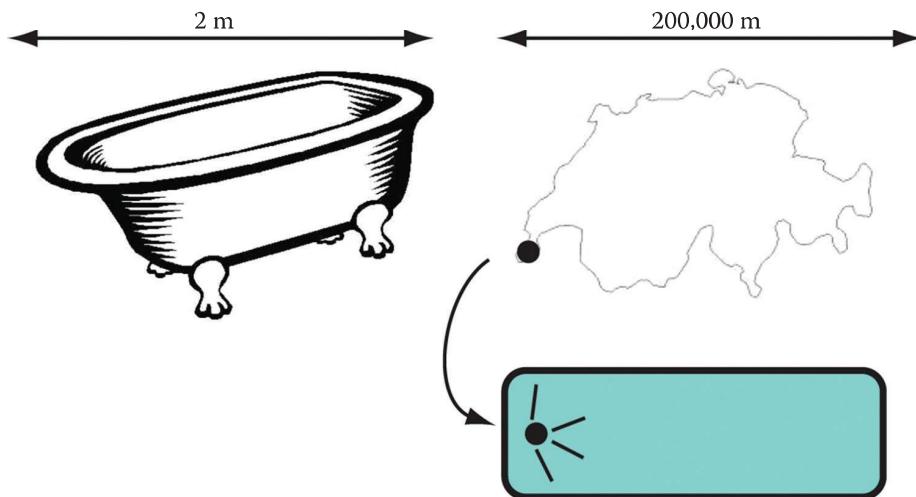


FIGURE 3.30 A bathtub and a pool the size of Switzerland.

How?

We just have to make sure that the Reynolds number of the bathtub is the same as the Reynolds number of the Swiss pool. Since the viscosities are the same and the lengths are given, we are left with the following condition on the velocities by equating their Reynolds numbers:

$$\text{Velocity of the bathtub} \times 2 = \text{velocity of the pool} \times 200,000.$$

Let's say you wanted to measure the effect of a *Genevois** cannonballing on the left side of the Swiss pool and calculate the effect that it causes. Let us additionally assume that the velocity created by the impact is 1 m/second in the gigantic Swiss pool. Thanks to Reynolds, we can reproduce this experiment in the bathtub. However, it is going to get messy. In order to get the same Reynolds number in the bathtub, the velocity caused by the impact would have to be

$$\text{Velocity of the bathtub} = 1 \times \frac{200,000}{2} = 100,000 \text{ m/second.}$$

That is wickedly fast and probably impossible to accomplish. Compare this velocity with the speed of sound which is approximately

* A Genevois is someone from the city of Geneva, which is located all the way to the left of Switzerland. The size of the Genevois is not drawn to scale. In the picture it seems like we are throwing the entire town of Geneva in the Swiss pool.

340 m/second.* Flows that are faster than the speed of sound are called *supersonic*.† Bullets (fired, not thrown) are supersonic, for example. I will not deal with such flows which are quite complicated. I like to keep things simple. There are a few depictions of supersonic fluid flows in Milton Van Dyke's book.

Another way to replicate the giant Swiss pool experiment is to fill the bath with a liquid that has a lower viscosity than water. Hence, in this case, the velocities and the lengths are fixed but we have to search for a fluid that has a much lower viscosity than water. How much lower? We can use the expression of the Reynolds number again. In this case, assuming that we stir both the gigantic pool and the bathtub with a swine's hair at 1 m/second, we get that

$$\frac{2}{\text{viscosity of the fluid in the bathtub}} = \frac{200,000}{\text{viscosity of water}}.$$

Therefore, the viscosity of this imaginary fluid that we will pour in the bathtub would have to be 100,000 times smaller than the viscosity of water. We could try air since it is a fluid but it is only a thousand times less viscous than water.

In conclusion, I do not think the bathtub will do the job to replicate the gigantic Swiss pool experiment.

What about an Olympic-sized pool which is 50 m in length? Now, we need to move the water in the bathtub to a velocity of 25 m/second. This is more manageable but still it is quite fast. Do not try this at your local community center.

To summarize: The Reynolds number characterizes one aspect of the behavior of a fluid. It helps to realize fluids with different scales and viscosities. But it is not always feasible in practice.

Now, let us get back to our lid-driven cavity experiment. Remember that the flow is two dimensional and has to be zero on three boundaries

* It is fast, but not as fast as the speed of light. Assuming your cell phone connection is speedy. Then it will take less time for your voice to travel to someone else's phone across the street than to yell at the same person directly.

† Here is another dimensionless number for you: the *Mach number*. It is the ratio between some speed and the speed of sound. Supersonic means that the Mach number is bigger than one. This is also known as the *sound barrier*. Chuck Yeager was the first pilot to go past the sound barrier by flying a custom made Bell X-1 in 1947. His Mach number was 1.06.

and it is stirred sideways at the top, a mathematician's concoction. This is a famous benchmark for fluid codes. It means that you spend weeks writing fluid code and now you want to know whether you are doing it the right way. You enter the lid-driven cavity boundary conditions in your program and then run it and check out the results. Next, it is important to check whether your results agree with the existing literature. These papers have been scrutinized and judged by other experts and eventually end up being published in journals.* The results in these papers are usually created using computer code. Some papers include actual physical experiments and these are approximations since 2D fluids, strictly speaking, do not exist in nature.

Figure 3.31 shows the results for the lid-driven cavity flow for different values of the Reynolds number from a respectable source—starting from a Reynolds number equal to zero (top left) to a Reynolds number equal to 10,000 (bottom right).

The lines in the pictures in this figure are *streamlines*: paths traced out by bugs or rubber duckies trapped in the liquid. They are *portraits* of the flow for different values of the Reynolds number. This is one of the many ways one can visualize a fluid. One thing we can notice is that more vortices appear as the Reynolds number is increased. This phenomenon is called a *Hopf Bifurcation*[†] by mathematicians.

Actually, for every value of the Reynolds number, the portraits in Figure 3.31 have an infinite amount of vortices. They just get exaggerated when the Reynolds number is increased.

Figure 3.32 shows a cartoon of the lid-driven cavity flow. Because of the discontinuity at the two lower corners, a cascade of vortices spinning in opposite directions is created. The vortices diminish in strength as well and there are an infinite number of them. In the right of Figure 3.32, we show a blowup of the right bottom. And this can be repeated *ad infinitum*. As with Bernoulli's infinite sums, this goes on and on. Of course in actual experiments, the smaller vortices eventually get blurred due to the effect of turbulence. Mathematically, they all still exist.

* This is the *peer reviewing* process. Publish or perish. These are not journals you can buy in your usual corner store. At any rate, these days everything is online. They are expensive though: about \$20 for an article. But if you are a student at some university or an adjunct professor as I am, it is free. So kids, stay in school!

[†] These bifurcations are named after the great German mathematician Heinz Hopf (1894–1971).

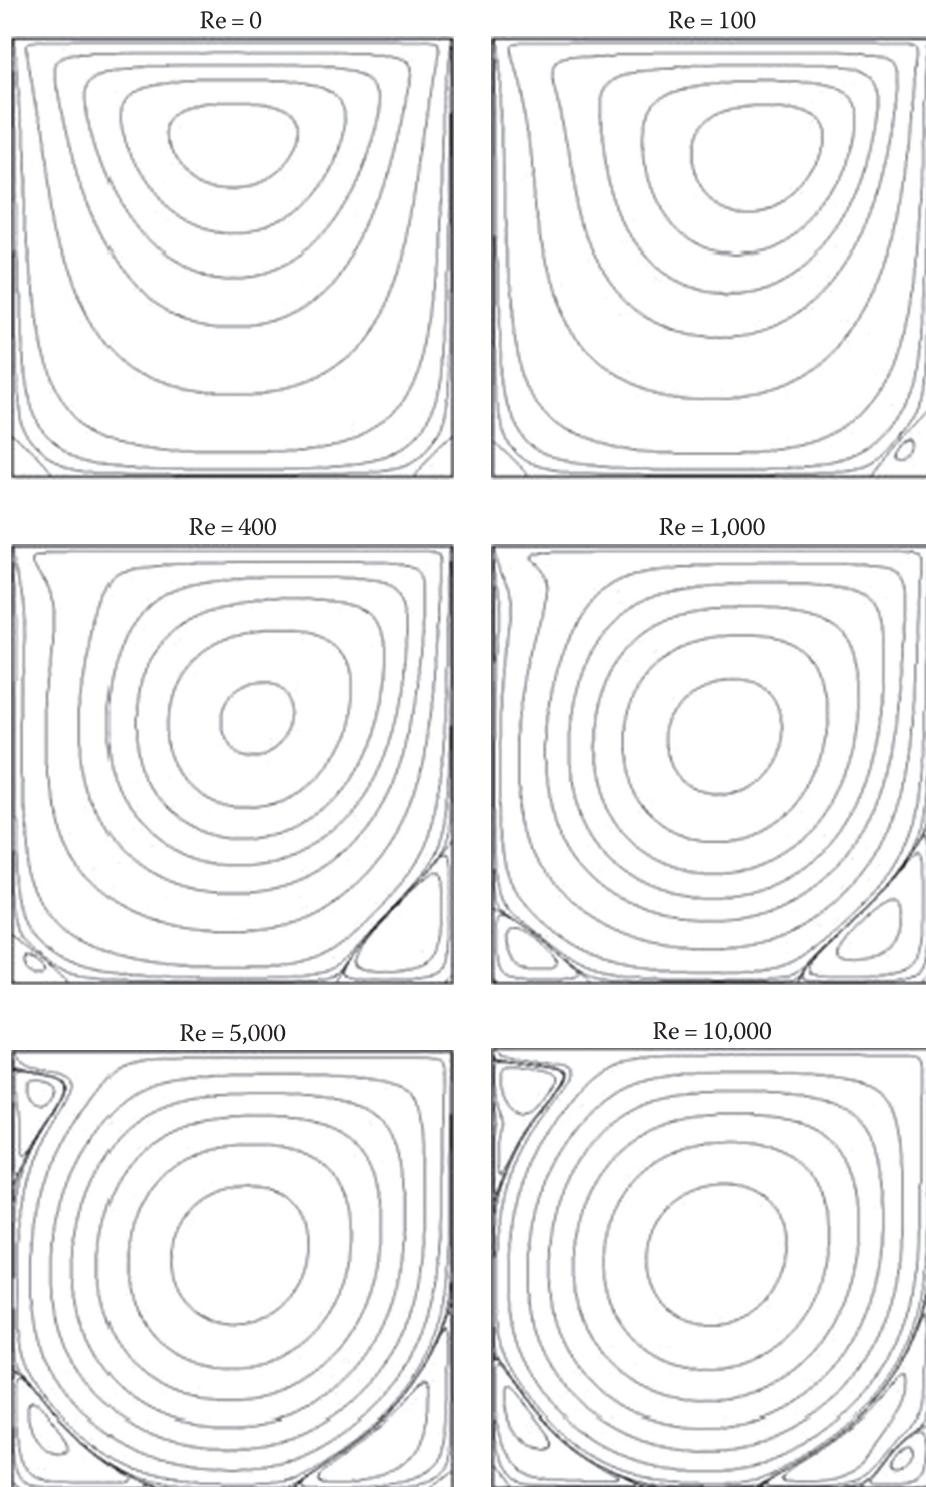


FIGURE 3.31 Lid-driven cavity solution for different values of the Reynolds number. (From Samin, M. and Owens, R.G., *Int. J. Numer. Methods Fluids*, 42, 66, 2003. With permission.)

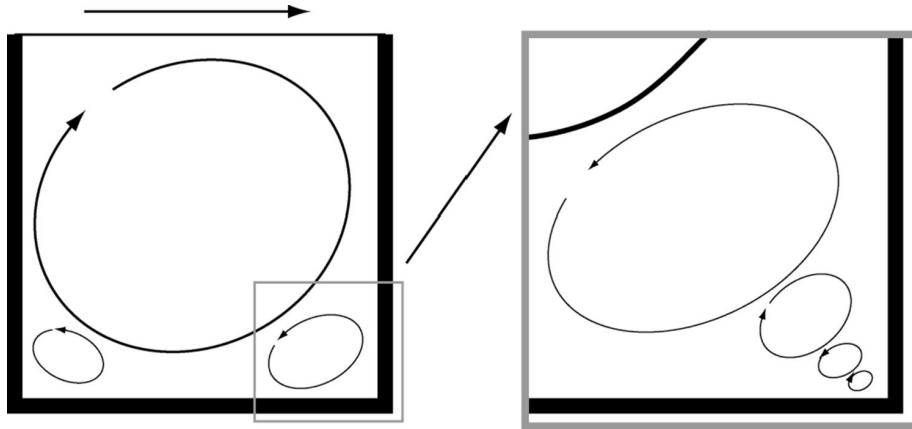


FIGURE 3.32 The lid-driven cavity flow generates an infinite number of counter-rotating vortices.

It is hoped that this example shows that boundaries are tricky and can give rise to exotic behaviors. Do not underestimate the role that boundaries play in trying to write fluid animation code. It is the hardest part.

To summarize: Boundary conditions are tricky. Do not underestimate them. They are only lower-dimensional boundaries but they affect the motion of the larger fluid contained in them in a crucial manner. In fact, they are in charge. The little guys control the bigger guys.

The Early Days of Computational Fluid Dynamics

Young man, in mathematics you don't understand things. You just get used to them.

JOHN VON NEUMANN (HUNGARIAN-AMERICAN
GENIUS MATHEMATICIAN)

Richardson was a very interesting and original character who seldom thought on the same lines as his contemporaries and often was not understood by them.

SIR GEOFFREY INGRAM TAYLOR (ENGLISH
MATHEMATICIAN AND PHYSICIST AND AN EXPERT IN
FLUID DYNAMICS)

Can com'uters be able to mult'ply: 999999999999999999999999
 9999999999 point 99999999999999999999999999999999
 99 by undredundredundredundredundred point 8889999
 99991212121414141414...

Yes! (A scientist intervenes)

Yo! Yo! I am not done. ...Undredundred and 9! Can we move on a little? When does technology goes horribly wrong?

ALI G. PLAYED BY SACHA BARON COHEN
 “INTERVIEWING” SCIENTISTS (ENGLISH ACTOR BEST
 KNOWN FOR PLAYING BORAT)

Typically, Computational Fluid Dynamics goes by the acronym of *CFD*. It stands for the process of creating the motion of a fluid on a computer. It involves computation. There are usually two components:

1. A representation of the fluid that our computer can understand.
2. A process to update this representation over time for each snapshot of the fluid.

The representation has to be *discrete*, unlike the mathematical models which are *continuous*. This means that the representation has to be described by a finite number of numbers. There are a lot of choices but the ones we will deal with here are either *grid based* or *particle based*. Either you keep your hand steady in the bathtub or you follow your rubber ducky floating along.

Once the representation is chosen, we need a discrete representation in time of the dynamical laws stated earlier to update the representation. I hope this clarifies the distinction between the continuous and discrete representations of fluids.

Where do we start with a brief account of computational fluid dynamics? Computing machines have been around since the time of Pascal and Leibniz. Did they compute the motion of fluids? As far as I can tell they didn't. Their early computer creations predated the Navier–Stokes equations.

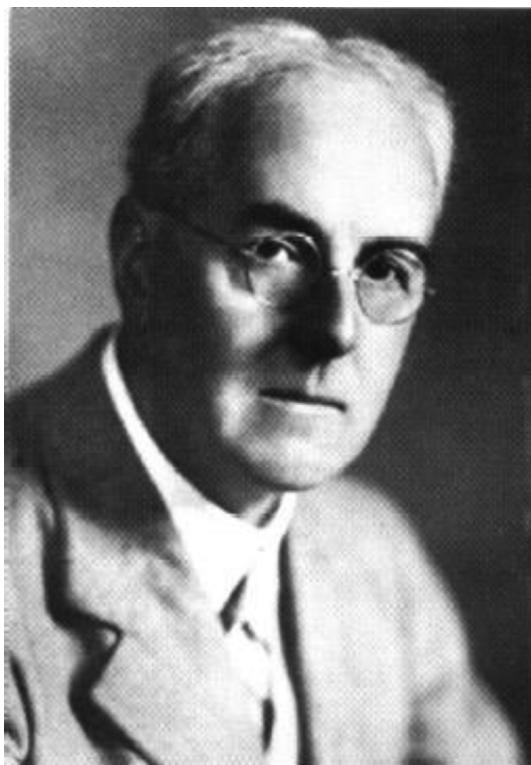


FIGURE 4.1 Lewis Fry Richardson.

The true pioneer in this area, I believe, is Lewis Fry Richardson (1881–1953),* he is portrayed in Figure 4.1. He was an English mathematician, a physicist, and a pacifist. But more importantly, he tried to put his talents into predicting the weather. This is easy in Britain: rain, rain, and more rain. In California it is sunshine, sunshine, and more sunshine. This is not strictly always true: sometimes there is sunshine in Britain and sometimes it rains in California.

Richardson wrote a treatise in 1922 called the *Weather Prediction by Numerical Process*. He was decades ahead of his time since computers as we know them today did not exist back then. No supercomputers, PCs, fancy laptops, iPads, or cloud-based computing.

A *computer* at the time was a person who would type in instructions in a calculator and report the results to another computer or themselves.

* There was however previous work published in 1904 by the Norwegian scientist Vilhelm Bjerknes (1862–1951) and the work by the Austrian scientist Felix Exner (1876–1930) in 1908. Their predictions were even worse than Richardson's. But they were all true pioneers.



FIGURE 4.2 A “computer” at work in 1922.

They were mostly women.* One of those *computers* is depicted in Figure 4.2. They were supposed to work in harmony. They would compute something and then pass it to another *computer*. This process would go back and forth. Obviously, this was quite a tedious and time-consuming task. A computer as we call it today is basically a dumb piece of hardware that does the same task, just way faster. But back then a computer was an aggregate of smart hardworking women cooperating together.

Richardson invented a representation of the fluid flow for Europe that is depicted in Figure 4.3. It is a coarse grid where every square is supposed to represent the velocity and pressure of the fluid. This is essentially a grid-based approach. Each cell corresponds to a woman-computer sitting in a square and passing notes to her four neighbors after doing a computation. This process would go on and on: iteratively. The prediction of this computation concerning weather prediction, by the way, was completely wrong.

This is what Richardson had to say about his first computation when he tried it in his office without the help of computers. Talk about dedication:

It took me the best part of six weeks to draw up the computing forms and to work out the new distribution in two vertical

* I know calling a woman a computer is degrading. Computers as we know them today are dumb, while women clearly are not. This terminology is wrong but that is what these smart women were called back then.

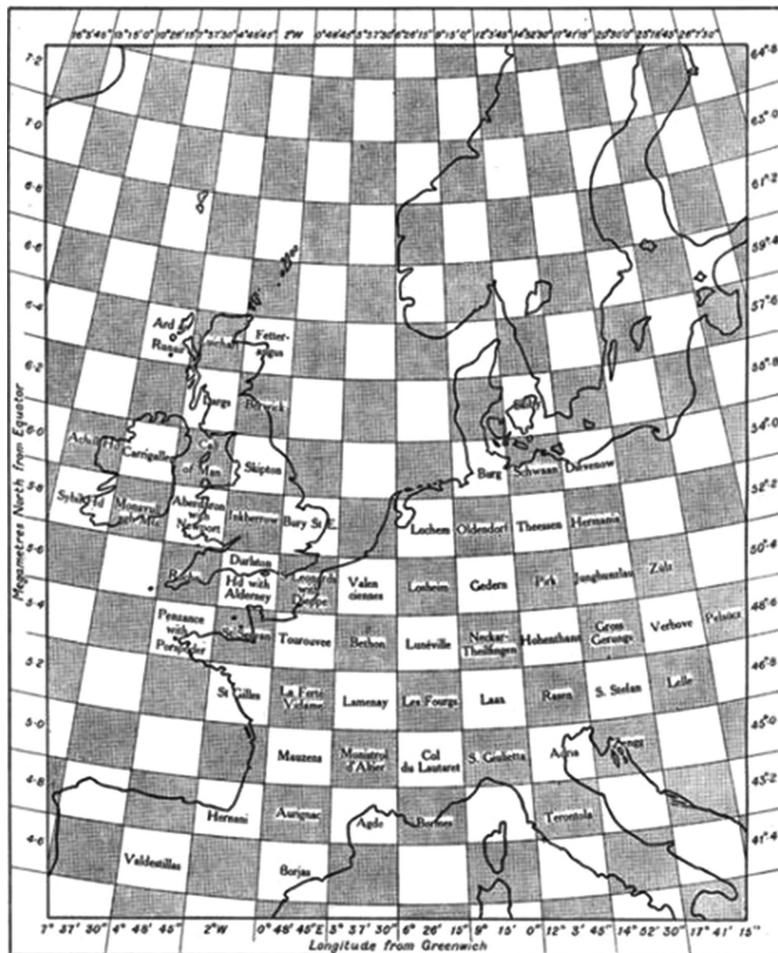


FIGURE 4.3 Richardson's first discretization of a weather prediction for Western Europe.

columns for the first time. My office was a heap of hay in a cold rest billet. With practice the work of an average computer might go perhaps ten times faster. If the time-step were 3 hours, then 32 individuals could just compute two points so as to keep pace with the weather.

Richardson was actually even more ambitious than that. He wanted to model the fluid circulation of the entire earth. Just like what global weather scientists are trying to do today. This is how he described his *fantasy*. The quote is a bit long but well worth reading in my opinion. Remember that this was in the 1920s, way before electronic computers existed. Richardson was a true visionary.

After so much hard reasoning, may one play with a fantasy? Imagine a large hall like a theatre, except that the circles and galleries go right round through the space usually occupied by the stage. The walls of this chamber are painted to form a map of the globe.

The ceiling represents the north Polar Regions; England is in the gallery, the tropics in the upper circle, Australia on the dress circle, and the Antarctic in the pit.

A myriad of computers are at work upon the weather of the part of the map where each sits, but each computer attends only to one equation or part of an equation. The work of each region is co-ordinated by an official of higher rank. Numerous little *night signs* display the instantaneous values so that neighboring computers can read them. Each number is thus displayed in three adjacent zones so as to maintain communication to the North and South on the map.

From the floor of the pit a tall pillar rises to half the height of the hall. It carries a large pulpit on its top. In this sits the man in charge of the whole theatre; he is surrounded by several assistants and messengers. One of his duties is to maintain a uniform speed of progress in all parts of the globe. In this respect he is like the conductor of an orchestra in which the instruments are slide-rules and calculating machines. But instead of waving a baton he turns a beam of rosy light upon any region that is running ahead of the rest, and a beam of blue light upon those who are behindhand.

Four senior clerks in the central pulpit are collecting the future weather as fast as it is being computed, and dispatching it by pneumatic carrier to a quiet room. There it will be coded and telephoned to the radio transmitting station. Messengers carry piles of used computing forms down to a storehouse in the cellar.

In a neighboring building there is a research department, where they invent improvements.

But there is much experimenting on a small scale before any change is made in the complex routine of the computing theatre. In a basement an enthusiast is observing eddies in the liquid lining of a huge spinning bowl, but so far the arithmetic proves the better way. In another building are all the usual financial, correspondence and administrative offices. Outside are playing fields, houses, mountains, and lakes, for it was thought that those who compute the weather should breathe of it freely.

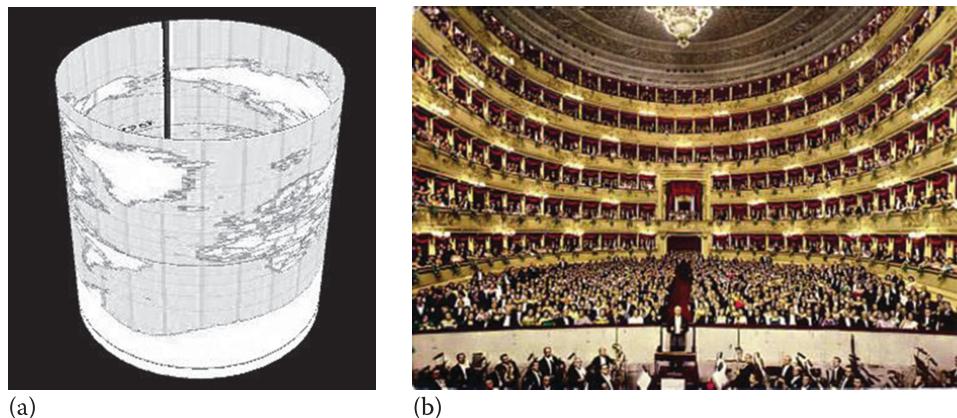


FIGURE 4.4 Richardson's fantasy. (a: From author's own image. b: From Tripnvisit.blogspot.com.)

Figure 4.4 shows the setup. It is basically a discrete cylindrical representation of the earth as depicted on Figure 4.4a. To implement Richardson's fantasy, one could use the venue of *La Scala* in Milan, Italy, depicted on Figure 4.4b. Usually, it is a venue for concerts and operas. But Richardson could have rented it out, placed a person-computer in each booth and had them compute. They would exchange their computations with the computers in the booths above, below, left, and right. *La Scala* is a fancy place really but it is just one-half of a cylinder. Oops. But through a clever mechanism one could in principle carry out these computations for the entire cylinder world. Two computers could be working side by side in each booth, for example. Or one computer could multitask.

Another option for Richardson would have been to rent the legendary *Wembley Stadium* in London, England, shown in Figure 4.5.* The stadium was inaugurated in 1923 and could seat 127,000 people. Of course it wouldn't be as classy and comfortable as *La Scala*. There would also be the challenge for Richardson to find 127,000 person-computers to do the task. Doing fluid computations in Wembley Stadium would roughly result in a 1270 by 100 2D cylindrical grid. Not bad actually even by current standards. If only enough computer-persons would show up to achieve this task.

Richardson is also the man who invented a technique to accelerate the convergence of series in a very clever way. It is a technique to get a computation from previous computations in a much faster manner. Not surprisingly,

* In 1996, I was lucky enough to see a soccer game live at Wembley Stadium (the new one) between England and Switzerland in the European Championship. It turned out to be a tie: 1–1. It was also the opening game of the tournament. Germany ended up winning the final. I happened to be in Hamburg on my way to Finland from Paris when the Germans won. Lots of beer and fireworks!



FIGURE 4.5 The old Wembley Stadium that opened in 1923.

it is called the *Richardson extrapolation*. We do not use this technique in our code so we won't elaborate on it. But it is worth looking up.

Then there was World War II. This catastrophic and tragic event, however, accelerated many technological advances in many areas such as aviation, explosives, and atom bombs but also computing. Examples include Konrad Zuse in Germany, Hewlett and Packard in the United States, the *Bomb* based on the work of Alan Turing,* the *Harvard Mark 1* created by Howard Aiken, and many more. This was a fertile time for the development of machines that compute instead of people-computers. Secret codes had to be cracked and ballistic trajectories had to be predicted in order to win the war. Computers were an essential part of the solution.

John Von Neumann was I think one of the biggest pioneers in computer theory and practice. He laid out the architecture of all computers we use today: the *Von Neumann machine*. A machine that can read in programs and is not hardwired for a specific problem. Turing of course envisioned this before. But in 1945, Von Neumann was involved in a project to build a better ENIAC computer, which stands for Electrical Numerical Integrator And Calculation. In Figure 4.6, John Von Neumann is shown next to the ENIAC computer. Many of the first programmers were women who were probably happy not to be called computers anymore. In Figure 4.7, we show two new programmers.

The ENIAC was one of the first computers to perform computational fluid dynamics, and it was based on Richardson's work from the 1920s.

* For a good account of Turing's contributions during the war and his other achievements read *Alan Turing: The Enigma* by Andrew Hodges. Sadly, Turing committed suicide because his homosexuality was viewed as a crime in Britain at the time. Such a waste. Here is a genius mathematician and a war "hero" who made major contributions to win a ridiculous war.

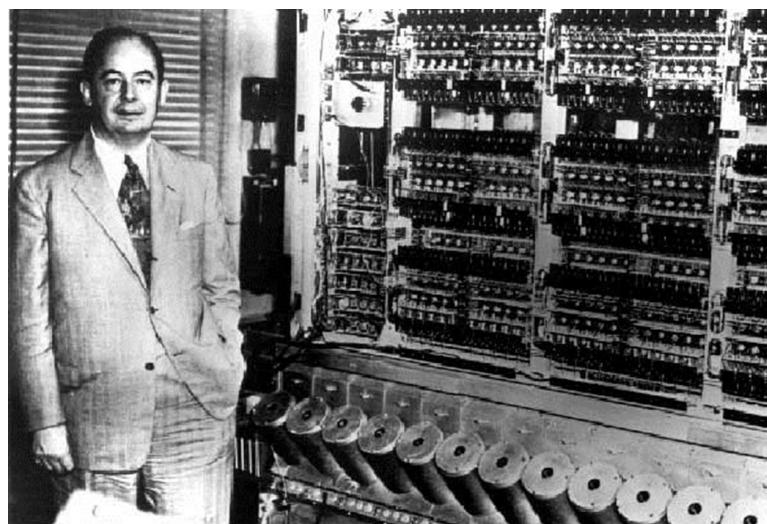


FIGURE 4.6 John Von Neumann (left) and part of the ENIAC computer (right).

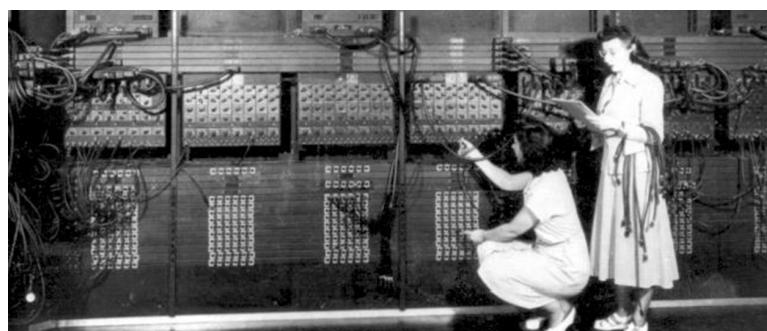


FIGURE 4.7 Computers became programmers programming a computer.

This time the computations were faster but still extremely slow compared to our computers today. The ENIAC could roughly perform 100,000 operations per second. That was huge back then and required a lot of power. The ENIAC was located in Philadelphia and legend has it that when the ENIAC was powered on all the lights in the city slightly dimmed. Nowadays, any average personal computer can operate roughly 1,000,000,000 instructions per second and if you boot up your PC no one will notice. Not even your neighbor.

Apart from his fundamental contributions to computer science, Von Neumann was a genius with a photographic memory. He advanced many fields including the mathematical foundations of *Quantum Mechanics*.

Most people know the town of Los Alamos in New Mexico. That is where the foremost scientists in the United States during World War II developed the Atom bomb.

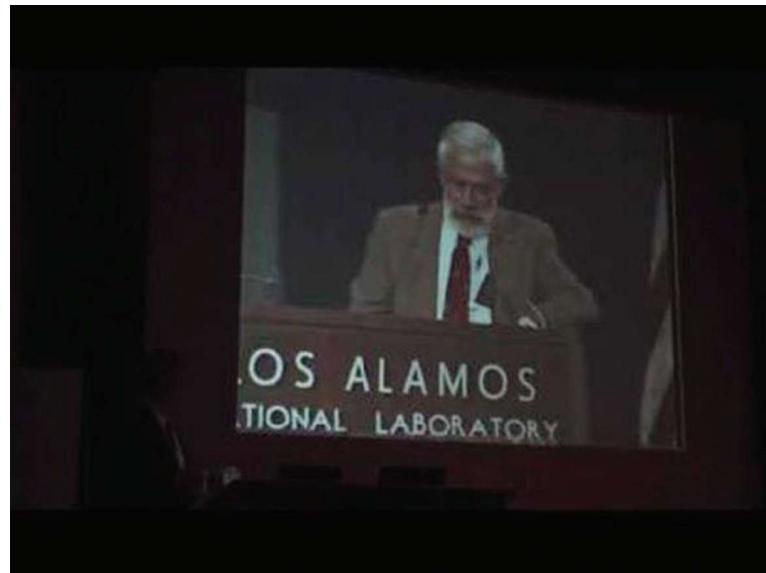


FIGURE 4.8 Frank Harlow.

After the war, Los Alamos Laboratories was a hotbed for research using mechanical computers to simulate fluids. One of the main characters involved in this research was Francis (Frank) Harley Harlow. Figure 4.8 shows a picture of him.* He devised, with the help of colleagues, many clever techniques to simulate fluids on a computer. Those were the *glory days* of computational fluid dynamics. He picked a lot of low hanging fruits. His creations are essential to how we model fluids numerically today.

After this period, there were of course tremendous improvements in computer power and software to improve the simulation of fluids. We will not delve into those improvements since they are well documented elsewhere.

To summarize: Richardson was a true pioneer in weather prediction, even though his predictions were way off. But he laid the foundations for computational fluid dynamics. Women had a crucial role in the early days of weather prediction: first as computers then as programmers of computers.

Let us briefly mention the topic of *turbulence*. This is a somewhat vague and controversial topic. That is what makes it interesting and cool.

* Sadly, I couldn't find a better picture of Harlow. This one I found on Google images extracted from a YouTube video.

Kolmogorov and Turbulence

... the smallest eddies are almost numberless, and large things are rotated only by large eddies and not by small ones, and small things are turned by small eddies and large.

LEONARDO DA VINCI (ITALIAN RENAISSANCE MAN)

Big whorls have little whorls,
which feed on their velocity;
And little whorls have lesser whorls,
And so on to viscosity.

LEWIS FRY RICHARDSON (ENGLISH MATHEMATICIAN,
PHYSICIST, AND PACIFIST)

When I meet God, I am going to ask him two questions: Why relativity? And why turbulence? I really believe he will have an answer for the first.

WERNER HEISENBERG (FAMOUS NOBEL LAUREATE,
GERMAN PHYSICIST)

NOT THE DUDE FROM *BREAKING BAD*

What is turbulence? No one really knows. It is a mess. But it involves fluids. For most people turbulence is often associated with air travel, fastening seat belts, being shaken around unpleasantly, and having coffee spilled all over your pants.

There are plenty of models of turbulence for fluids.

If you are in this turbulence business, you will be in this business for a long time. That is as long as no experiment proves that your business is wrong. But then you will just change your business model. And you are back in business.

Really, there is no proper understanding of turbulence as of yet. We already mentioned earlier that the dynamics of fluids are nonlinear. That means that large-scale features influence small-scale features. But also the other way around: smaller scales influence larger scales.

The so-called Arab Spring event is a good analogy of this nonlinear turbulence behavior.

Totalitarian regimes (large scale) controlled the masses (small scale) but then a small tragic event in Tunisia (small scale) ripples through the Arab world and challenges these totalitarian regimes (small scales influence large scales). The totalitarian regimes (large scale) step down because of revolts (small scales becoming large scales) and subsequently elections (an aggregate of small scales) bring in a new party (large scale), which is then brought down by the military (medium scales influencing large scales). This is an ongoing saga as of the writing of this book: large scale versus small scale and small scale versus large scale.

Bottom line: turbulence is a mix or a battle between scales.

My description of the Arab Spring is just an analogy of course and an oversimplification of the real situation.

We still haven't precisely defined what turbulence is.

One pathetic attempt at a definition could be "turbulence is the small stuff we cannot model at large scales." It is the complicated stuff that happens at the smaller scales of fluids. On the other hand, these smaller scales influence the larger scales. That is why turbulence is a hard nut to crack.

One of the most famous models of turbulence was created by the great Russian mathematician Andrey Kolmogorov (1903–1987). Figure 5.1 shows a picture of Kolmogorov in action explaining some fundamental concepts using sketches and pictures. Recall the quote mentioned above in the Preface of this book. To humbly paraphrase: well understood mathematics is best explained with no fancy mathematics at all.



FIGURE 5.1 Kolmogorov in action. Cool, just diagrams.

His work on turbulence is called the *K41 Theory*. This is because he published three influential papers on turbulence in 1941 and the “K” stands for Kolmogorov of course. By far, this work is not his only contribution to science and mathematics.

For example, he defined complexity as follows. The complexity of a problem is equal to the number of bits of the smallest program that encodes the solution to the problem. Got it? The smaller the program that solves the problem is, the less complex the problem is. I love it because it makes it so concrete.

A program that prints the number 1 a trillion times is less complex than printing a trillion digits of the number π .*

Compare the size of a program that prints:

* The record for memorizing π was achieved by Akira Haraguchi who recited π from memory to 100,000 decimal places in 16 hours on October 3, 2006. Each digit requires about 4 bits. So, the total amount of bits is $4 \times 67,890 = 271,560$. The length of his program is what he memorized. That is a lot of bits. For most people π is 3.14. Indiana House Bill #246 introduced on January 18, 1897 claimed that $\pi = 3.2$. A small complexity declared by fiat. Of course computer programs that compute the digits of π do not do it from something stored in memory. There are many clever ways of generating these digits. As of the writing of this book, the record is held by a program called "y-cruncher." In 2013, it computed 12,100,000,000,050 digits of π in 94 days.

To a program that prints:

3.14159265358979323846264338327950288419716939937510582097494
4592307816406286...

A truly random sequence cannot be computed by a program smaller than itself.

What was Kolmogorov's contribution to turbulence?

The starting point of Kolmogorov's theory of turbulence is Richardson's poem quoted at the beginning of this chapter. Kolmogorov assumed that there is a cascade of energy transfer from large structures to smaller structures and that energy is eventually dissipated at the smallest scale because of viscosity. This setup assumes no boundary conditions at all. Figure 5.2 is a cartoon depiction of this scenario. The circles indicate the size of a *structure*. Think of these structures as vortices or any discernible pattern in the flow. This is somewhat sketchy but there is really no precise definition of turbulence. Structures of different sizes are clearly visible in Da Vinci's drawing (see Figure 3.2b). The wiggly line in our cartoon figure is the scale where the features "pop" and disappear because of viscous dissipation.

This picture and the theory is really an approximation of what really goes on. But abstracting complex problems like this can lead to useful results and insights. Kolmogorov's theory assumes that the energy of the fluid is transferred from large scales to smaller scales. As we mentioned earlier, this is not realistic because energy also gets transferred from small scales to large scales: a fluid is nonlinear after all.

In Figure 5.2 we show two arrows: one on the left and one on the right. Why? The right one is obvious and shows that as you go upward the circles

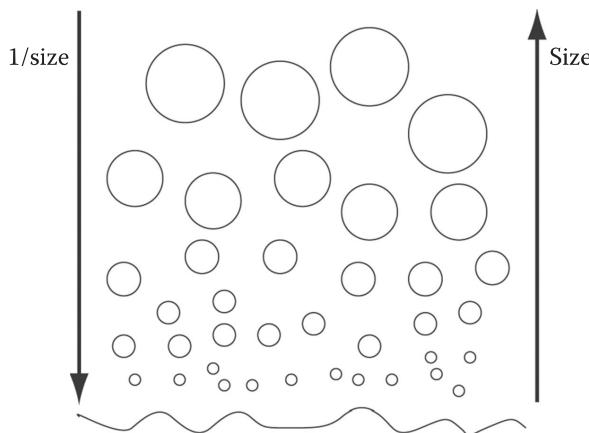


FIGURE 5.2 A cartoon of Kolmogorov's cascade. These are not bubbles!

get bigger. Their size increases. Kolmogorov uses the left-hand arrow. This arrow goes in the opposite direction and shows the increases of what is called the *wave number* of a structure.* We will see more about these wave numbers when we discuss about the Fourier transform, but really in our figure it just stands for “one divided by the length.” Small structures have larger wave numbers than larger structures which have smaller wave numbers. Confusing? A big number is big, but if you divide a small number like 1 by a big number then the result is tiny. Compare 1,000,000 with $1/1,000,000 = 0.000001$.

Kolmogorov first assumes that the fluid has constant density, and therefore we can ignore mass. This means that energy has units of $[L^2/T^2]$. Kolmogorov derived an equation of how the energy depends on the size of these whimsical structures.

Actually, Kolmogorov derived an equation not for the energy but for the *energy per wave number*. This is the energy for each structure in the flow and we denote it by E . You can think of E as a density of the total energy for each structure of a specific size. The wave number is usually denoted by k and it is related to 1/size.

Kolmogorov postulated that the energy density E should only depend on some nondimensional constant C ,[†] the rate of flow of energy which we call *flow*, and on the wave number k . This density of energy per wave number has units of $[L^3/T^2]$. Kolmogorov used dimensional analysis to come up with his formula. We mentioned this technique above to derive Bernoulli’s law. The *flow* variable is energy per unit time. Therefore *flow* has units of $[L^2/T^3]$. A first attempt at a law could be

$$E(k) = C \times \text{flow} \times k.$$

What about the units? Do they match? Let’s try it out: $[E] = [C] \times [\text{flow}] \times [k]$. Thus $[L^3/T^2]$ has to be equal to $[L/T^3]$. That unfortunately doesn’t work.

Here enters the *coup de génie* of Kolmogorov.

He introduced two exponents a and b to be determined such that

$$E(k) = C \times \text{flow}^a \times k^b.$$

* Hopefully, I will make it clear why they are called wave numbers instead of *reciprocal sizes* or *one over size* later in this book.

[†] This constant can be determined through experiments. It is just a number that depends on the fluid.

Now we have an equation for the exponents a and b using dimensional analysis. $[L^3/T^2] = [L^{2a}/T^{3a}] \times [L^{-b}]$. This gives us two equations for a and b since the exponents have to match on both sides: $3 = 2a - b$ and also $2 = 3a$. And the only solution is that a must be $2/3$ and that b must be $-5/3$.^{*} And *ta da!* From this follows Kolmogorov's law:

$$E(k) = C \times \text{flow}^{2/3} \times k^{-(5/3)}.$$

This is also known as Kolmogorov's "5/3 Law." Given the assumptions stated earlier to this argument holds for any fluid.

That is the beauty of it.

Of course this is a far cry from settling the problem of turbulence. The reason I mention this particular result of Kolmogorov's is that this formula and its derivation are relatively simple. No need for any fancy mathematics.

More importantly, we use this expression when we introduce numerical models for turbulence in computer animation.

To summarize: Turbulence is far from being understood. But Kolmogorov gave us a simple formula that characterizes the cascade of energy from large scales to small scales. It is a power law. Like *fractals!*

* For those of you not familiar with exponents we have the following facts. Let X be any number then we have that: $1/X^a = X^{-a}$ and $X^a \times X^b = X^{a+b}$. For example: $X^2 \times X^3 = X \times X \times X \times X \times X = X^5$. That is why exponents are awesome: they turn multiplication into addition. Unless you are a prodigy, I think most people would rather add numbers than multiply them in their head.

Introduction to Fluid Animation

If it sounds good, it is good.

DUKE ELLINGTON (FAMOUS AMERICAN JAZZ
MUSICIAN)

We now focus on the problem of animating fluids. This is a research area of its own that of course relies heavily on all the previous fluid work described earlier.

Research is a process that constantly evolves.

Generally in computer animation the motto is:

If it looks good, it is good.*

That is the bottom line. It might seem simplistic and unscientific at first glance.

Anything goes really, but not really. It is hard to make things look good. In that sense it is unscientific. But that doesn't make it less interesting. Making things that look good with computers involves using both principled models and using artistic intuition.

* I only heard of Duke Ellington's quote from a fellow researcher in a restaurant in Fukuoka, Japan in 2012.

My approach in fluid animation research has always been to use any technique that creates good-looking fluids. I am a model agnostic. Whatever model gets the job done is a friend of mine. Quite often I meet people who equate models to reality (without ironic quotation marks). To them any model for fluids *has* to use the Navier–Stokes equations. Everything else is garbage, whatever. But with this kind of attitude you are going to miss out on a lot of fun and cool stuff.

Besides, computer animation is not about *reproducing* reality but about the *creation* of an imaginary and controllable virtual reality *inspired* in part by conventional physical models of *reality*. I do not own an apology to anyone. I do what I do.

Computational fluid dynamics (CFD) is of course a great tool. It has derivations that are principled, it has fancy equations, and it has a lot of successes that are backed up by rigorous experiments. I love CFD, and I am using some of their concoctions. Mathematical models, once you understand them and tried them, can be highly addictive. But we should not be carried away by their allure.

In fluid animation we can augment and take liberties to improve the visual quality of the fluids. Just taking existing respectable fluid codes and visualizing their results doesn't always do the trick in fluid animations. In fact a lot of CFD papers do not depict fluids at all. In some cases, they are just 2D plots of some quantity. Da Vinci would not be pleased.

To summarize: Enough said. Let's turn to the business of writing code that simulates fluids for computer animation.

6.1 DISCRETIZE! BUGS, GRIDS, AND BUGS MOVING THROUGH GRIDS

It is better to destroy one's own errors than those of others.

DEMOCRITUS (GREEK PHILOSOPHER KNOWN FOR HIS
THEORY OF DISCRETE MATTER)

Before we write code we have to decide how to turn our perfect continuous fluids into a discrete representation that our computers can handle. Remember everything has to be crammed into a finite number of bits that the computer can process. The continuous models are just fictions made up by mathematicians. But, and this is a big but, these models are extremely useful to writing code. How cynical is that?

6.1.1 Bugs

There are, as is known, insects that die in the moment of fertilization. So it is with all joy: life's highest, most splendid moment of enjoyment is accompanied by death.

SOREN KIERKEGAARD (DANISH PHILOSOPHER AND THEOLOGIAN)

One model is to use bugs running around. However, this time they just cannot crawl around just as they please. They have to cooperate: no climbing over each other. They are more social than our diffusion bugs mentioned earlier. Ideally, they have to remain at a certain distance from each other. This distance is usually fixed by the programmer.* If they get too close, a *spring* is attached to them temporarily so they don't bang into each other. The density of the bugs doing their wild dance is then approximately constant. And their velocities are approximately incompressible. Remember: flow in equals flow out. Figure 6.1 shows this situation. On the left the bugs are free to move but on the right they are connected with springs to keep them apart. These springs come and go.

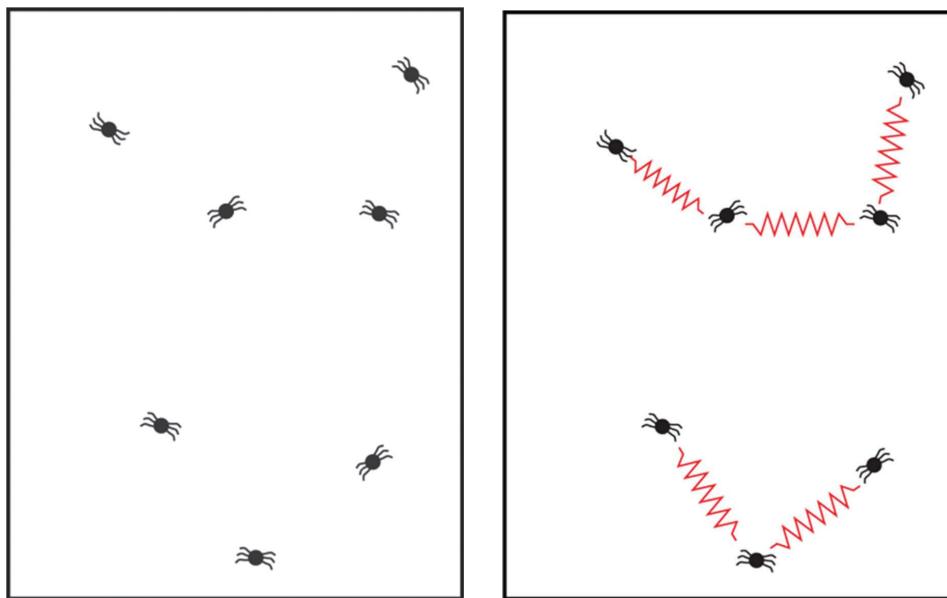


FIGURE 6.1 Free bugs (left) and constrained bugs with *springs* attached (right).

* In the case of humans, this distance varies from culture to culture and from country to country. In my experience, in most of Europe that distance is smaller than in North America.

This is not what these methods are normally called. They often go by many names like *griddles*, *particle based*, *smoothed particle hydrodynamics*, and so on. A bug is usually called a particle in these models. There are various models for the *springs*.

A famous model for the *spring* is the so-called *Lennard-Jones force* between particles. I use the word particles because this force was originally designed for atoms and molecules, not bugs. Figure 6.2 shows the strength of the Lennard-Jones force as a function of the distance between particles. For small distances the force is highly repulsive. Actually, the force is equal to infinity when the two particles share the same location. This is problematic when you try to code this model. Computers do not like infinities. Remember they can only handle a finite number of bits. The force is zero for a certain given distance r_0 . This is the *comfort zone*. For distances larger than r_0 the force is attractive but much weaker than the repulsive force.

Another popular model is the *hard sphere model* for fluids. In this case, each bug is put into a sphere that protects them from other bugs. Their spheres collide when they get too close. So, it is like a wild game of billiards with a bug in each ball. This situation is depicted in Figure 6.3. One can also replace the bugs with hamsters. Anyone who owned a hamster knows that you can buy those transparent balls that allow them to roam through your house. Do this on the first floor. You do not want your hamster

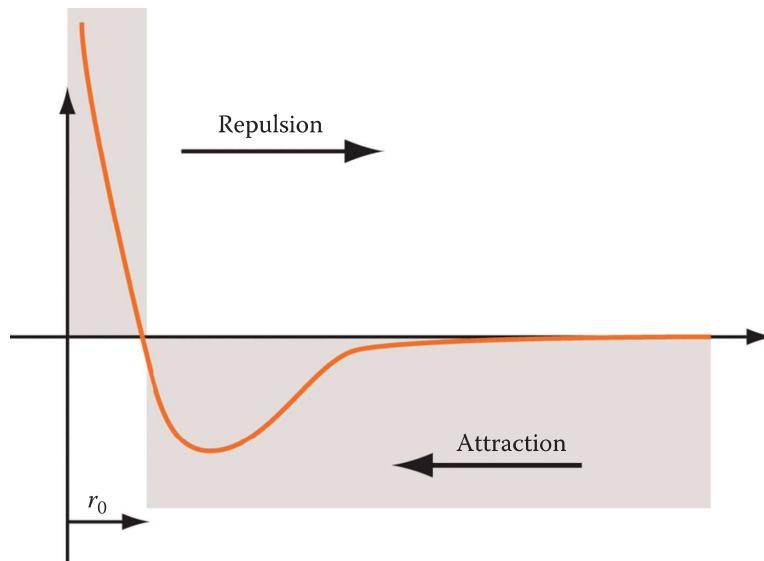


FIGURE 6.2 The Lennard-Jones force between two particles as a function of distance.

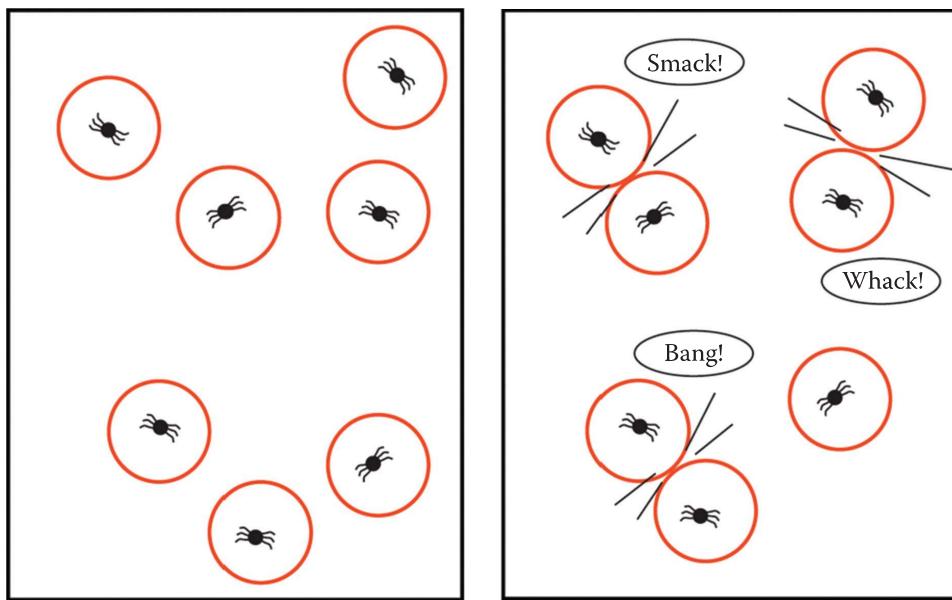


FIGURE 6.3 The hard sphere model of a fluid.

tumbling down the stairs. Now, imagine 1,000,000,000,000 hamsters in their mighty spheres all banging into each other. That would be kind of a realistic implementation of a hard sphere fluid model.

Alternatively, there is a model called *SPH*, which stands for *smoothed particle hydrodynamics*. It is quite popular these days in computer animation and was originally developed to model the dynamics of galaxies. These models try to explain how galaxies form, collapse, merge, or disappear in a black hole. This means the model is perfect for fluids that compress and diverge.

Briefly, the way SPH works is as follows:

Unlike our hamsters, each bug now has an *aura* around them that determines the interactions with other bugs. This is depicted in Figure 6.4. The poor bugs are now in the dark driven only by their aura. Once the auras of the bugs overlap they interact. That is the smoothed part in SPH. Bugs are not trapped in hard spheres but have this glowing aura around them. Kind of like a new age energy thing. In computer graphics, the official nomenclature is that every bug is surrounded by a *blob*.^{*} The blob is strongest at the bug and then it drops off to zero after a fixed distance. If you add up the blobs then you get a fuzzy function everywhere: a global aura of energy. The velocity of each bug is determined by its own velocity

* Or if you are in Japan, *blobs* are called *meta-balls*. Sometimes they have been misspelled in papers as *meat-balls*.

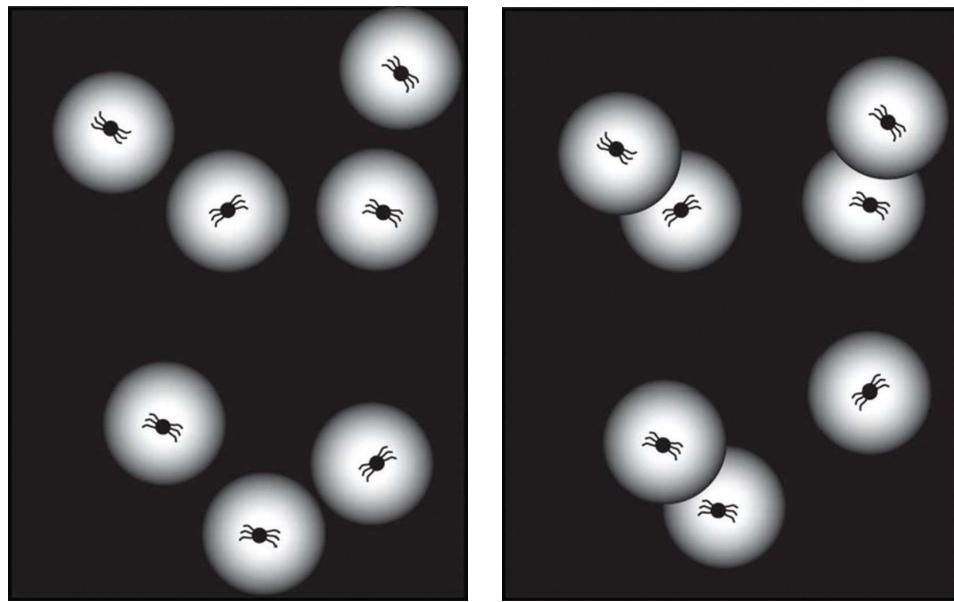


FIGURE 6.4 In the SPH method each bug is surrounded by a blob. Interactions occur when blobs overlap.

and the velocities of the neighboring new age bugs whose auras overlap. Once you have the velocity of the bug it now knows where to move next: It is the same for the other new-ager bugs. Through their interactions, an emergent fluidlike behavior emerges.

That is SPH in a nutshell.

What about incompressibility?

That is the *Achilles heel* of SPH.

However, a lot of people have extended SPH to incompressible fluids like air and water. The basic approach is to have repulsive/attractive forces between the particles when their auras overlap. If incompressibility is to be enforced, very stiff spring forces have to be introduced, almost like a rigid link.

In practice, this means that you have to be careful about the *time step*: the amount of time separating two snapshots of the roaming bugs in the dark. If you are not careful the bugs will *blow up*. They will wander all over the place as far as they can from each other. This creates infinite values.

This is not cool for the computer Slave since it can only handle a finite number of bits. So, the way the Slave handles these cases is to just give up. “Geez these numbers are too big, take it easy Master. Instead of your big numbers I will call them *NaNs* and you deal with them. Even better if you

fix up your code so they do not show up at all.” A number called an *NaN*, which stands for Not a Number is the kiss of death. Once a *NaN* appears when running your code it will slow it down and do all kinds of funny stuff. But there are ways to catch *Nan*s and make sure they do not happen.

Our users do not like *Nan*s. Mathematicians call this an *instability*. Coders call this “useless garbage and go fix your code dude.” Coders have an affinity to the Slave who has to digest their code.

An interesting side effect of SPH is that fluids can be bouncy like Jell-O. But less damped and more lively. This is not very realistic but this effect can be cool in fluid animation. The biggest issue with a simple vanilla implementation of SPH is that it can be unstable. But a lot of researchers have tackled this problem.*

This book will not include SPH code. Sorry. You can find the code on the Internet, however. Or you can write your own. Have fun.

I just briefly described a few so-called Lagrangian approaches to fluid motion. Remember you can measure the change by following the flow. Go with the flow.

To summarize: The interaction of the bugs determines a velocity for each bug and their crawling follows that velocity field. At the same time, these bugs have to get along: not too close and not too far.

Now let’s turn our attention to grids.

6.1.2 Grids

The best advice I’ve received is to be yourself. The best artists do that.

FRANK GEHRY (FAMOUS TORONTO-BORN ARCHITECT)

Grids are just a way to partition space into cells. Think of grids as one of those hideous apartment buildings with identical tiny apartments. Usually, they are built to house the poor and newly arrived immigrants. In French, they are called *cages à lapin* (rabbit cages). Now each bug is trapped in its tiny apartment, and it can only interact with its neighboring bugs through its six walls. Yes we also include the floor and the ceiling. This is depicted in Figure 6.5

* I implemented the standard SPH in 1991 and it was incredibly unstable. I gave up. I like big time steps. Kudos to the people who made it work. Back then I asked J.J. Monaghan from Australia who co-invented the technique to send me some of his papers. He sent them to “Jos Stam, Toronto, Canada.” Amazingly his papers got to me. Thanks Monaghan! I sincerely mean it. Toronto is the third largest city in North America for those of you who do not know this fact. We have more humans than wolves and bears in this fine city even though it is located in funky Canada.

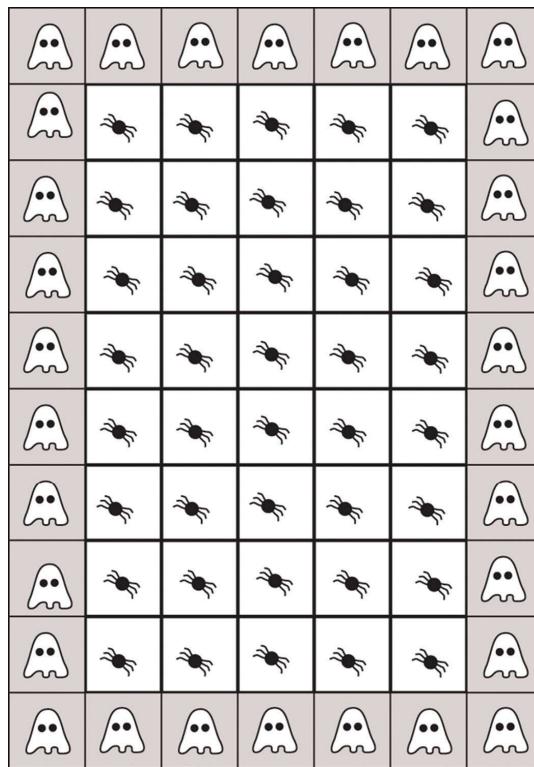


FIGURE 6.5 Bugs stuck in their apartments surrounded by ghosts

for bugs trapped in a 2D apartment building. In this case, the bugs can only interact with their four neighbors.

They interact exactly like the women computers mentioned earlier. Let's say each wall, ceiling, and floor has a tiny hole so the bugs can pass messages back and forth. For example, "My velocity arrow is such and such what is yours?" No privacy for these bugs. Not true for the lonely bugs that live in "internal boundary apartments." They can sleep all day and night.

We also included a layer of *ghost apartments*. They do not really exist. But we added them so the bugs with an actual view also have four neighbors. The ghost apartments with their ghost bugs will not block their views.* But it will make it easier to handle what happens at the boundaries. The ghost in an adjacent apartment is sometimes like a mirror image of the bug. It does everything the bug does but in reverse. Actually, it is kind of freaky to have a ghost neighbor like that. But this is not always the case. The ghost can act in different ways. Another example is that the ghost neighbor could be blowing air in the bug's face all day. Or the ghost can grab the poor bug and throw him out of the window.

* What about the ghosts? Well they do not exist anyway, so who cares.

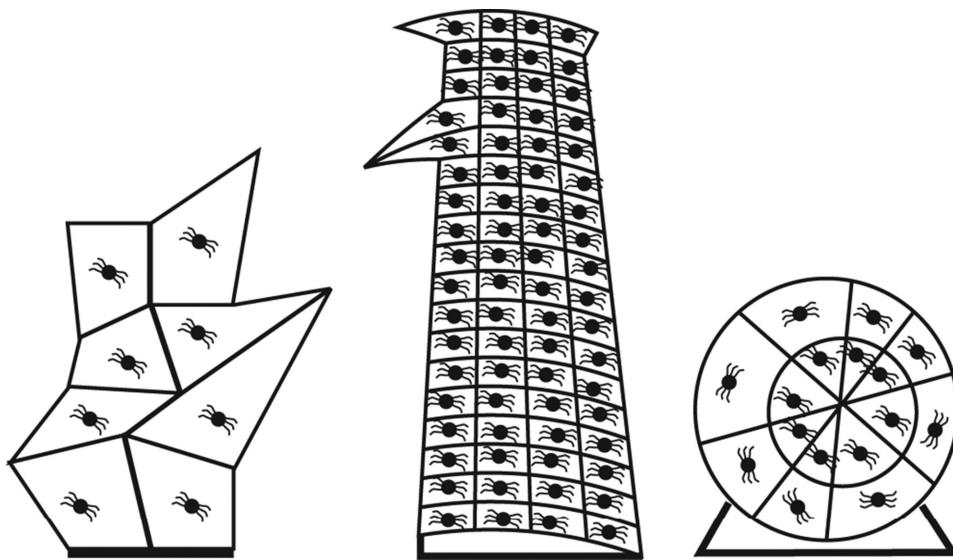


FIGURE 6.6 Three funky hipster bug apartment buildings.

Grids do not have to be as regular as this. We can imagine funkier apartments for the bugs. Three examples are depicted in Figure 6.6. These apartments are for hipster bugs and are most likely built by famous architects like Gaudi or Libeskind. We could have come up with even crazier grids. But I think I made my point. Notice that the grids can have curved walls *à la* Frank Gehry.

In a grid representation, the bug is actually not allowed to move at all. The bug is just sitting in the center of his apartment on a revolving stool.* So it can face any angle. The bug is also holding an arrow whose length can change over time. The evolution of the bug's arrow is a function of the arrows of their bug neighbors. This is shown in Figure 6.7. The top three bugs show the evolution of a single bug's orientation and arrow length over time. Notice that the color of the bug can also change over time as well. This could be indicative of the amount of heat in the apartment, the levels of carbon dioxide and what not.

In the bottom of the figure, we show how bugs exchange information with their neighbors through their walls.

This grid representation is called *cell centered* or *collocated* because both the bug's color and their velocity arrow are defined at the center of the cell. This representation makes the computer implementation easier

* In three dimensions, the stool would have to be more complicated. Like something you can find at an amusement park.

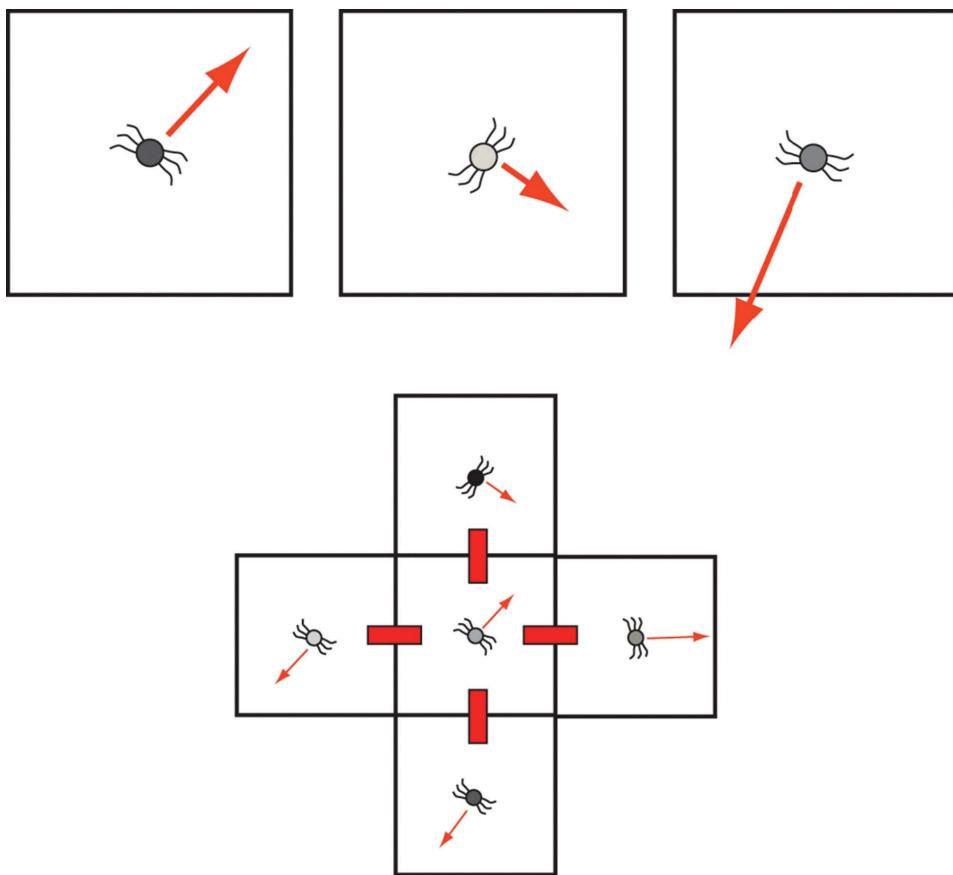


FIGURE 6.7 Evolution of a bug's arrow and color (top). Interactions between neighbors through the walls.

and simpler. Indeed, we will give a full implementation in code for the case of these poor bugs that are stuck in the *cages à lapin* surrounded by ghosts.

Another grid representation has the lateral velocities defined at the wall boundaries. This representation is actually more natural and can be understood as fluxes defined at the walls of the bug's apartment. This representation is usually known as a *staggered grid* or a *MAC grid* in fluid dynamics speak. This is an acronym for Marker And Cell. MAC is normally the preferred representation, and we have used it to explain continuity and the incompressibility conditions earlier (see Figures 6.9 and 6.10). This is the representation that is most commonly used for grids. We used MAC in our computer implementation of the fluid animation feature called *Fluid Effects* in our MAYA animation software.

This representation is depicted in Figure 6.8, where each bug has a horizontal and a vertical velocity defined at the walls that it shares with its

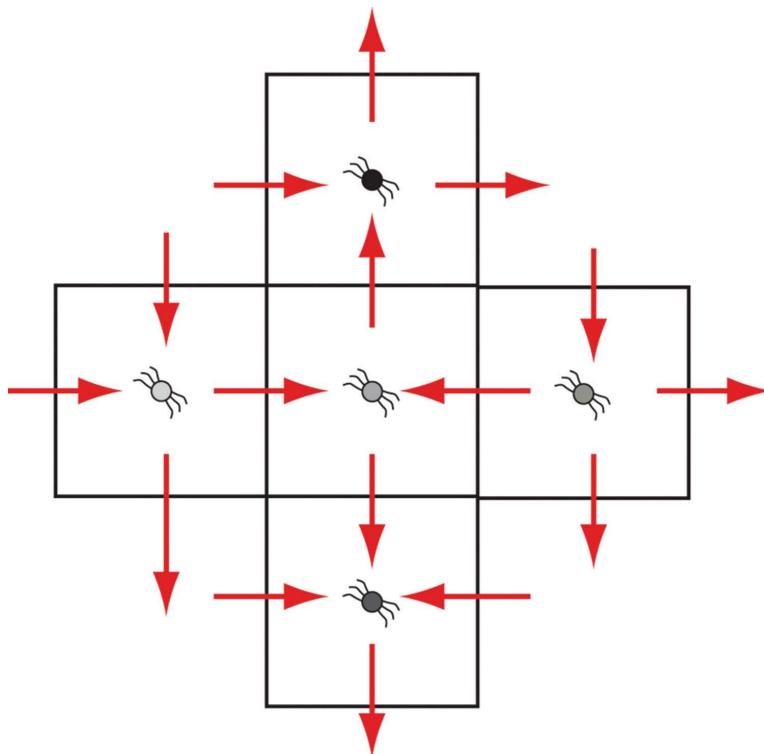


FIGURE 6.8 In a staggered grid, velocities are defined at the walls and other quantities are defined at the bugs.

neighbors. The bugs are being very neighborly. The staggered approach also extends to the more hipster apartments depicted in Figure 6.6 earlier, even though the walls have funkier shapes.

To summarize: In a grid representation, bugs are stuck in their apartments and can only exchange information with their neighbors. Properties of the fluid are only defined in each bug's apartment.

Now we will briefly discuss cases where bugs are allowed to leave their apartments. In some cases the bugs have to return home but in other cases the bugs can take over their neighbor's apartments.

6.1.3 Bugs Moving through Grids

Most people are aware of the particle/wave duality in quantum mechanics. Particles and waves are mathematical abstractions. Sometimes using particles makes sense and sometimes using waves makes sense. They are just models that complement each other.

But really when you think about it: what is a particle that has zero size and what is a wave made of? It is all mathematical abstract nonsense that happens to be useful in physics. Take a second to think about that.

To paraphrase John Von Neumann: no one understands what a particle or a wave is. We just get used to the math and then pretend that we understand it. But it works!

This analogy is just a dandy (pretentious?) way to illustrate that one can combine models that describe a phenomenon from different points of view. That is why I describe bugs crawling between apartment buildings. I can feel and see them and therefore I can describe them.

Earlier we talked about free-roaming polite bugs and poor neighborly bugs stuck in their apartments. Why not combine these two situations into one? That is when the bugs can move from one apartment to the other. Like some sort of commune of the 1960s in California or a kibbutz in Israel. I like this approach as it is social and model agnostic at the same time.

We will present two such approaches.

6.1.4 Semi-Lagrangian

I will start with my favorite technique which is officially called *semi-Lagrangian*: half of a French mathematician and philosopher. This method is also called the *method of characteristics*.

This is how the semi-Lagrangian technique works.

As shown in Figure 6.9, we have two identical apartment buildings. The one on the left is filled with bugs and the building on the right is initially empty. Let's say the landlord decided to relocate all the bugs because the

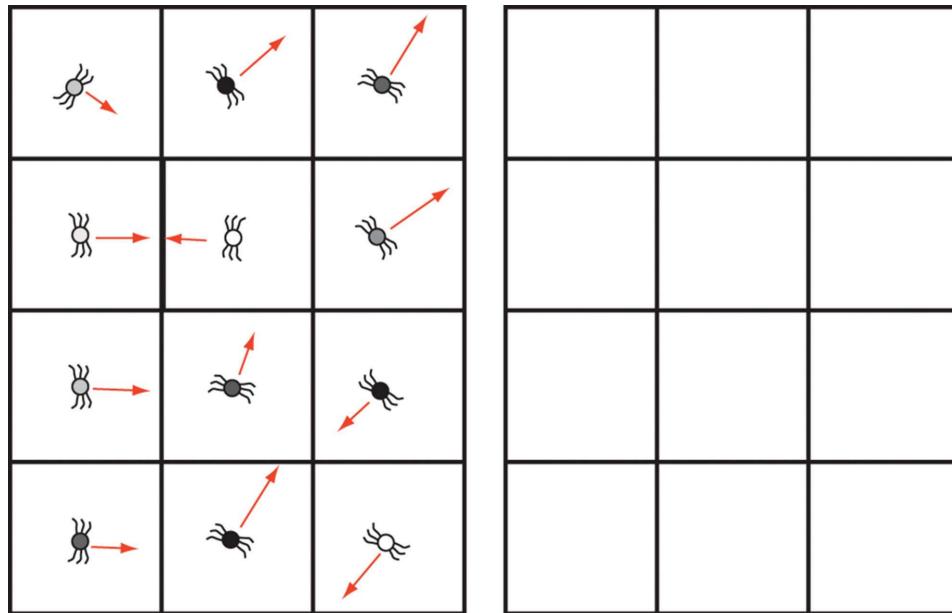


FIGURE 6.9 Semi-Lagrangian transport of bugs between two apartment buildings.

first building has an asbestos problem. The bugs have to eventually end up in their matching apartment in the second building, which is asbestos free. The landlord is clearly a bit of a psychopath and a control freak. But he offers a lower rent in the newly cleanly built apartment building if the bugs follow his request.

The landlord's crazy request is as follows.

You guys are now free to leave your stool. And roam around. One condition, however: follow your arrow multiplied by some number in the opposite direction. You will end up in another apartment. Now mix the colors of the bugs in the neighboring apartments and recolor yourself. Also compute your new arrow from the neighbor's arrows. And *voilà!* Welcome to your new apartment that is asbestos free and cheaper.

This is a transport mechanism.

Colors and vectors are transported from the *old* apartments to the *new* apartments.

Figure 6.10 shows how this works for one of the bugs. But it will work for all bugs as long as each bug leaves a note behind describing their former color and arrow. The bugs' new arrow and color is a blend of the neighboring ones. The blend is a function of how close the bug is to its respective new neighbors. As you can see from Figure 6.7, the bug now has a different

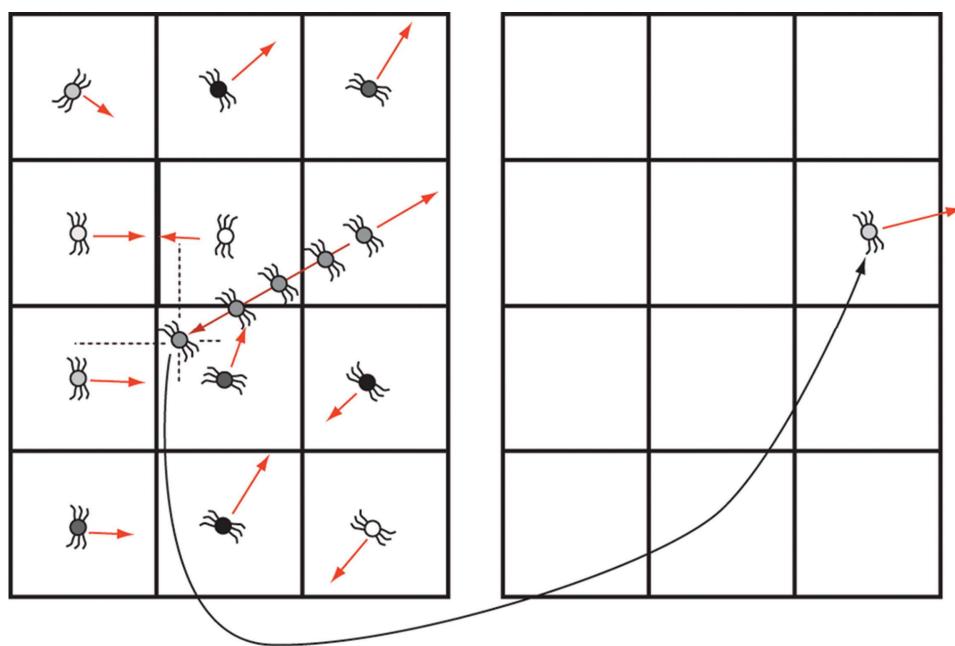


FIGURE 6.10 One bug crawls into neighboring apartments and blends both the arrows and the colors of the final neighbor. The bug then moves into its new apartment.

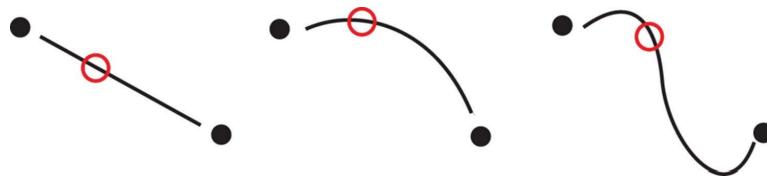


FIGURE 6.11 Three different ways to interpolate a point in between two known data points.

color and a different arrow. Mathematicians call this blending procedure *interpolation*.*

Interpolation just means: you have two values and you want to know the values in between. This is shown in Figure 6.11. The black dots are the known values and the red one is the unknown whose value we are trying to figure out. As can be seen in the figure, many options are available. We usually like the simplest one shown on the left of the figure, which is just a straight line. On the other hand, if information of the rate of change is available at the known points, then one can use a more clever interpolation scheme as shown in the middle and right-hand side of Figure 6.11. Really it is about relating the information at some particular points to points in between them.

That is how we get the unknown from the known.

For example, we have that

$$1.2 = 0.8 \times 1.0 + 0.2 \times 2.0$$

Which means that 1.2 lies between 1.0 and 2.0. 1.2 is closer to 1.0 than 2.0. Therefore the number that multiplies 1.0 is bigger than the number that multiplies 2.0. These numbers are called *weights*. Notice that the weights satisfy the following: $0.8 + 0.2 = 1.0$ and that they are always positive and smaller than 1.0. In fact, $0.8 = 1.0 - 0.2$. In this case, we used linear interpolation as shown on the left-hand side of Figure 6.11. In general, for linear interpolation we have that

$$\text{value in between} = (1 - \text{distance to 1}) \times \text{value 1} + (\text{distance to 1}) \times \text{value 2}.$$

In the case of our bugs that crawl in 2D apartments they end up with four neighbors: b_1 , b_2 , b_3 , and b_4 . Now we have a 2D problem and the interpolation is as follows:

* When classmates and I first heard this word in a math class in Geneva, Switzerland, we all chuckled. Whoa INTERPOL? Cool like in *International Police*.

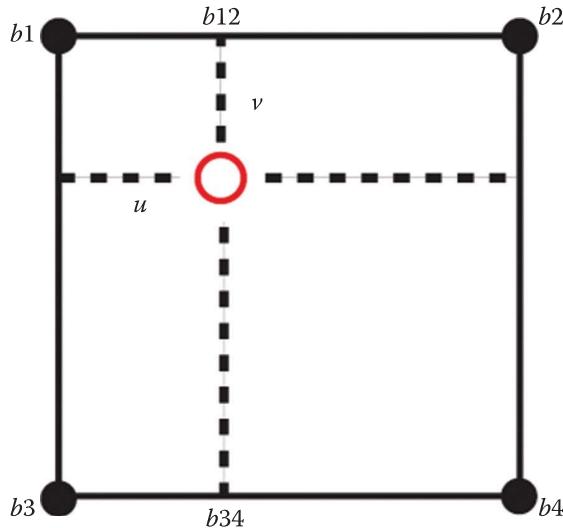


FIGURE 6.12 Linear interpolation in two dimensions.

$$\text{bug color} = (1 - u) \times (1 - v) \times b1 + u \times (1 - v) \times b2 + (1 - u) \times v \times b3 + u \times v \times b4.$$

This is shown in Figure 6.12. This formula is really an application of the 1D interpolation formula applied twice. In Figure 6.12, u denotes the horizontal distance to $b1$ and v denotes the vertical distance to $b1$. To see the link between 2D interpolation and 1D interpolation, we can introduce the following two intermediary 1D interpolations:

$$\begin{aligned} b12 &= (1 - u) \times b1 + u \times b2 \\ b34 &= (1 - u) \times b3 + u \times b4 \\ \text{Bug color} &= (1 - v) \times b12 + v \times b34 \end{aligned}$$

This procedure can be carried out in three dimensions as well using the same principle. Try it. In fact it works for any dimension. The 1D interpolation can be extended to any dimension through a process that goes by the fancy name of a *tensor product*. Coders will also notice that the second version is more efficient than the first one.

Why?

6 \times 's versus 8 \times 's, 3 $+$'s versus 4 $+$'s and 3 $-$'s versus 4 $-$'s. The less \times 's, $+$'s and $-$'s, the faster your code is. Usually the less code, the faster your program will run.* Software called *Optimizers* that reads your code will

* Yes, yes, yes, I know about loop unrolling. But Mister Optimizer automatically takes care of that. Can you imagine if you had to unroll your loops yourself?

sometimes take care of this for you. Optimizers take your code and output other code that supposedly runs faster. But remember an Optimizer is only as smart as the person who wrote it. That is an upper bound. Even worse is when Mister Optimizer creates code that does something different than you intended it to do: “Cool. It is faster but it does not do what my Pet wanted it to do.”

This can be the source of many so-called Heisenbugs named after Werner Heisenberg, one of the founders of quantum mechanics quoted previously. The Heisenberg principle states that the velocity and the position of a particle cannot be measured exactly at the same time in the quantum world.

The reason being is that to localize a particle you have to bombard it with some wave with a specific velocity. That is going to mess up the position. And similarly when measuring the velocity, you are going to mess up the velocity by bombarding it with particles. Bottom line: if you look at something it is going to affect whatever you are looking at. That is the Heisenberg principle in a nutshell.

In the world of computer coding, one sometimes gets a different behavior when Mister Optimizer is involved. Your code all of a sudden does not do what you want it to do. But as soon as you try to inspect the code: insert a breakpoint or add a print statement, the behavior is kosher again. This is frustrating since breakpoints and prints slow down the code. But anyone who has gone through this experience will understand what a *Heisenbug* is and the art it takes to deal with them. It helps to work with many smart coders who are just down the hall and who are willing to help you out.

To summarize: In a semi-Lagrangian framework, bugs leave their apartments to invade neighboring apartments following their arrows back in time. They then gather information through interpolation from their neighbors and change their colors and vectors accordingly. At the end, every bug is happy and lives in an asbestos-free and cheaper apartment.

Let’s now turn to another hybrid approach.

6.1.5 PIC

PIC stands for *Particle In Cell*. In our terminology it would be *BIA*: *Bugs In Apartments*.

PIC involves both bugs and apartments just like the semi-Lagrangian methods described earlier.

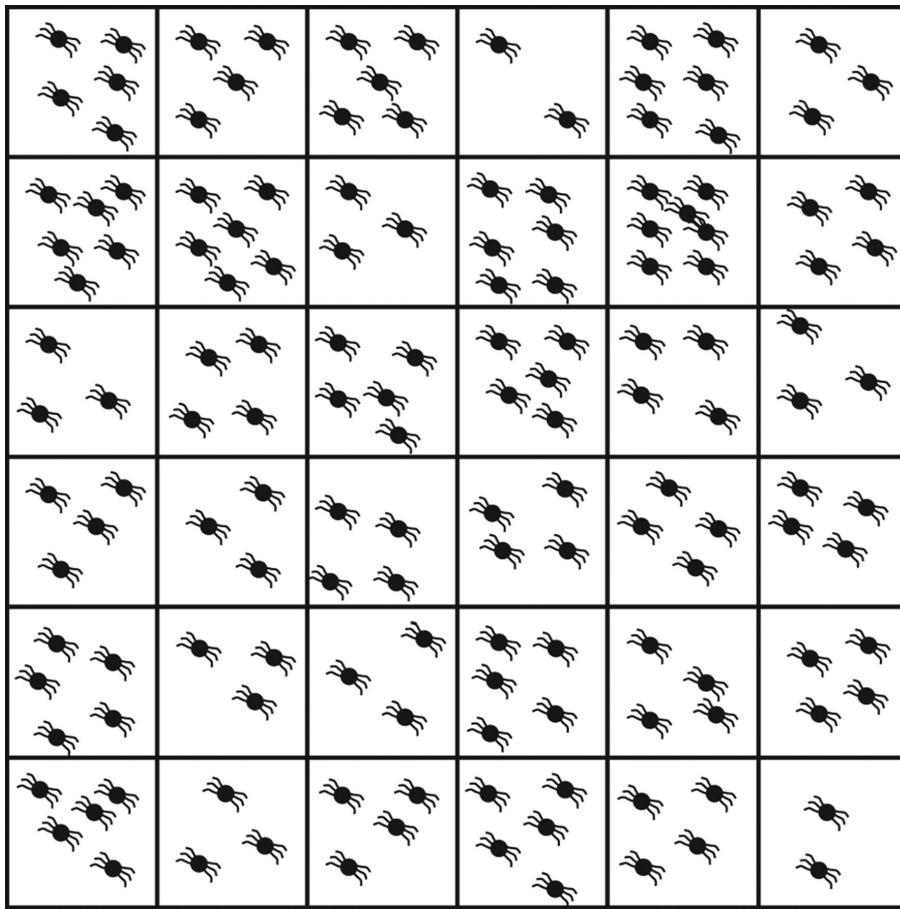


FIGURE 6.13 Bugs in cells.

Frank Harlow came up with the PIC method for fluids in the 1960s for (there was early work for plasmas). This method requires many bugs per apartment for it to work well. How many? Well that is up to the Master again who writes the code. Figure 6.13 shows the situation. Many bugs are now sharing apartments. Their number is not fixed just bounded by the Master. Some apartments have more bugs than others.

This is how PIC works.

Velocities are defined on the grid, either collocated or staggered. Just like in the grid-based methods we discussed previously. Each bug interpolates its velocity from its neighboring apartments. The bug then uses this velocity to wander to another apartment.

This time forward in time. No psychopathic landlord involved.

The bug then alters the velocity of its new neighbor apartments. This is depicted in Figure 6.14. We assigned a different color to each bug so they are easier to keep track of. The top left shows their initial configuration

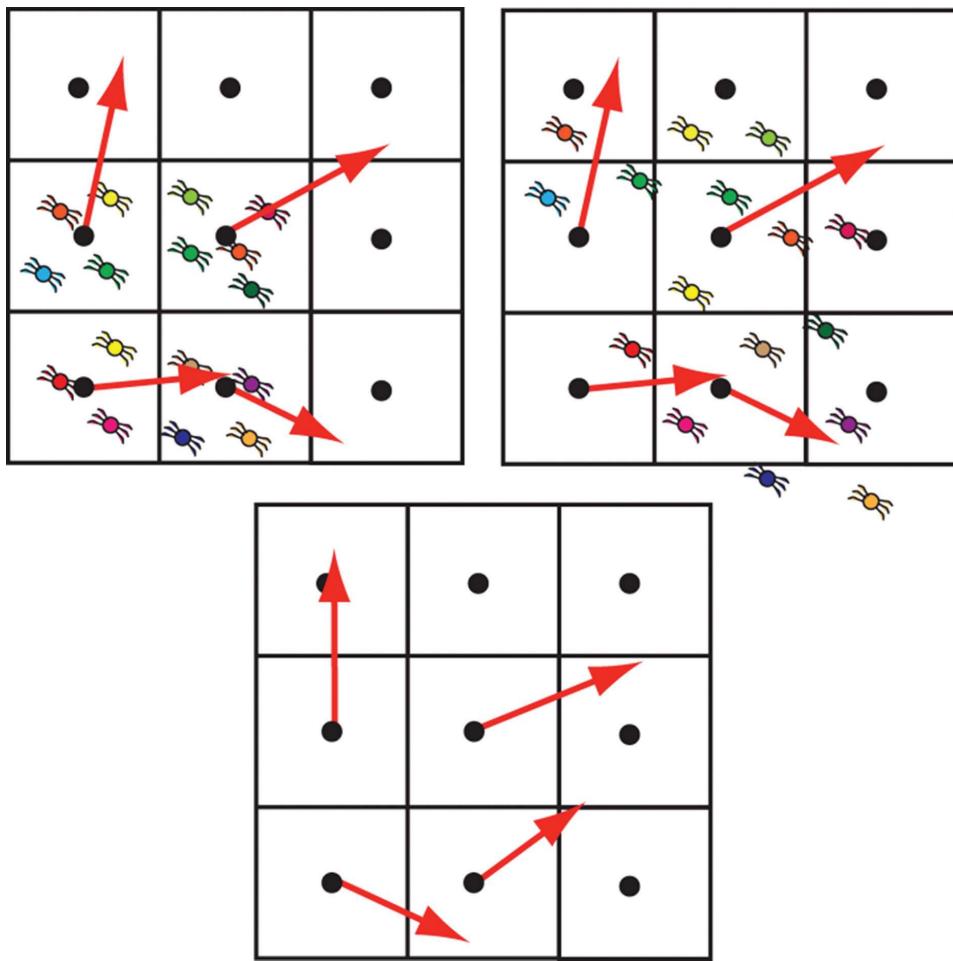


FIGURE 6.14 PIC's three steps.

with bugs in their apartments and their velocities. The top right shows their track to other apartments based on the velocities. And the bottom shows how the velocities are updated in each apartment due to the arrival of the new (or old) bugs with their new velocities. Of course new bugs (not shown) can come from apartments not depicted in Figure 6.14. They could come from the left wall, for example. Also two bugs have left the portion of the building that we are showing.

The bottom line is: we really only care about the velocities. You can think of the bugs as being little minions that carry around velocities from one place to another. However, in computer animation, the bugs can be useful to define the surface when simulating a liquid-like water. Unlike smoke, a liquid has a well-defined surface that separates the air from the liquid. Both are fluids but usually only the motion of the liquid is animated. So the bugs are floating and swimming in the liquid but not flying in the air.

This is kind of like the semi-Lagrangian approach except that there are more bugs and they move forward in time.

Again this is a transport problem: values in the apartments are updated using crawling bugs.

The grid, on the other hand, takes care of continuity (what flows in has to come out) and diffusion (stuff gets blurred). We will get to how to solve these effects on a grid below.

I have never implemented a PIC solver or its prettier cousin the FLIP solver. And of course there are many variants with their secret sauces. There is code to get you started on the Internet. Since I did not write it, I will not include it in this book.

To summarize: Moving bugs are great for accounting for the transport of a fluid, but grids are better for accounting for continuity and diffusion. By combining these techniques we get PIC.

Intermezzi



Intermezzo is an Italian word which stands for some sort of performance that fits in between the main acts of an opera or a concert. This usually gives people time to go to the washroom, stretch their legs, or get another drink. The other people who do not have these needs remain entertained by frivolous acts and performances. But sometimes they are more fun than the main acts.

The following are four *Intermezzi*. Enjoy!

7.1 INTERMEZZO UNO: LINEAR SYSTEMS



FROM DR. SEUSS'S BOOK: *ON BEYOND ZEBRA*

In order to understand the following fluid solver code, it is helpful to grasp the concept of a linear system of equations.

If you are familiar with this material feel free to skip this *Intermezzo* and go to the washroom.

If you are not familiar with linear systems, no worries linear systems are one of the simplest problems to solve in mathematics.

For starters I think anyone who is reading this book can solve this single linear equation without any effort

$$5 \times x = 7.$$

This equation asks which unknown thing x multiplied by 5 equals 7.

The solution for the unknown thing is obviously the following fraction:

$$x = \frac{7}{5}.$$

Now the unknown thing becomes a known thing.

This result can also be written as follows:

$$x = \frac{1}{5} \times 7 = (5)^{-1} \times 7.$$